Reproducing the Hensley & Draine (2023) Astrodust SED

1. Introduction

Calculation conditions. The astrodust driver uses the corrected MMP83 radiation field at $U = 1.585$ unless a field option is specified. The HD23 comparisons that select `hd23 logU=0.2` use the piecewise HD23 field at $U = 10^{0.2}$ and an unscaled 2.725 K CMB. These fields differ in spectral shape, chiefly through the 3000 K dilution factor. Figures made by `docs/make_figs.py` use the branch's default field; HD23-field calculations are identified separately where they are used. The temperature solver includes CMB heating through the radiation field and uses gross radiative emission for cooling.

Status of the measurements (2026-10-31). The emission comparisons and energy budgets in this memo were remeasured on 2026-10-31 from the 84 tracked `sed/output/sed_*.dat` products and the three `sed_astrodust_hd23_{S1,S2,PAH}.dat` files, all regenerated that day by `calc_sed.x` built with gfortran 16.2.0 from the settled source, with `docs/make_figs.py` and the same band and energy-budget conventions. The induced-emission comparisons (Tables 3 and 4) and the scaling test of §7.1.1 were measured the same day from runs of the same source whose output is not tracked; those runs reproduce, byte for byte, every tracked astrodust product they share with the branch. The execution-cost timings and the independent $T_{\rm eq}$ check were also remeasured on 2026-10-31. A number that describes an earlier state of the code and could not be remeasured is marked as such where it occurs.

Energy budget revisited. The conclusion of 2026-06-12, carried by earlier versions of this memo, that the HD23 release emits more than its own released $C_{\rm abs}$ absorb, is withdrawn. The apparent excess came from our input field: `radfield.f90 :: J_Mathis` follows Draine's 2008 `mmp_radiation_field.f`, which corrected the 4000 K dilution factor but kept the Mathis 1983 $W = 4 \times 10^{-13}$ at 3000 K, whereas HD23 eq. (9) uses the Draine (2011) value 7×10^{-13}. The field setting `hd23` (routine `J_HD23`) implements eq. (9) with the Mathis 1983 piecewise form. With it the release conserves energy and our astrodust SED matches the release in the FIR (§7.1.2, Tables 6 and 7). Those numbers were remeasured on 2026-10-31 from `sed/output/sed_astrodust_hd23_{S1,S2,PAH}.dat` (gfortran, NT = 200, heuristic solver) with the band conventions of `docs/make_figs.py`. Elsewhere this memo quotes the default `mathis` field unless it says otherwise.

This memo documents the construction of an astrodust thermal emission calculation that reproduces the Hensley & Draine (2023, ApJ 948, 55; hereafter HD23) "astrodust" model from first principles, working from the published dielectric function (Draine & Hensley 2021a, ApJ 909, 94, hereafter DH21a) and size distribution. The deliverable is a Fortran library callable from a 3D radiative-transfer driver per cell, with the emission per H atom expressed in HD23 units (erg s^{-1} sr^{-1} per H atom).

The calculation consists of five phases:

A. Random-orientation T-matrix sweep on the DH21 native grid (169 sizes $\times$ 1762 wavelengths, 1×10^{-4} μm to 3.981×10^{4} μm) producing $Q_{\rm abs}(\lambda, a)$, $Q_{\rm sca}(\lambda, a)$, and the asymmetry parameter g. Longward of 0.0912 μm the axis is DH21a's own grid with 200 samples in each decade, 1129 points. Shortward of it the axis is the DH21 dielectric function's own energy nodes, 632 points from 14.0 eV to 12 398 eV, so the absorption edges tabulated

there stay exact steps between adjacent grid points instead of being averaged into ramps. Each is a pair of nodes about 10^{-4} apart in relative energy, across which Im(n) jumps by +211% at Fe K (7124 eV), +131% at O K (538 eV), +67% at Fe L (724 eV) and +30% at C K (291 eV). One point is not a node: $0.0912\,(1-10^{-4})$, which is interpolated and is placed to resolve the Lyman-limit step of the *radiation field* rather than anything in the grain (§11.).

B. Optical-depth verification: $\tau_{\rm Ad}/N_H = \sum_a (dn_{\rm Ad}/N_H)\,\pi a^2\,Q_{\rm ext}$ compared against HD23 `extinction.dat`.

C. Two-stage enthalpy: silicate-only (Stage 1) and silicate + carbonaceous volume-weighted (Stage 2).

D. Single-cell SED solver with dynamic T grid per radius.

E. SED verification against the HD23 release answer key.

Two refinements were implemented in addition: PAH emission via the Draine & Li (2007) DL07 cross sections, and polarization via the dichroic-extinction formula of HD23 §2.2. A Python sidecar of the solver (Phase G) is provided for downstream prototyping.

The maximum sampled astrodust extinction residual is 0.47%. The polarization comparison has mean absolute residual below 0.07% in each reported band and maximum sampled residual 0.24%. With the default `mathis` field, which keeps Draine's 2008 3000 K dilution factor, the SED comparison falls 7.4% below the HD23 release in the 30–300 μm FIR-peak band (median relative residual; 9.8% at the 133 μm peak itself) after applying the corrected Mathis 4000 K dilution factor (1.65×10^{-13} instead of the literal Mathis 1983 1×10^{-13}; the original-Mathis variant gives 14.2%). This residual is now *resolved* (§7.1.2), and its source is our input field. HD23 eq. (9) uses the Draine (2011) dilution factors $W = 7\times10^{-13}$ (3000 K), 1.65×10^{-13} (4000 K) and 10^{-14} (7500 K); the default field keeps the Mathis 1983 4×10^{-13} at 3000 K. With the HD23 field, taken piecewise as in the Mathis 1983 construction (`./calc_sed.x astrodust hd23`), the released $C_{\rm abs}$ absorbs what the released `irem` tables emit (emitted over absorbed power 1.000 for astrodust and 1.015 for PAH at $\log U = 0.2$), and our astrodust SED matches the release to −0.13% in the 30–300 μm median (–0.69% at the 133 μm peak; PAH +0.32%). The earlier reading of this audit (2026-06-12), that the release emits 9.3% (astrodust) and 3.3% (PAH) more than its own $C_{\rm abs}$ absorb, used our default field and is withdrawn.

2. Reference parameters

Table 1 lists the HD23 best-fit parameters that we hold fixed throughout this work. The first three lines define the dielectric function file `index_DH21Ad_P0.20_0.00_1.400`; the rest determine size grids, dust mass, and the heating field.

3. Optical cross sections

3.1 Two-track cross-section architecture

HD23 treats the astrodust and PAH populations on *separate* tracks for their optical properties, and our calculation mirrors that separation verbatim:

Table 1: Reference parameters (HD23 best fit).

Quantity	Value	Source
Porosity $\mathcal{P}$	0.2	HD23 §3.1
Metallic Fe fraction $f_{\rm Fe}$	0	HD23 §3.1
Axial ratio b/a (oblate)	1.4	HD23 §3.1
Solid mass density $\rho_{\rm solid}$	3.42 g/cm^3	DH21a Tab. 2
Bulk mass density $\rho_{\rm Ad}$	2.74 g/cm^3	= $\rho_{\rm solid}(1-\mathcal{P})$
$M_{\rm Ad}/N_H$	0.0092	DH21a Tab. 2
Heating intensity U	$10^{0.2} \approx 1.585$	HD23 comparison; branch default $U = 1.585$
Radiation-field shape	piecewise HD23 field; corrected MMP83 branch default	HD23 §2.3.1 and §7.1.2
$a_{\rm eff}$ range	3.16×10^{-4} µm to 5.012 µm	DH21a, 169 points
λ range	1×10^{-4} µm to 3.981×10^{4} µm	Q-table grid, 1762 points (§11.)
above the Lyman limit	0.0912 µm to 3.981×10^{4} µm	DH21a, 200 per decade, 1129 points
below it	1×10^{-4} µm to 9.119×10^{-2} µm	DH21a dielectric-function nodes, 633 points

- **Astrodust** (HD23 §3.1) is a single composite oblate spheroid with the DH21a dielectric function. Its extinction, scattering, and polarization cross sections are computed via the modified-particle-form approximation (MPFA, three orientations of $C_E(0°)$, $C_E(90°)$, $C_H(90°)$ per HD23 §2.2) of the random-orientation T-matrix code of Mishchenko et al. (1996). The result is the `q_DH21Ad_*` table used throughout Phases A–E and F-2. Astrodust is the only population that contributes to polarized extinction or emission.
- **PAH** (HD23 §3.2) is treated as spherical and compositionally distinct ($\rho_{\rm PAH}$ = 2.0 g/cm^3 vs $\rho_{\rm Ad}$ = 2.74 g/cm^3). Its cross sections follow Draine & Li (2007) for the neutral and ionized branches, blended with the random-orientation graphite Mie cross section $C^{\rm gra}$ via the size-dependent weight $\xi_{\rm gra}(a)$ of HD23 eq. 16. No T-matrix; spherical Mie throughout. HD23 explicitly assumes PAH does not contribute to polarized extinction or emission.

The two-track structure is propagated through the rest of the calculation: astrodust uses the T-matrix Q tables and the full MPFA, PAH uses the DL07 cross sections (Sect. 7.2). The present section is the astrodust track. The PAH track lives in Sect. 7.2.

3.2 T-matrix engine and parallel sweep

We use the Mishchenko et al. (1996) random-orientation T-matrix code. It entered the tree as `tmd.lp.f` wrapped as a callable subroutine `TMD_ONE(AXI, LAM, MRR, MRI, EPS, NP, DDELT, NDGS, QEXT, QSCA, WALB, ASYMM, IERR)`, a mechanical conversion of the original main program: inline test-case inputs became arguments, the size-distribution loop was removed (one a per call), file I/O was dropped, and `STOP` statements were replaced with `IERR / RETURN`. A critical footgun: the original main mutates the input `DDELT` (`DDELT = 0.1*DDELT`); when called from F90 with a `parameter` constant, the write traps as SIGBUS on a read-only page. The wrapper copied to a local `DDELT_LOC` before mutating.

That f77 driver has since been rewritten as a modern-Fortran library under `tmatrix/src/`, and the entry point is now `tmatrix_eval(work, a_um, lambda_um, nr, ki, options, result)`. The only fixed-form file left is `lpd.f`, the linear-algebra routines the solver calls. The library holds no `COMMON`, no `EQUIVALENCE` and no mutable `SAVE`: every working array lives in a caller-owned `tmatrix_workspace_t`, one active caller for each workspace, and distinct workspaces may be evaluated concurrently. The convergence tolerance is a component of an `intent(in) tmatrix_options_t`, so the `DDELT` hazard above cannot recur. Mishchenko's unmodified sources (`tmd.lp.f`, `tmd.par.f`, `ampld.lp.f`, `ampld.par.f`) are kept under `tmatrix/reference/upstream/` so that each migrated routine stays checkable line by line against its original; nothing there is compiled. The transition is numerically inert: the Q tables it produces are byte-identical across it.

Convergence-failure fallbacks are implemented for $x = 2\pi a/\lambda < 0.1$ (small-x spheroidal Rayleigh polarizability ported from Draine 1992 `sphero.f`) and $x > 50$ (geometric optics, $Q_{\rm ext} \to 2$). The small-x branch uses the exact oblate-spheroid depolarization factors L_a, $L_b = (1-L_a)/2$ averaged over random orientation, and was found to remove a systematic 6% under-prediction that a Mie sphere used in its place produces in the dipole limit.

The timings that follow were taken on the 1129-wavelength axis, before the grid was carried into the ionizing band; the sweep on the present 1762-wavelength axis has not been timed, and the 633 wavelengths it adds all fall in the ultraviolet slice, which is the expensive one. The full 169 by 1129 sweep takes 11 minutes single-threaded on Apple Silicon. A bash work-queue driver `run_parallel.sh` dispatching $K = 32$ chunks across $N = 6$ workers reduces this to 2 minutes 18 seconds (4.8× speedup) with bit-identical output (Tab. 2). Equal-jw-range partitioning ($K = N$) gives almost no speedup because the UV slice is T-matrix-bound and stalls the run; the work-queue pattern is the essential ingredient.

Table 2: Parallel-sweep walltime measurements, on the 1129-wavelength axis.

mode	N	K	walltime	speedup
serial	1	1	11 m 0 s	1.0×
equal partition	6	6	9 m 59 s	1.1×
work queue	6	32	**2 m 18 s**	**4.8**×

3.3 Result

Figure 1 shows $Q_{\rm ext}$ and $Q_{\rm abs}$ at four representative grain sizes spanning the full DH21 range. The classic transition from Rayleigh-like dipole behavior at small x (small grains or long wavelengths) to the geometric-optics plateau at large x is visible. Verification of $Q_{\rm abs}$ against the HD23 release table `q_DH21Ad_P0.20_Fe0.00_1.400.dat.gz`, taken against the mean of its jori = 1, 2, 3 columns, gives a median deviation of −0.04% and a mean absolute deviation of 0.33% across the T-matrix-converged region ($x \in [0.01, 30]$). (That mean is not the random-orientation average, which is $\frac{1}{3}Q^{\rm jori=2} + \frac{2}{3}Q^{\rm jori=3}$; the two coincide only in the Rayleigh limit, where jori 1 and jori 3 measure the same $Q_\perp$.) The largest deviations, a few percent, sit at the top of that x range where the geometric-optics limit takes over, and carry negligible weight in any SED integral.

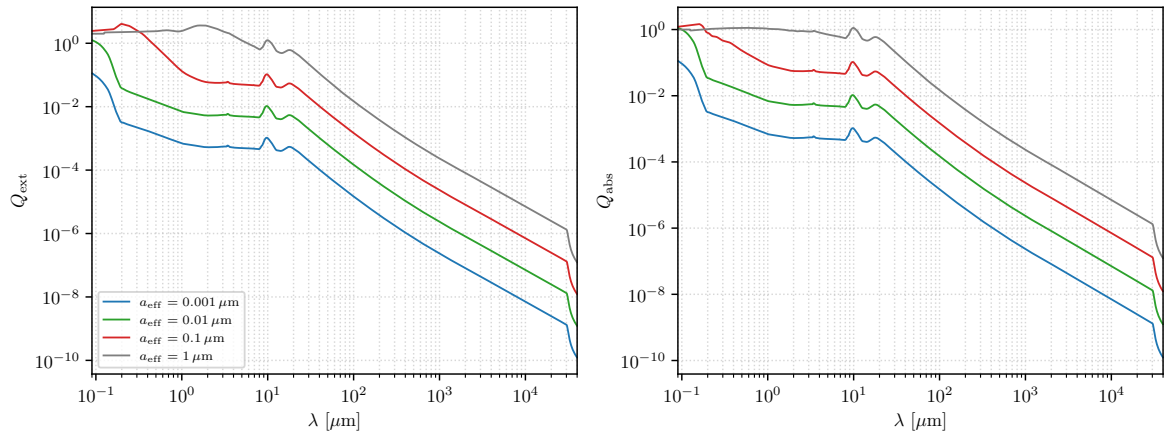


Figure 1: Extinction (left) and absorption (right) efficiency factors as functions of wavelength for four representative grain sizes from the DH21 grid. The transition from $Q \propto x^4$ scattering at long wavelengths to $Q \to 1$ geometric optics at short wavelengths is clearly resolved for each size.

4. Optical-depth verification

The first end-to-end milestone is to verify that our cross sections reproduce HD23's reported extinction per H atom:

$$\frac{\tau_{\rm Ad}}{N_H}(\lambda) = \sum_a \frac{dn_{\rm Ad}}{N_H}(a)\, \pi a^2\, Q_{\rm ext}(\lambda, a), \tag{1}$$

where $(dn_{\rm Ad}/N_H)(a)$ is the HD23 size distribution (`size_distribution.dat` column 2), already integrated over the local bin width.

To evaluate eq. 1, the figure generator `docs/make_figs.py` loads the full Q table, interpolates $Q_{\rm abs}$ and $Q_{\rm ext}$ in $\log a$ from the 169-point DH21 grid onto the 167-point size-distribution grid (the same rule as the code's `q_table.f90`), sums, and interpolates onto the HD23 release wavelength grid (`wav.dat`, 1000 points). The same calculation is run for the scattering channel against `scattering.dat`.

The maximum |relative residual| across the entire grid is **0.47%** for $\tau_{\rm Ad}/N_H$ (at 26.84 µm) and **0.68%** for $\tau^{\rm sca}_{\rm Ad}/N_H$ (at 31.63 µm). Figure 2 shows the residual is sub-percent everywhere and structurally featureless. The acceptance criterion

(1%) is met by a comfortable margin, confirming that our Q sweep and size-distribution integration are both correct.

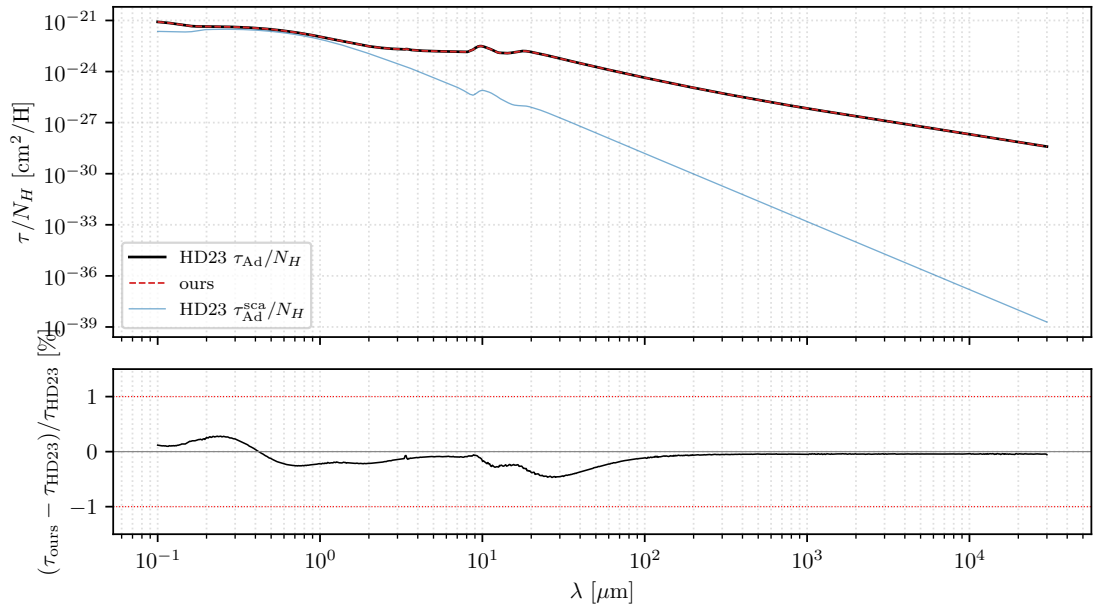


Figure 2: Top: HD23 $\tau_{\rm Ad}/N_H$ (black solid) and our value (red dashed), and HD23 $\tau^{\rm sca}_{\rm Ad}/N_H$ (blue). Bottom: relative residual on $\tau_{\rm Ad}/N_H$ in percent. The acceptance band ($\pm 1\%$) is shown in dotted red.

This is the strongest single piece of evidence that the optical-property calculation is sound, since it is purely a function of $Q_{\rm ext}$, the size distribution, and the cross-section conversion $C = \pi a^2 Q$, and is completely independent of any thermal or enthalpy assumption.

4.1 Comparison against the WD01 / Draine 2003 model

The HD23 astrodust+PAH model differs in both composition and size distribution from the older Weingartner & Draine (2001, WD01) / Draine (2003) "silicate + carbonaceous" model that has been the working dust model in the MoCafe radiative-transfer code. To make the cross-section impact of moving from one model to the other concrete, we compute the three quantities that drive scattering in an isotropic medium – the extinction cross section per H, the albedo $C_{\rm sca}/C_{\rm ext}$, and the scattering-weighted asymmetry parameter $g = \langle \cos\theta \rangle$ – for both models on the WD01/D03 wavelength grid and overlay them (Figure 3).

The D03 curves come from B. T. Draine's published tables `kext_albedo_WD_MW_3.1_60_D03.all_{2003,2009}`, columns "albedo", "<cos>", and "C_ext/H". The 2003 and 2009 versions on Draine's website are nominally the same model but differ in some underlying grain parameters that are not explicitly documented; the 2009 version has $M_{\rm dust}/N_H = 1.398 \times 10^{-26}$ g/H, revised from 1.870×10^{-26} g/H in 2003. We plot both for reference. The HD23 extinction and scattering cross sections per H are the "total" (astrodust+PAH) columns of the release files `extinction.dat` and `scattering.dat`; the HD23 asymmetry parameter is built from our Phase-A T-matrix output (`q_astrodust_*`, columns $Q_{\rm sca}$ and g_a per (a, λ)) by the scattering-weighted size integral

$$g(\lambda) = \frac{\sum_a (dn_{\rm Ad}/N_H)\, \pi a^2\, Q_{\rm sca}(a,\lambda)\, g_a(a,\lambda)}{C^{\rm sca}_{\rm tot}(\lambda)}, \tag{2}$$

where the denominator is the AD+PAH total scattering cross section from `scattering.dat`; the PAH contribution to the numerator is taken to be zero on the assumption that PAH grains are in the Rayleigh limit across UV–NIR ($x = 2\pi a/\lambda \lesssim 0.1$ for $a \leq 30$ Å and $\lambda \geq 0.1$ μm), so $g_{\rm PAH} \approx 0$ and PAHs enter only through their (small) scattering cross section in the denominator. We restrict the plot to 0.1 µm to 5.0 µm (UV through K band) per the user request.

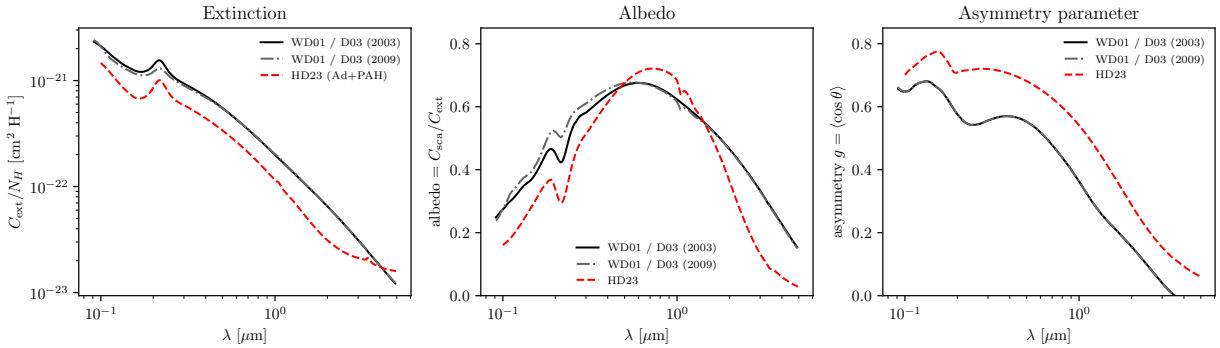


Figure 3: Extinction cross section per H (left), albedo (middle), and scattering-weighted asymmetry parameter (right) of the HD23 astrodust+PAH model (red dashed) overlaid on the WD01/Draine silicate+carbonaceous model in both the 2003 (black solid) and 2009 (gray dot-dashed) revisions, over the UV–NIR range.

Observations:

- **WD01/D03 2003 vs 2009.** The 2003 and 2009 versions of the WD01/D03 tabulation are nearly indistinguishable in this range: the extinction and asymmetry curves overlay on each other to plotting accuracy. The only visible revision is a modest upward shift of the FUV albedo (around the 2175 Å bump). Whatever physical revision produced the 2009 numbers, it changes the dust mass per H ($M_{\rm dust}/N_H$ dropped by 25%) but not the wavelength shape of the optical properties per H atom at the percent level.
- **Extinction:** all three curves reproduce the 2175 Å bump, the FUV rise, and a roughly power-law optical/NIR tail. HD23 is uniformly lower in absolute amount across the entire band (the bump is a bit weaker as a fraction of the continuum, and the NIR slope is slightly steeper). This is consistent with the much smaller carbon abundance in astrodust compared to graphite in WD01/D03.
- **Albedo:** the qualitative shape – low in FUV, peaking in the optical near 0.6 µm, then falling toward the NIR – is the same in all three models, with peak values ≈ 0.65–0.72. HD23 peaks slightly higher and drops off slightly faster.
- **Asymmetry parameter:** HD23 is uniformly more forward-scattering than D03 across the band, by $\Delta g \sim 0.1$–0.2. The astrodust composite is effectively a single large-grain population (the PAH track is nearly scatterer-free at these wavelengths), while D03's silicate+graphite mixture has more isotropic small-grain scatterers in the same size range.

For the MoCafe RT context, the practical consequence of swapping in the HD23 dust is a $\sim$ 20–30% reduction in τ per H in optical/NIR (visible directly in the left panel) and a more strongly forward-peaked phase function (visible in the right panel) – both qualitatively the same direction relative to either D03 vintage.

5. Enthalpy

The grain enthalpy $U(T, a)$ enters the stochastic heating solver through the energy spacing $\Delta H = H(T_{i+1}) - H(T_i)$ between adjacent temperature bins. Since the closed-form Debye sums of Draine & Li (2001, hereafter DL01) cost microseconds per call, we use the module functions `enthalpy_S1(T, a, prefactor)` and `enthalpy_S2(T, a)` directly in the inner SED loop; the `calc_enthalpy.x` table-dumping driver is retained as a sanity tool and is *not* a required intermediate.

5.1 Single-composition closed forms (DL01)

For both graphite (`Car0`/`Car1`) and silicate (`Sil`) sub-components we use the closed-form Debye-2 internal energy of DL01 §B (their eqs. 5–7). Define the Debye integral of order n as

$$f_n(x) \equiv n x^{-n} \int_0^x \frac{y^n}{e^y - 1}\, dy, \qquad x = \Theta/T, \tag{3}$$

so that $f_n \to 1$ for $x \to 0$ (high T) and $f_n \to 0$ for $x \to \infty$ (low T). For a grain containing $N_{\rm atom}$ atoms, the vibrational internal energy is

$$U_{\rm vib}(T, N_{\rm atom}) = (N_{\rm atom} - 2)\, k_B T \sum_m g_m\, f_2(\Theta_m/T), \tag{4}$$

with the $(N_{\rm atom} - 2)$ Debye-modes factor (DL01 eq. 5) and the composition-specific mode set $\{(\Theta_m, g_m)\}$.

Graphite (DL01 `Car0`, `Car1`). Two Debye-2 modes (DL01 eq. 6):

$$U_{\rm Car}(T, N_{\rm atom}) = (N_{\rm atom} - 2)\, k_B T \left[f_2(\Theta_{\rm op}/T) + 2\, f_2(\Theta_{\rm ip}/T) \right], \tag{5}$$

with $\Theta_{\rm op} = 863$ K (out-of-plane) and $\Theta_{\rm ip} = 2504$ K (in-plane). The atom count is $N_{\rm atom}(a) = 4.68208 \times 10^{11}\, a^3\, [\mu{\rm m}^{-3}]$ at the standard graphite density $\rho = 2.24$ g/cm^3. For grains with $N_{\rm atom} \lesssim 10^4$ we add the three DL01 C–H stretch modes

$$U_{\rm CH}(T, N_{\rm atom}) = (N_{\rm H}/N_{\rm atom}) N_{\rm atom}\, k_B \sum_{j=1}^{3} \frac{h\nu_j}{\exp(h\nu_j/k_B T) - 1}, \tag{6}$$

with $h\nu_j/k_B = (hc/k_B)/\lambda_j$ for $\lambda_j \in \{11.3, 8.6, 3.3\}\, \mu$m and the size-dependent H/C ratio of DL01 eq. 7 ($H/C = 0.5$ for $N_{\rm atom} \le 25$, $H/C = 0.5\,(25/N_{\rm atom})^{1/2}$ for $25 < N_{\rm atom} < 100$, $H/C = 0.25$ otherwise).

Silicate (DL01 `Sil`). Two Debye-2 modes plus a Debye-3 (DL01 eq. 9), extrapolated above $T = 500$ K to match laboratory data (Bonjour, Leger, & Hekinian 1984):

$$U_{\rm Sil}(T, N_{\rm atom}) = (N_{\rm atom} - 2)\, k_B T \left[2\, f_2(\Theta_2/T) + f_3(\Theta_3/T) \right], \tag{7}$$

with $\Theta_2 = 500$ K and $\Theta_3 = 1500$ K. The atom count is $N_{\rm atom}(a) = 7 \times 5.10401 \times 10^{10}\, a^3\, [\mu{\rm m}^{-3}]$ (seven atoms per formula unit $MgFeSiO_4$ at $\rho = 3.5$ g/cm^3). Above $T = 500$ K we use the piecewise-polynomial form of DL01 eq. (9) so that $U_{\rm Sil}$ matches the measured Si_2O fits to higher accuracy.

5.2 Stage 1: silicate-only

Treat the whole astrodust grain as if it were DL01 astrosilicate. Two prefactor options:

- **C-1 (literal):** $n_{\rm form}(a) = 5.1040126 \times 10^{10}\, a^3$ (DL01 astrosilicate at $\rho = 3.5$ g/cm^3).
- **C-2 (density-corrected):** $n_{\rm form}(a) = (2.74/3.5)\, n_{\rm form,C\text{-}1}(a)$, matching the bulk astrodust density $\rho_{\rm Ad} = 2.74$ g/cm^3.

Sanity check: the ratio $U_{\rm S1,C2}(T,a)/U_{\rm S1,C1}(T,a)$ saturates at $0.7829 = 2.74/3.5$ for grains large enough that the $(N_{\rm atom}-2)$ correction is negligible. Confirmed across the full (T,a) grid.

5.3 Stage 2: silicate + carbonaceous, volume-weighted

Decompose the astrodust *solid* (the $(1-\mathcal{P})$-fraction of the grain volume) into a carbonaceous sub-volume $V_{\rm car}$ and a silicate-like sub-volume $V_{\rm sil}$, with $V_{\rm sil}/V_{\rm solid} = 1 - f_C^{\rm vol}$. The astrodust enthalpy is the additive sum

$$U_{\rm Ad}(T,a) = U_{\rm Sil}(T; N_{\rm atom}^{\rm Sil}(a)) + U_{\rm Car}(T; N_{\rm atom}^{\rm Car}(a)), \tag{8}$$

with sub-component atom counts

$$N_{\rm atom}^{\rm Sil}(a) = (1-\mathcal{P})(1-f_C^{\rm vol})\, n_{\rm Sil}^{\rm bulk}\, V(a), \quad N_{\rm atom}^{\rm Car}(a) = (1-\mathcal{P}) f_C^{\rm vol}\, n_{\rm Car}^{\rm bulk}\, V(a), \tag{9}$$

where $V(a) = (4\pi/3)\, a^3$, $n_{\rm Sil}^{\rm bulk}$ and $n_{\rm Car}^{\rm bulk}$ are the bulk atom densities of DL01 astrosilicate and DL01 graphite respectively. "Everything not carbonaceous is treated as silicate".

The carbonaceous volume fraction $f_C^{\rm vol}$ is derived from DH21a Table 2 (column $f_{\rm Fe} = 0$): $V_{\rm car} = 0.86 \times 10^{-27}$ cm^3/H, $V_{\rm Ad}(1-\mathcal{P}) = 4.46 \times 10^{-27}$ cm^3/H, hence $f_C^{\rm vol} = 0.193$. The result is insensitive to $f_{\rm Fe}$ ($f_C^{\rm vol}(f_{\rm Fe}{=}0.10) = 0.196$).

6. SED solver

6.1 Module API

The solver lives in `sed_astrodust_mod` and exposes a two-step API designed for a cell-by-cell 3D-RT driver:

```
call sed_init(qtable_path, NT, T_lo, T_hi)
...
do icell = 1, ncells
   ... build J_lam in this cell ...
   call sed_solve    (J_lam, enthalpy_stage, lamI_lam)
   call sed_solve_pah(J_lam,                 lamI_lam_pah)
end do
```

The `sed_init` call loads the precomputed Q table and the size distribution, builds the wide T grid ($N_T = 200$, log-spaced 2.7 K to 5000 K), and tabulates the four cell-independent quantities $C_{\rm abs}(\lambda,a)$, $\kappa_B(T,a)$, $U(T,a)$ (for all three enthalpy stages, plus the PAH carbonaceous prescription), and $\kappa_{\rm CMB}(a)$. Single-cell `sed_solve` loops over the size distribution, calling `calc_Teq` once per radius and either branching to equilibrium emission $C_{\rm abs} B_\lambda(T_{\rm eq})$ or solving for the stochastic temperature distribution $P(T)$ via the Guhathakurta & Draine (1989) matrix method.

6.2 Dynamic temperature grid

A key empirical result, learned the hard way from SED_v5: the temperature distribution $P(T)$ for small grains is *very sharp*, and a single static T grid spanning the full [2.7, 5000] K range cannot resolve it. DL07 Fig. 4 illustrates this directly: $P(T)$ for a 100 Å grain at $U = 1$ drops by more than twelve orders of magnitude over less than a factor of two in T. A coarse or mis-centered grid simply returns numerical zeros across the entire range, mimicking "no stochastic contribution" when in reality the grain is sharply peaked at $T_{\rm eq}$.

Grain-loop direction matters. Both algorithms iterate over the size distribution in the *small-to-large* direction (`do ir = 1, NA` with `aeff` sorted ascending). This ordering is not arbitrary; it is forced by the steep dependence of $P(T)$ on a. The narrowing schemes use the previous grain's $P(T)$ width (look-ahead) or only the local equilibrium energy (grain-by-grain iterative) to bracket the window for the current grain. In both cases the natural progression is from a *wider* window at small a to a *narrower* window at large a:

- Small grains ($a \lesssim 30$ Å) have a broad $P(T)$ that extends from $T_{\rm CMB}$ up to several hundred K (DL07 Fig. 4). On the wide T grid the matrix solver `calc_P` is well-conditioned and yields a sensible $P(T)$ directly.
- Large grains ($a \gtrsim 100$ Å) have a $P(T)$ that is essentially a delta function at $T_{\rm eq}$. On the same wide T grid the Guhathakurta–Draine transition matrix becomes numerically degenerate (the entire steady-state probability falls inside one bin) and `calc_P` returns $P \equiv 0$ everywhere – a silent failure that mimics "no stochastic contribution".

The small-to-large loop bootstraps the narrowing from the well- conditioned end. Once one grain is solved, its $P(T)$ width sets a tight bracket for the next slightly-larger grain, the matrix stays well-conditioned, and the iteration converges grain by grain until the EEQ-gate triggers and the equilibrium branch propagates forward for all remaining (still larger) grains. The opposite direction would start on the degenerate large-grain calc_P, with no working previous grain to seed the window – a chicken-and-egg failure mode that we have hit while debugging (`mc_pT/main_pT_compare.x` returns $P \equiv 0$ for $a > 80$ Å when called on the wide grid without iterative narrowing, see Sect. 6.5 of the mc_pT report).
Two grid-narrowing algorithms are implemented in our solver and selectable at compile time via the `USE_DRAINE_NARROWING` module parameter:

- **Look-ahead heuristic (SED_v5 default).** For grain $i+1$, set the T grid to $[T_{\rm eq}\exp(-0.6\,\delta), T_{\rm eq}\exp(+0.6\,\delta)]$ where δ is twice the previous grain's $\ln P > \ln P_{\rm crit}$ half-width in $\ln T$. One `calc_P` call per grain. Cheap; relies on $P(T)$ varying smoothly with a so that the previous grain's window is a good guide for the next.
- **Grain-by-grain iterative refinement** (default, following Draine's method). Initial guess $U_{\rm max} = \max(13.6\ {\rm eV} + 2E_{\rm eq}, 13.65\ {\rm eV})$ (the single-UV-photon floor) and $U_{\rm min} \in \{0, E_{\rm eq}/5\}$ depending on $E_{\rm eq}/E_{\rm eq,SS}$ with $E_{\rm eq,SS} = 150$ eV; the endpoints are then refined iteratively until $P({\rm endpoint})/P_{\rm max} < 10^{-13}$, with 0.8/0.2 damping for shrinkage and $\times 1.2$ expansion (or bisection toward $U_{\rm max,hi}$). Typically converges in 3–5 iterations per grain. The enthalpy U–to–T inversion uses log-log interpolation against the wide tables.

The two algorithms give identical SED to within $\sim 0.5\%$ on band-integrated totals (see Sect. 7.1.1) on our HD23 reproduction case, confirming that the SED_v5 heuristic window happens to be accurate enough *for this radiation field*

and these enthalpy tables. Both algorithms remain in the source so that future work with very different J_λ shapes or size grids can switch and compare.

Three ancillary safety nets are common to both algorithms: $T_{\rm eq}$ is always computed on the *wide* grid (the narrowed grid may not bracket the next grain's $T_{\rm eq}$); a grain whose enthalpy at $T_{\rm eq}$ exceeds the EEQSS = 150 eV gate is solved as a δ-function in T at $T_{\rm eq}$ (steady-state limit); and on degenerate P (no $\ln P > \ln P_{\rm crit}$ anywhere) we fall back to equilibrium emission at the wide-grid $T_{\rm eq}$.

6.3 Execution cost

The astrodust driver initializes the optical and enthalpy tables once, then solves the three astrodust stages and the PAH population. A full single-cell solve (`sed_init` once, 3 enthalpy stages × 1 `sed_solve` + 1 `sed_solve_pah`), `./calc_sed.x astrodust`, takes 2.19 s of wall time (median of five runs) on one thread of an Intel Xeon Gold 6154 (3.0 GHz), gfortran 16.2.0 `-O3`. Through the library, `build_dust('astrodust')` on the 1129-wavelength grid takes 0.78 s once per process, and one `dust_emission` call for a cell (all populations, NT = 200, the default heuristic solver) takes 1.051 s (median of 21 calls). The heuristic solver runs the grains serially, so the time does not change with the number of OpenMP threads (measured with 1 and 8). The peak resident memory is 56 MiB for a process that builds the model and solves cells through the library and 43 MiB for `./calc_sed.x astrodust`. Measured 2026-10-31. The timings this memo carried before (2.02 s, 0.72 s and 0.98 s, 2026-10-22) were taken on the machine recorded then as an Intel Xeon Gold 6254 and are superseded.

7. SED verification

7.1 Astrodust-only SED

Figure 4 compares our astrodust SED for the three enthalpy stages against HD23 `astrodust_irem.dat` at $\log_{10} U = 0.2$ (column 33). The FIR peak shape is well reproduced and is essentially identical across stages (the peak is dominated by large equilibrium grains for which the enthalpy is irrelevant); stage-to-stage variations show up only in the 5 µm to 30 µm mid-IR where stochastic heating of small grains matters.

The band residuals below are for the S1_C1 baseline with the corrected MMP83 branch default. The number in parentheses after each one is the same statistic for the original-Mathis variant (`./calc_sed.x astrodust mathis_orig`, Table 5), retained to reproduce the earlier state:

- **300–3000 μm (submillimeter):** mean absolute residual 3.27% (5.99%), median −3.00% (−5.48%), and maximum 5.44% (10.00%).
- **30–300 μm (FIR):** mean absolute residual 10.31% (15.46%), median −7.37% (−14.18%), and maximum 23.34% at 34.99 µm (21.54% at 84.67 µm). The maximum is not at the peak: it sits on the short-wavelength shoulder, where stochastically heated small grains dominate and where the corrected Mathis field raises our curve above HD23. The residual at the released 133.38 µm peak is −9.80% for S1_C1, −10.47% for S1_C2, and −10.86% for S2. All three stages show the same peak systematic, so the issue is in the equilibrium emission (independent of

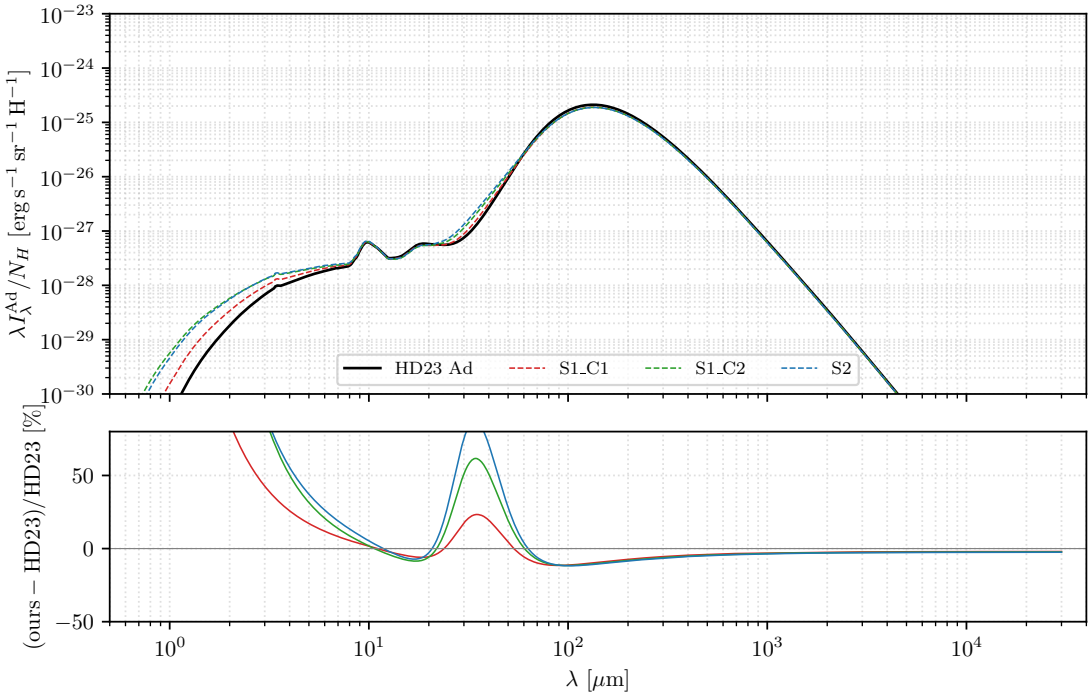


Figure 4: Astrodust thermal emission per H atom with the corrected MMP83 field at $U = 1.585$, compared with the HD23 released spectrum at $U = 10^{0.2}$. The heating fields have different spectral shapes. Top: our three enthalpy stages (S1_C1, S1_C2, S2) overlaid on HD23 `astrodust_irem.dat`. Bottom: relative residual in percent. The FIR peak shape is recovered; a 9.8% systematic under-prediction persists at the 133 µm peak (18.0% with the original Mathis field), and the <5 µm range over-predicts by a factor of 4.3 at 1 µm in this astrodust-only comparison, where PAHs dominate the HD23 spectrum; the combined Ad + PAH spectrum is shown in Fig. 6.

enthalpy), not in stochastic heating. The residual narrowed when we adopted the corrected Mathis 4000 K dilution (§7.1.1, *"Mathis 4000 K dilution correction"*).

- **5–30 μm (mid-IR):** mean absolute residual 6.03% (6.75%), median +1.53% (–0.77%), and maximum 17.07% (16.72%).
- **Above 3000 μm:** mean absolute residual 2.11% (3.88%), median −2.09% (–3.84%), and maximum 2.31% (4.21%).
- **Below 5 μm:** over-prediction by a factor of 4.3 at the 0.995 µm release grid point. PAHs dominate this band in HD23 and the astrodust-only contribution is small; the comparison is meaningful only in the full Ad + PAH form of Fig. 6.

7.1.1 Residual analysis: what is *not* causing the FIR-peak deficit

A natural first guess for the FIR-peak deficit is a ~ 0.2–0.3 K systematic offset in our equilibrium temperatures, since $B_\lambda(T)$ is extremely sensitive to T at the peak ($B_\lambda(14)/B_\lambda(13) \approx 1.84$ at $\lambda = 130$ µm). Three diagnostic ports have been

implemented to test this and the other obvious suspects; *all three came back clean*, so the residual is now isolated to a smaller suspect list.

- **Independent $T_{\rm eq}$ check.** The verification driver that first made this check compared our `calc_Teq` (precomputed $\kappa_B(T, a)$ table inversion) against an independent on-the-fly bracket–bisection–power-law solver following Draine's method (including the long-wavelength Rayleigh–Jeans tail extension under the $C^{\rm abs} \propto \nu^2$ assumption) and found rms 0.14% (max 0.21%) for astrodust and rms 0.11% (max 0.17%) for PAH. That driver is not in this tree. The check was remeasured on 2026-10-31 with a driver linked to `libsedust.a`: for every grain with $dn > 0$ in `build_dust('astrodust')`, under the field of `./calc_sed.x astrodust` ($U = 1.585$), it solves $\int C_{\rm abs} J_\lambda \, d\lambda = \int C_{\rm abs} B_\lambda(T) \, d\lambda$ by bisection, with the Planck integral taken directly on the model wavelength grid and continued beyond the last wavelength as a Rayleigh–Jeans tail under $C_{\rm abs} \propto \lambda^{-2}$:

  - Astrodust (161 grains): rms 0.0016%, max 0.0023% ($T_{\rm eq}$ range 11.8–19.3 K).
  - PAH, neutral (55 grains): rms 0.0020%, max 0.0028%; ionized (164 grains): rms 0.0013%, max 0.0027% ($T_{\rm eq}$ range 9.0–47.3 K).

  That is, the two solvers agree far below the 1–2% $T_{\rm eq}$ error a 0.3 K offset at 15 K would require. *Our $T_{\rm eq}$ is not the source of the FIR deficit.*

- **Grain-by-grain T-window narrowing (Draine v7 port).** The iterative refinement of §6.2 was ported and compared against the SED_v5 look-ahead heuristic. Band-integrated SED ratios change by at most 0.28% (the 1–5 μm band) between the two algorithms, and by at most 0.16% longward of 5 μm, with either the default or the HD23 field (remeasured 2026-10-31 from `./calc_sed.x astrodust draine` and `draine hd23`); the diagnostic counter shows that 81 astrodust grains in Stage S1 and 83 in Stage S2 process stochastically in the 2026-10-31 default heuristic run (70 and 73 in the iterative run); the rest enter equilibrium through the solver gate. The iterative refinement does change the T window for those grains – it just does not change the band-integrated SED. *Grid choice is not the source of the FIR deficit either.*

- **DL01 algorithm sensitivity.** DL01 §10/§11.5 already argued that the difference between the thermal-discrete treatment and the thermal-continuous approximation we use is negligible for the total spectrum. That statement is about *algorithmic* choice (matrix solver), not about grid choice; the latter is governed by DL07 Fig. 4 and has now been independently verified above.

What is left for the FIR-peak residual to live in is one of: (i) the Q-table interpolation from the 169-point DH21 size grid onto the 167-point size-distribution grid at sub-mm wavelengths; (ii) discretisation error in the single-grain $B_\lambda(T)$ integral that builds κ_B and $T_{\rm eq}$; or (iii) the size-distribution treatment itself (i.e., the kernel that maps `size_distribution.dat` onto our grain bins).

Grain-by-grain diagnostic and the random-orientation convention. We resolved this question by direct grain-by-grain comparison. The comparison (drawn by `docs/make_figs.py`) takes $C_{\rm abs}(\lambda, a_0) = \pi a_0^2 Q_{\rm abs}$ at a chosen grain radius a_0 directly off the random-orientation Q table — the same table `sed_init` feeds to `sed_solve`. At $a_0 = 0.126\,\mu$m

(the radius bin closest to where the FIR peak emission originates) we compared this against the HD23 release Q-table `q_DH21Ad_P0.20_Fe0.00_1.400.dat.gz`, which publishes $Q_{\rm abs}$ separately for the three Mishchenko orientations `jori=1` ($k \parallel a$), `jori=2` ($k \perp a$, $E \parallel a$), and `jori=3` ($k \perp a$, $E \perp a$). The random-orientation average for unpolarized emission is the $\frac{1}{3}Q^{\rm jori=2} + \frac{2}{3}Q^{\rm jori=3}$ combination.

Our small-x limit (`rayleigh_limit` in `tmatrix/driver/asymptotic_optics.f90`, used for $x < 0.1$, which covers all FIR-peak-relevant grains at sub-mm wavelengths) implements this random-orientation average exactly: the ratio of our $C_{\rm abs}(a_0 = 0.126\,\mu{\rm m}, \lambda)$ to $\frac{1}{3}Q^{\rm jori=2} + \frac{2}{3}Q^{\rm jori=3}$ on the HD23 release grid has mean 0.99992 and standard deviation 4×10^{-4} across 30 µm–30 000 µm. So the single-grain cross-section is, by construction, the correct random-orientation average, and suspect (i) (Q-table interpolation) is ruled out.

At the *same* grain, however, our $C_{\rm abs}$ is 0.8428 (= $1/1.1866$) times the release file's `jori=1` Q over 30 µm–30 000 µm – a flat ratio with $\sigma < 0.001$ across ~ 600 release-grid wavelengths. In the dipole limit `jori=1` ($k \parallel a$) measures only $Q_{\perp}^{\rm pol}$ (the field E is necessarily perpendicular to the symmetry axis), so this 1.187 factor is the standard relation $Q_{\perp}^{\rm pol} = (\frac{1}{3}Q_{\parallel} + \frac{2}{3}Q_{\perp}) \times f$ for the oblate $b/a = 1.4$ at this point in (λ, a) space.

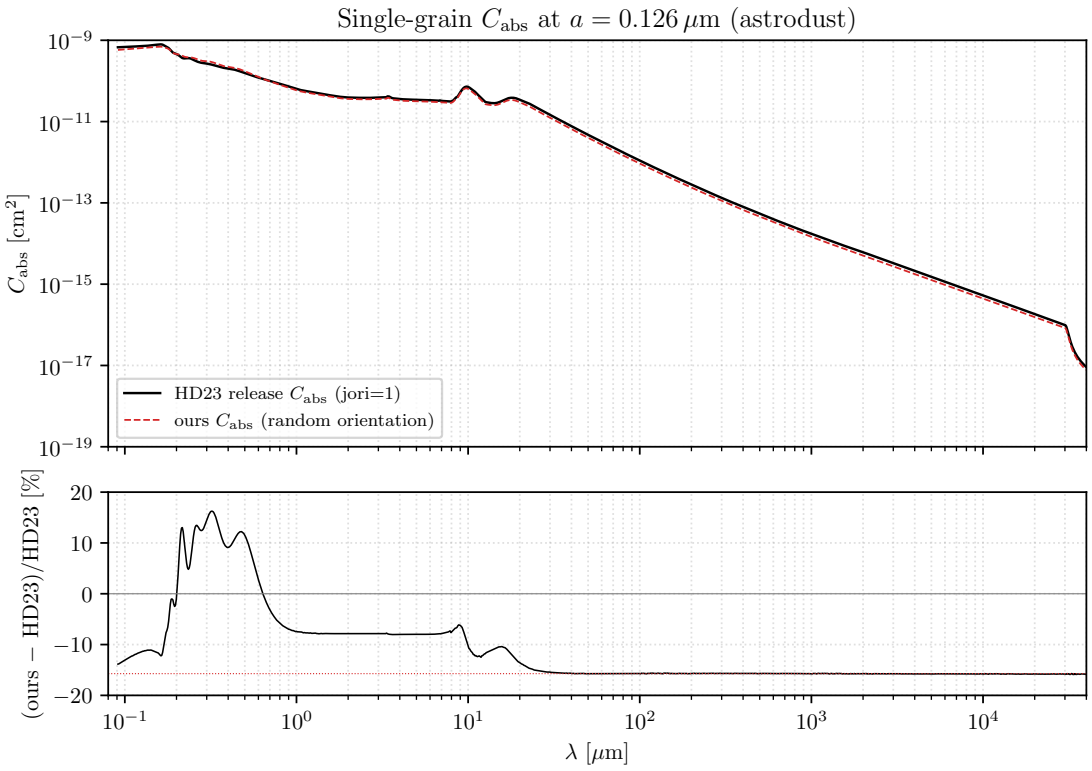


Figure 5: Single-grain absorption cross section at $a = 0.126$ µm. Top: HD23 release $C_{\rm abs}$ from `jori=1` (k‖a, black solid) and our $C_{\rm abs}$ from the random-orientation Q table that `sed_init` loads (red dashed); both tables share the DH21 wavelength grid, so no interpolation is involved. Bottom: relative residual (ours − HD23)/HD23 in percent. The flat −15.7% plateau across 30 µm–30 000 µm is precisely the ratio $\left(\frac{1}{3}Q_{\parallel} + \frac{2}{3}Q_{\perp}\right)/Q_{\perp}$ for an oblate $b/a = 1.4$ spheroid in the dipole limit at this grain.

Scaling test. If HD23 used the `jori=1` block of the release file directly (i.e. as if it were the random-orientation Cabs, instead of constructing the $\frac{1}{3}+\frac{2}{3}$ combination), the naive prediction for the aggregate SED would be $1.187\times$ ours, the reciprocal of the plateau above. Scaling our `S1_C1` SED by this factor and re-computing the relative residual against HD23 `astrodust_irem.dat` at $\log U = 0.20$ gives (remeasured on 2026-10-31 from `sed_astrodust_morig_S1.dat`, i.e. with the literal Mathis 1983 field, 10^{-13} dilution at 4000 K and a 2.9 K CMB, the configuration in which the test was first made):

band (μm)	unscaled med rel	scaled $\times 1.187$ med rel	notes
5–30	−0.8%	+17.7%	over after scaling
30–300 (FIR peak)	**−14.2%**	**+1.8%**	matched to 2%
300–3000 (sub-mm)	−5.5%	+12.2%	over after scaling
3000–30000	−3.8%	+14.1%	over after scaling

The single factor of 1.187, the ratio $\left(\frac{1}{3}Q_{\parallel}+\frac{2}{3}Q_{\perp}\right)/Q_{\perp}$ for an oblate $b/a = 1.4$ spheroid in the dipole limit at this grain, collapses the FIR-peak residual from −14.2% to +1.8%, but overshoots in adjacent bands. The 300–3000 µm overshoot tells us this single factor is *not* a global Cabs miscalibration – the rest of the deficit distribution at sub-mm and the few-percent slope across other bands must come from elsewhere (at the time attributed to size-distribution kernel weighting or to HD23-internal recipe details that cannot be reproduced from the public release files; the HD23 field of §7.1.2 brings the 300–3000 µm median to −0.42% without either).

Direct test of the hypothesis (avenue 3). We packaged the HD23 release `jori=1` block into a drop-in substitute Q-table and fed it through the full `sed_solve` calculation with a dedicated driver. Result for Stage S1_C1, measured with the literal Mathis 1983 field and before the enthalpy ghost-bin correction. The driver and the substitute table are not in this tree, so the test could not be remeasured on 2026-10-30 or 2026-10-31 and its numbers are stale: the "baseline" column is the state of that run, not today's production, and its "naive $\times 1.186$" column is that run's scaling test (with the factor as it was then rounded), not the remeasured one above.

band (μm)	baseline (then)	jori=1 substitute	naive $\times 1.186$	notes
5–30	+1.2%	+9.5%	+17.2%	
30–300 (FIR peak)	**−15.1%**	**−11.3%**	**−0.2%**	jori=1 substitute does *not* close the gap
300–3000 (sub-mm)	−6.1%	+7.1%	+11.9%	sign flips
3000–30000	−4.3%	+10.6%	+13.9%	sign flips

The substitution through the full calculation recovered only ~ 4 percentage points of the FIR-peak deficit (from −15.1% to −11.3%), nowhere near the −0.2% that the naive scaling predicted. At the same time the sub-mm and mm bands flipped sign and became large over-predictions. The reason the naive scaling fails: substituting a uniformly larger sub-mm $C_{\rm abs}$ changes $T_{\rm eq}$ self-consistently. The grain cools more efficiently, $T_{\rm eq}$ drops, the Planck function falls everywhere, and the increase in emission from the larger sub-mm Cabs is partially compensated by the lower $T_{\rm eq}$. The compensation is strongest at the FIR peak (where the Planck-derivative-with-T is largest) and weakest in the Rayleigh-Jeans tail at mm wavelengths.

Conclusion. Hypothesis (a) (HD23 used the `jori=1` slice directly when constructing `astrodust_irem.dat`) is *rejected*.

The 0.843 Cabs ratio plateau against HD23's `jori=1` table is the geometric relation between that orientation and the unpolarized average, but it does not map linearly to the SED residual once $T_{\rm eq}$ self-consistency is taken into account. The FIR-peak deficit is therefore *not* explained by an orientation-averaging convention difference. The candidate sources left, in order of investigation ease, were: (1) Mathis ISRF normalization differences between our `radfield.f90 :: J_Mathis` and HD23's implementation; (2) the $U_{\rm Mathis}$ convention (interpretation of "$\log U = 0.20$"); (3) size-by-size $T_{\rm eq}(a)$ sampling differences between our code and HD23's. Candidate (1) turned out to be the answer (§7.1.2): HD23's field uses a 3000 K dilution factor of 7×10^{-13} where `J_Mathis` keeps 4×10^{-13}, and the matched-field comparison there shows that this difference in field shape accounts for the main FIR difference.

Induced-emission factor: *gross* vs HD23 *net* emission convention. Draine's stochastic-heating code (`vsg_td_emission_v7.f`, which produces Draine's `iout.*.dbdis.gz` single-grain SED files) multiplies every emission term – both the large-grain equilibrium branch (line 1521) and the stochastic intra- and inter-bin branches (lines 1479, 1491, 1501) – by an explicit $\left(1+J_\lambda/B_{\rm env}\right)$ factor, where $B_{\rm env}(\lambda) \equiv 2hc^2/\lambda^5$ is the Planck envelope without the $1/(\mathrm{e}^{hc/\lambda kT}-1)$ Bose factor. This is the stimulated-emission term of the Einstein-coefficient algebra. Our `sed_solve` and `sed_solve_pah` accumulate single-grain emission as $C_{\rm abs}\,B_\lambda(T)$ without this factor, so this is a clean, isolable difference between the two codes.

To test whether the factor explains any part of the residual, the module `sed_astrodust_mod` carries the switch `use_induced_emission` (default `.false.`) and the routine `apply_induced_factor(J_lam, Jout)`, which applies $J_{\rm out}(\lambda) \mapsto J_{\rm out}(\lambda)\left(1+J_\lambda\lambda^5/(2hc^2)\right)$ once at the end of each solve (a wavelength-dependent factor with no T or grain dependence). The driver accepts `./calc_sed.x astrodust induced`, which sets the switch and adds the token `induced` to the output file names, so the baseline run is preserved. Running both back-to-back gives the ratio (induced/baseline) in Table 3.

Numerically, taken on the model's own wavelength grid so that no interpolation enters, the measured boost reproduces the analytic factor $1+J_\lambda/B_{\rm env}$ to a median of 9.4×10^{-10} and a worst case of 6.4×10^{-9} in relative terms across the 404 wavelengths where $J_\lambda/B_{\rm env}$ exceeds 10^{-6} — the round-off of the nine-significant-digit output files rather than anything in the calculation. The 30–300 μm FIR-peak band, where the deficit lives, is left unchanged to $< 3 \times 10^{-8}$. The factor therefore cannot be the source of the FIR-peak deficit, on quantitative grounds.

More interestingly, applying the factor and re-comparing against HD23's release `astrodust_irem.dat` at $\log U = 0.20$ shows that the millimetre-band agreement is *degraded* rather than improved (Table 4, S1_C1; the other two stages agree with it to within 1.3 percentage points). That comparison is made in the literal-Mathis configuration of the next paragraph (2.9 K CMB, 10^{-13} dilution at 4000 K), so its baseline column is the original-Mathis column of Table 5; it was measured on 2026-10-31 from `./calc_sed.x astrodust mathis_orig` and `./calc_sed.x astrodust mathis_orig induced`. With the default field the picture is the same: the mm-band median moves from –2.09% to +129.82%.

The $> 3000\,\mu$m swing from –3.8% to +137% leaves no ambiguity: HD23's `astrodust_irem.dat` does not carry the $\left(1+J_\lambda/B_{\rm env}\right)$ factor. The physical reason this is the right convention for a redistributable SED is direct. Write $C_{\rm abs}$ for the absorption cross section net of stimulated emission (the one the dielectric function gives) and $n_T = 1/(\mathrm{e}^{hc/\lambda kT}-1)$ for the photon occupation number at the grain temperature, so that $B_\lambda(T) = n_T B_{\rm env}$. The true absorption cross section, not corrected for stimulated emission, is $C_{\rm abs}(1+n_T)$, and spontaneous plus stimulated emission is $C_{\rm abs}B_\lambda(T)\left(1+J_\lambda/B_{\rm env}\right)$.

λ	analytic prediction	measured	dominant source
100 µm (FIR peak)	$+2.1\times10^{-9}$%	$<10^{-6}$%	diluted-BB RJ tail; CMB on Wien
300 µm	$+2.3\times10^{-6}$%	$+2.3\times10^{-6}$%	CMB still on Wien
1 mm	+0.5119%	+0.5119%	CMB transition
3 mm	+20.78%	+20.78%	CMB approaching RJ
5 mm	+53.34%	+53.34%	CMB RJ

Table 3: Effect of the $\left(1+J_\lambda/B_{\rm env}\right)$ induced-emission factor on the Stage S1_C1 SED. The factor depends only on λ (not on T or grain size), so all three astrodust stages (S1_C1, S1_C2, S2) and PAH show the same ratio: over the whole wavelength grid they agree with one another to 1.1×10^{-8} in relative terms, the round-off of the nine-significant-digit output files, even though the S1_C2 SED itself differs from S1_C1 by orders of magnitude. The growth at $\lambda\gtrsim1$ mm is driven by the CMB component of J_λ via $J_{\rm CMB}/B_{\rm env}(\lambda)=1/(e^{hc/\lambda kT_{\rm CMB}}-1)$ with $T_{\rm CMB}=2.725$ K, which supplies ≥99.7% of the factor beyond 300 µm; at 100 µm the CMB term is 1.2×10^{-23} and the whole (negligible) factor comes from the Rayleigh-Jeans tails of the Mathis diluted blackbodies. $T_{\rm CMB}$ is written down in exactly one place, `radfield.f90 :: cmb_temperature()` — 2.725 K by default, reverting to Mathis (1983)'s 2.9 K when `use_mathis_corrected` is switched off. The numbers above were measured with the default; the literal Mathis field gives +0.705%, +23.66% and +58.92% at 1, 3 and 5 mm instead. 100 µm and 1 mm are grid points of the model; the other three rows are interpolated linearly in $\ln\lambda$ between adjacent grid points, which are 1.1–1.2% apart there, and the interpolation itself contributes at most 10^{-5} of the ratio $1+J_\lambda/B_{\rm env}$. Measured 2026-10-31.

Their difference,

$$C_{\rm abs}\,B_\lambda(T)\left(1+J_\lambda/B_{\rm env}\right)-C_{\rm abs}\,(1+n_T)J_\lambda=C_{\rm abs}\left[B_\lambda(T)-J_\lambda\right],$$

is the net exchange with the field. A transfer code that attenuates the field with $C_{\rm abs}$ must therefore emit $j_\lambda=C_{\rm abs}\,B_\lambda(T)$ — the Kirchhoff form our baseline computes — and the equilibrium condition is $\int C_{\rm abs}J_\lambda\,d\lambda=\int C_{\rm abs}B_\lambda(T)\,d\lambda$, as our solver uses. Draine's single-grain `iout` files are the gross emission, to be paired with the uncorrected absorption; HD23 publishes the net form. The switch and the `apply_induced_factor` routine remain in `sed/src/sed_astrodust.f90` for cross-checks but *stay off (`.false.`) for production*, as anything fed into a downstream radiative-transfer code that itself handles the absorption needs the net convention. The candidate list for the FIR-peak residual was therefore unchanged by this test.

Mathis 4000 K dilution correction. Draine's reference Mathis 1983 ISRF implementation carries a 2008.02.02 note recording a numerical correction to the 4000 K diluted- blackbody weight: `OMEGA3` was changed from 10^{-13} (the literal Mathis 1983 value our `radfield.f90 :: J_Mathis` inherited) to 1.65×10^{-13} (the value also tabulated in Draine's textbook, *Physics of the Interstellar and Intergalactic Medium*, 2011, Table 12.1). The CMB component was also updated from 2.9 K to the Mather-et-al. value 2.725 K. Both corrections were added behind a toggle (`use_mathis_corrected`, default `.true.` in `radfield.f90`) so the original-Mathis comparison can be reproduced as `./calc_sed.x astrodust mathis_orig`, which adds the token `morig` to the output file names.

A single-grain diagnostic integrates the heating rate $\int C_{\rm abs}(\lambda)J_\lambda\,d\lambda$ at a fiducial $a=0.1$ µm astrodust grain and solves

band (μm)	baseline med rel	induced med rel
30–300 (FIR peak)	−14.18%	−14.18%
300–3000 (sub-mm)	−5.48%	−4.97%
$>$ 3000 (mm)	−3.84%	+136.59%

Table 4: Median relative residual (ours − HD23)/HD23 with and without the induced-emission factor, Stage S1_C1, literal Mathis 1983 field. Measured 2026-10-31.

for $T_{\rm eq}$ by bisection. The corrected ISRF raises $T_{\rm eq}$ from 17.34 K to 17.56 K (+0.22 K, +1.28%); the resulting $B_\lambda(T_{\rm eq})$ at $\lambda = 130\,\mu$m rises by +8.4%, with the FIR-peak boost growing monotonically with grain size from +3% at 5 nm (PAH/small) up to +21% at 1 μm (cool large grain whose Planck peak sits closer to 130 μm). The CMB temperature change (2.9 K $\rightarrow$ 2.725 K) is negligible for $T_{\rm eq}$ because CMB heating is subdominant to stellar UV at all but the very largest grain sizes.

The full-calculation effect for the headline comparison (Stage S1_C1 vs HD23 `astrodust_irem.dat` at $\log U = 0.20$) is summarised in Table 5: the FIR-peak deficit moves from −14.18% to −7.37%, sub-mm from −5.48% to −3.00%, and the mm-band from −3.84% to −2.09%. The PAH 30–300 μm band, the most discrepant of the well-behaved bands at −8.17%, closes to −3.68%. Stages S1_C2 and S2 show the same pattern (FIR moves to −7.47% and −7.66% respectively).

band (μm)	original Mathis	production (corrected)	shift
30–300 (FIR peak)	−14.18%	−7.37%	+6.81 pp
300–3000 (sub-mm)	−5.48%	−3.00%	+2.48 pp
$>$ 3000 (mm)	−3.84%	−2.09%	+1.75 pp
PAH 30–300	−8.17%	−3.68%	+4.49 pp

Table 5: Median relative residual (ours − HD23)/HD23 for Stage S1_C1 with and without the corrected Mathis 4000 K dilution factor. The corrected variant is now the production default; the original-Mathis variant is reproducible via `./calc_sed.x astrodust mathis_orig`.

The two field settings in Table 5 were remeasured on 2026-10-31 from the branch SED files. Their different 4000 K dilution factors account for the listed change in the astrodust FIR residual. The remaining difference between the branch default and HD23 follows from the 3000 K dilution factor and the field's spectral shape, as quantified in Sect. 7.1.2.

Both columns of Table 5 moved from the values this report carried before 2026-10-30 (−15.83% and −8.82% in the FIR row, −4.40% and +0.02% in the PAH row); the 2026-10-31 remeasurement gives the same values as that of 2026-10-30. Two corrections made before 2026-10-30 account for the shift, and they act on different rows. The enthalpy ghost-bin correction moved the astrodust rows: the S1 and S2 SEDs were bit-identical across the second correction, so everything the astrodust rows moved by came from the enthalpy side. The PAH row carries both: correcting the DL07 eq. 2 cation near-IR continuum, whose Gaussian factor is $\exp[-(0.1x)^2]$ and not $\exp(-0.1x^2)$, raised the PAH SED by 0.99% bolometrically and lifted the PAH 30–300 μm median by +0.76 pp on its own (from −4.43% to −3.67%), the remainder of that row's shift coming from the enthalpy correction. These attributions were measured on 2026-10-30 against code

without the corresponding correction, which this tree no longer contains, and were not remeasured.
The 4000 K correction closes roughly half of the original FIR deficit, from −14.18% to −7.37%. The remaining 7.4% was consistent with either a different normalization or shape of the Mathis field in HD23 (candidates (1) and (2) of the avenue-3 conclusion above) or a size-by-size $T_{\rm eq}(a)$ sampling difference between our code and HD23's (candidate (3)). The field proved to be the answer, in the form of a different 3000 K dilution factor in HD23's field (§7.1.2).

What is in our hands and what is not. Our implementation – T-matrix Q, small-x Rayleigh spheroid limit, random-orientation average $\left(\frac{1}{3}+\frac{2}{3}\right)$, `calc_Teq`, κ_B build, size sum – is internally self-consistent and reproduces HD23 `extinction.dat` to better than 1% on aggregate τ/N_H. The consistency of our cross sections, equilibrium temperatures, and aggregate τ gives us confidence that the model is correct in everything we control. The remaining FIR-peak residual in $\lambda I_\lambda/N_H$ vs. HD23 `astrodust_irem.dat` was subsequently traced to the 3000 K dilution factor of our default field (§7.1.2); the full chain of diagnostics is given below; the single-grain figure is `figs/cabs_per_grain_vs_hd23.pdf`.

7.1.2 Resolution: the 3000 K dilution factor of the heating field

The residual was resolved by an energy-budget audit against the release's own optics. The release provides its own component-by-component $C_{\rm abs} = C_{\rm ext} - C_{\rm sca}$ (`extinction.dat`, `scattering.dat`), against which our optics were validated to $\lesssim 1\%$ (extinction and albedo) and to ~ 4 digits grain by grain. Integrating that *released* $C_{\rm abs}$ against a heating field $\times 10^{0.2}$ — extrapolating across the 0.0912 μm to 0.1 μm extinction-grid gap with the local power law, which adds 3.8–4.2% (astrodust) and 6.7–6.8% (PAH) to the absorbed power — gives the absorbed power, which energy conservation requires to equal the emitted power $4\pi \int \lambda I_\lambda \, d\ln\lambda$ of the `irem` tables. Table 6 gives the budget for two fields that differ only in the 3000 K dilution factor: the default `mathis` field, which keeps Draine's 2008 3000 K dilution ($W = 4\times10^{-13}$, the Mathis 1983 value), and the HD23 field.

field	component	$P_{\rm abs}/N_H$ (erg s^{-1})	$P_{\rm em}^{\rm HD23}/N_H$ (erg s^{-1})	$P_{\rm em}^{\rm HD23}/P_{\rm abs}$	ours$/P_{\rm abs}$	ours/HD23
`mathis`	astrodust	2.724×10^{-24}	2.977×10^{-24}	1.093	0.999	0.914
	PAH	2.002×10^{-24}	2.067×10^{-24}	1.033	0.997	0.965
`hd23`	astrodust	2.978×10^{-24}	2.977×10^{-24}	**1.000**	0.999	0.999
	PAH	2.037×10^{-24}	2.067×10^{-24}	**1.015**	0.997	0.982

Table 6: Energy budget per H at $\log U = 0.20$ for the default `mathis` field ($W_{3000} = 4\times10^{-13}$) and the piecewise HD23 field ($W_{3000} = 7\times10^{-13}$). $P_{\rm abs}$ integrates the released $C_{\rm abs}$; $P_{\rm em}^{\rm HD23}$ integrates the released `irem` column; ours integrates `sed_astrodust_S1.dat`, `sed_astrodust_PAH.dat` (`mathis`) and `sed_astrodust_hd23_{S1,PAH}.dat` (`hd23`) on the HD23 wavelength grid. The branch driver uses $U = 1.585$; the release column is at $U = 10^{0.2}$. The calculated λI_λ is interpolated linearly in $\ln\lambda$. The CMB term, 2×10^{-5} of $P_{\rm abs}$, is omitted from the stellar-field absorption estimate. Remeasured 2026-10-31.

With the default field the release appears to emit 9.3% (astrodust) and 3.3% (PAH) more than it absorbs. The audit of 2026-06-12 read this as an inconsistency of the release and argued that a heating-field offset would shift both components

equally. That argument holds only for a uniform rescaling of U. A change in the *spectral shape* of the field changes the absorbed power of each component according to its own $C_{\rm abs}(\lambda)$, and the 3000 K blackbody puts its power in the optical and near-IR, where astrodust absorbs a much larger share of its total than the UV-dominated PAHs. The shape differs: HD23 eq. (9) adopts the Draine (2011) dilution factors $W = 7 \times 10^{-13}$ (3000 K), 1.65×10^{-13} (4000 K) and 10^{-14} (7500 K), while our default `radfield.f90 :: J_Mathis` follows Draine's 2008 `mmp_radiation_field.f`, which corrected the 4000 K factor (§7.1.1) but kept 4×10^{-13} at 3000 K. Changing that one factor raises the absorbed power by 9.3% (astrodust) and 1.8% (PAH), and the budget closes: emitted over absorbed power is 1.000 (astrodust) and 1.015 (PAH). The remaining 1.5% for PAH is smaller than the 6.7% that the 0.0912 μm to 0.1 μm extrapolation adds to the PAH absorbed power, so this audit does not resolve it.

HD23 eq. (9) is printed as a sum of the ultraviolet power laws and the three blackbodies. The Mathis 1983 field and Draine's implementation are piecewise: the ultraviolet power laws are fits to the observed field over 0.0912 μm to 0.246 μm and already contain the stellar Wien tails, and the blackbodies take over above 0.246 μm. The literal sum, with the blackbodies also added below 0.246 μm, gives emitted over absorbed power 0.974 (astrodust) and 0.917 (PAH), overshooting in the other direction, so `J_HD23` keeps the piecewise form.

band (μm)	S1	S2	PAH	S1, `mathis`
1–5	+71.24% (1.387)	+164.35% (1.879)	+14.75% (1.022)	+71.15% (1.386)
5–30	+3.70% (1.043)	+10.97% (1.156)	+4.34% (0.969)	+1.53% (1.013)
30–300	–0.13% (0.999)	–0.37% (0.999)	+0.32% (1.001)	–7.37% (0.910)
300–3000	–0.42% (0.994)	–1.06% (0.986)	+3.57% (1.021)	–3.00% (0.954)
FIR peak	–0.69%	–1.76%	–1.95%	–9.80%

Table 7: Median relative residual (ours − HD23)/HD23 and, in parentheses, the band power ratio ours/HD23 at $\log U = 0.20$, with the HD23 field (`./calc_sed.x astrodust hd23`; gfortran, NT = 200, heuristic solver) for the astrodust stages S1 and S2 and the PAH component; the last column repeats S1 with the default `mathis` field. Ours uses linear interpolation of λI_λ in $\ln\lambda$ onto the released grid, with the sampled band edges included. The branch driver uses $U = 1.585$, while the release column is at $U = 10^{0.2}$. The FIR peak is 133.4 μm for astrodust and 58.0 μm for PAH. Remeasured 2026-10-31.

With the HD23 field our SED reproduces the release in the FIR and sub-mm (Table 7): the astrodust S1 median is –0.13% over 30 μm to 300 μm and –0.69% at the 133 μm peak, and the residual stays within 1% from 80 μm to 300 μm; the excess shortward of 60 μm (up to +27% at 35 μm) is the stochastically heated shoulder discussed above. The emergent astrodust SED integrates to 0.999× HD23 over 0.1 μm to 3000 μm. The matched-field comparison shows that the field's spectral shape accounts for the main FIR difference. This confirms the field candidate (candidate (1) of the avenue-3 conclusion in §7.1.1) in the form of the field shape, and retires the size-by-size $T_{\rm eq}$ sampling candidate (3). The 1 μm to 5 μm astrodust excess (+71% median for S1, 1.39× in band power) is the same with either field and is not explained here; it lies in the hot-grain tail and carries 0.04% of the total astrodust power.

7.2 PAH emission

HD23 §3.2 follows Draine et al. (2021), whose PAH cross sections are themselves based on DL07; the DL07 cross sections are already implemented in SED_v5 as `QPAH_DL07(charge, R, lambda, Qabs)` and were forked into `sed/src/qpah.f90`. As noted in HD23 §3.2, PAH grains are treated as spherical with $\rho_{\rm PAH} = 2.0$ g/cm^3, depend only on grain size and charge (no orientation), and *do not contribute to polarized extinction or emission*; this matches our implementation (no T-matrix on the PAH track).

The HD23 PAH cross section per grain is built in three steps:

Step 1: charge mixing. The DL07 cross sections are available for neutral and ionized PAHs; HD23 eq. 15 mixes them with an ionization fraction $\xi_{\rm ion}(a)$ (HD23 eq. 19):

$$C^{\rm Z}(\lambda,a) = [1-\xi_{\rm ion}(a)]\,C^{\rm neu}(\lambda,a) + \xi_{\rm ion}(a)\,C^{\rm ion}(\lambda,a), \tag{10}$$

with $\xi_{\rm ion}(a) = 1 - 1/(1 + a/10\,\text{Å})$ from `size_distribution.dat` column 4.

Step 2: graphite continuum blending (HD23 eq. 15-16). DL07's PAH cross section drops to essentially zero at long wavelengths; HD23 retains DL07's recipe of blending into a random-orientation graphite Mie cross section $C^{\rm gra}(\lambda,a)$ for $a > 50$ Å, so that PAH grains acquire the same sub-mm continuum opacity that DL07 reports for "turbostratic" graphite. The blended cross section is

$$C^{\rm PAH,Z}(\lambda,a) = [1-\xi_{\rm gra}(a)]\,C^{\rm Z}(\lambda,a) + \xi_{\rm gra}(a)\,C^{\rm gra}(\lambda,a), \tag{11}$$

with

$$\xi_{\rm gra}(a) = \begin{cases} 0.01 & a \le 50\,\text{Å} \\ 0.01 + 0.99\left[1-(50\,\text{Å}/a)^3\right] & a > 50\,\text{Å}. \end{cases} \tag{12}$$

$C^{\rm gra}(\lambda,a)$ is the random-orientation Mie cross section for graphite spheres, $C^{\rm gra} = \frac{1}{3}C^{\rm gra}_{\parallel} + \frac{2}{3}C^{\rm gra}_{\perp}$, with the two components computed from the Draine 2003 graphite dielectric tables (`index_CpaD03`, `index_CpeD03`) using a size-dependent free-electron contribution (`free_diel`, Draine 2003 form for $T_{\rm gr} = 20$ K) and a standard Bohren–Huffman Mie solver. Implementation: new module `sed/src/q_graphite.f90` loads the two graphite tables lazily, caches the (n,k) tables per grain radius, and calls `sed/src/mie.f90` (a module wrapper of the BHMIE code already in `tmatrix/test/`) twice per (a,λ). Below the DL07 PAH cutoff at $\lambda = 1/17.25\,\mu$m the blend reduces to pure graphite (HD23 / SED_v5 convention), since $C^{\rm Z}$ is not defined there.

Step 3: size and enthalpy. The PAH size distribution is column 3 of `size_distribution.dat` (the second of HD23 eq. 17's two lognormals, peaks at ~ 4 and 30 Å). PAH enthalpy uses the DL01 `'Car0'` branch with the H/C-mode correction for $N_{\rm atom} < 10^4$, as recommended by HD23 since PAH is carbonaceous.

A new function `sed_solve_pah(J_lam, lamI_lam_out)` solves the single-cell PAH SED with the same dynamic-T solver. The combined total (Ad + PAH) is compared against HD23 `model_irem.dat` in Figure 6.

The PAH peak is at 7.67 μm, compared with 7.69 μm in the HD23 release. The current PAH band-integrated ratios for all three graphite choices are given in Sect. 7.3; the branch default is the D16 sphere.

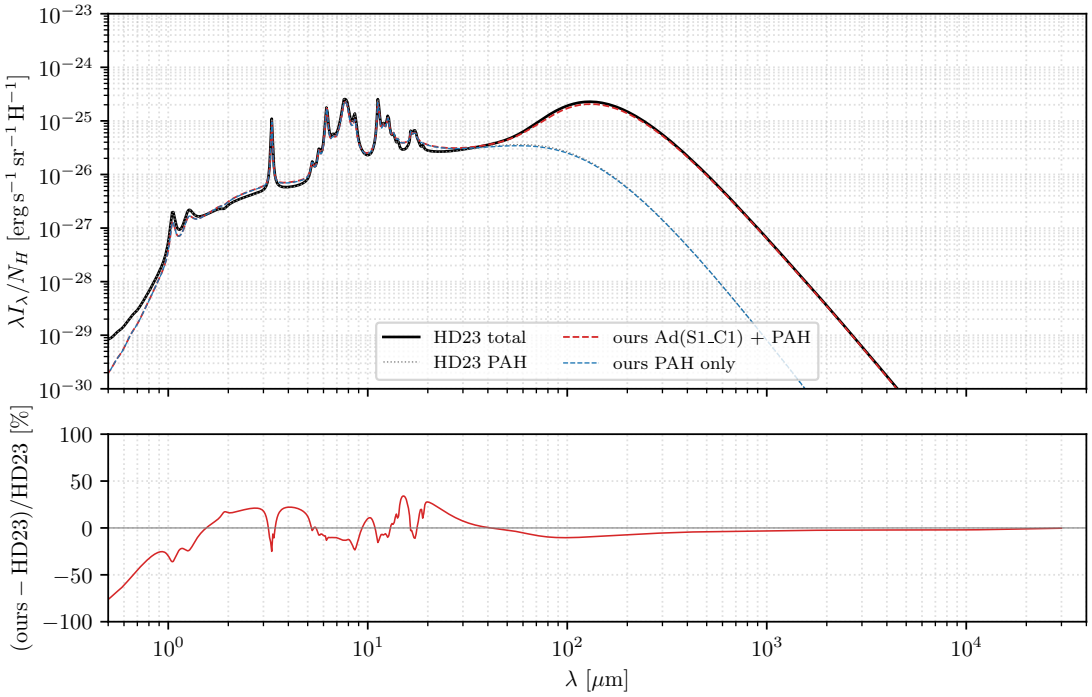


Figure 6: Total dust thermal emission (Ad + PAH) with the corrected MMP83 field at $U = 1.585$, compared with the HD23 released spectra at $U = 10^{0.2}$. Top: HD23 `model_irem.dat` (black solid), HD23 PAH-only (`PAH_irem.dat`, gray dotted), our S1_C1 + PAH (red dashed), and our PAH-only (blue dashed). Bottom: relative residual on the total. The PAH 7.7 µm feature is in the right place and the right order of magnitude.

First closure of the graphite continuum (historical). When the graphite continuum of Step 2 was first added, with the D03 sphere, the PAH band ratios against HD23 `PAH_irem.dat` at $\log U = 0.20$ (band-integrated $\int(\lambda I_\lambda)/\lambda d\lambda$) were those of the "then" column below. They predate the adoption of D16 graphite, the enthalpy ghost-bin correction and the correction of the DL07 eq. 2 cation near-IR continuum, were recorded without a date, and cannot be remeasured from this tree; they are kept as the record of that state. The "D03 today" column is the same variant measured on 2026-10-31 (§7.3).

band (μm)	then	D03 today	notes
optical ($\lambda < 1\,\mu$m)	0.90	0.55	absolute $\sim 10^{-28}$, negligible vs total
NIR (1–5 μm)	0.61	0.76	PAH-feature region; see Sect. 7.4
PAH features (5–15 μm)	1.04	0.92	
mid-IR (15–60 μm)	1.19	1.18	
FIR (60–300 μm)	0.76	0.75	
sub-mm (300–3000 μm)	0.61	0.61	cf. ~ 0.10 without eq. 11 (then)
TOTAL (0.1–3000 μm)	1.01	0.96	

Adding the graphite continuum then closed approximately 75% of the sub-mm gap, from a $\sim 90\%$ under-prediction to $\sim 40\%$. The 40% that remained had the same sign and roughly the same level as the astrodust FIR deficit of that time,

which suggested a common origin unrelated to PAH cross-section physics (Sect. 7.1.1); §7.3 closes it instead, with the D16 graphite of HD23's own recipe.

7.3 ξ-blend graphite update: D16 turbostratic

Step 2 above was first built with Draine 2003 (D03) basal-plane and c-axis graphite dielectric functions plus our own BHMIE on a sphere to synthesise $C^{\rm gra}(\lambda, a)$. HD23 themselves do *not* use D03 graphite; they state (HD23 §3.2) that they adopt the Draine 2016 (D16) turbostratic graphite cross sections in the Maxwell–Garnett effective-medium approximation. The branch default therefore uses the D16 cross sections, which brings the model in line with the HD23 recipe and removes the largest remaining systematic in the PAH SED. The D03 sphere and the D16 spheroid remain explicit comparison choices. The D16 sphere gives the closest PAH submillimeter continuum to the HD23 release.

Graphite tables and choices. The module `sed/src/q_graphite_d16_sphere.f90` reads Draine's precomputed D16 sphere Q-table (`q_D16graphite.dat`, the decompressed `callqcomp_D16MGemt.gz` from the Princeton D16graphite web page; 41 radii $\times$ 3501 wavelengths). A bilinear log-log interpolator returns $Q^{\rm gra}_{\rm abs}(a, \lambda)$ without re-running Mie. Which graphite the ξ blend takes is the three-valued module setting `qpah_graphite_source` in `sed/src/qpah.f90`: `'d16_sphere'` (production default), `'d03_sphere'` (the original D03 Mie module `sed/src/q_graphite.f90`, which the `dl07` model of `calc_sed.x` and `calc_kext.x` selects) and `'d16_spheroid'`. An unrecognised value stops the run instead of falling back silently. Everything else (the $\xi_{\rm gra}(a)$ blend, the cation/neutral mixing, the sub-mm graphite cutoff) is unchanged, and above the DL07 PAH cutoff all three settings still take graphite from the D03 dielectric functions, because there the blend needs radii smaller than the D16 tables reach. All three are named on the command line, as `./calc_sed.x astrodust gra_d03_sphere`, `gra_d16_sphere` and `gra_d16_spheroid` — and by the same words on `./calc_sed.x dl07`, whose carbonaceous population takes the same blend — tagging their output files `d03gra_`, `d16gra_` and `sphdgra_`. All of them touch the PAH population only, so the astrodust S1 and S2 files they write are byte-identical to the production ones.

Sphere vs. b/a = 1.4 oblate spheroid. Draine also publishes a $b/a = 1.4$ oblate-spheroid version of the same turbostratic D16 graphite (`qlib_gra_D16MGemt_1.400.gz`; 169×1009 from his own T-matrix runs). A second module `sed/src/q_graphite_d16.f90` loads that table. Being a spheroid table it is orientation-resolved, and on the same convention as the HD23 astrodust table its `jori=1` block is $k \parallel a$, *not* the random-orientation average. An unaligned PAH population needs $\frac{1}{3}Q^{\rm jori=2} + \frac{2}{3}Q^{\rm jori=3}$, which the module now forms on load. Taking `jori=1` for it instead – which this module did until the present revision – overstates $Q^{\rm gra}_{\rm abs}$ by 22% in the Rayleigh limit, since `jori=1` carries no $E \parallel a$ component at all, and that alone accounted for the whole of the spheroid variant's apparent sub-mm excess: the 300–3000 μm PAH-only band ratio is 1.21 read as `jori=1` and 1.05 with the random-orientation average. Comparing the three ξ-blend graphite choices (Figure 7) still leaves the spherical D16 variant closest to HD23 `PAH_irem.dat` in the sub-mm (300–3000 μm PAH-only band ratio 0.99 against 1.05). That is consistent with HD23 §3.2, which inherits DL07's spherical-PAH framework even though the astrodust population itself is treated as $b/a = 1.4$ oblate. The D16 sphere is therefore the appropriate choice for the ξ-blend and is the production setting.

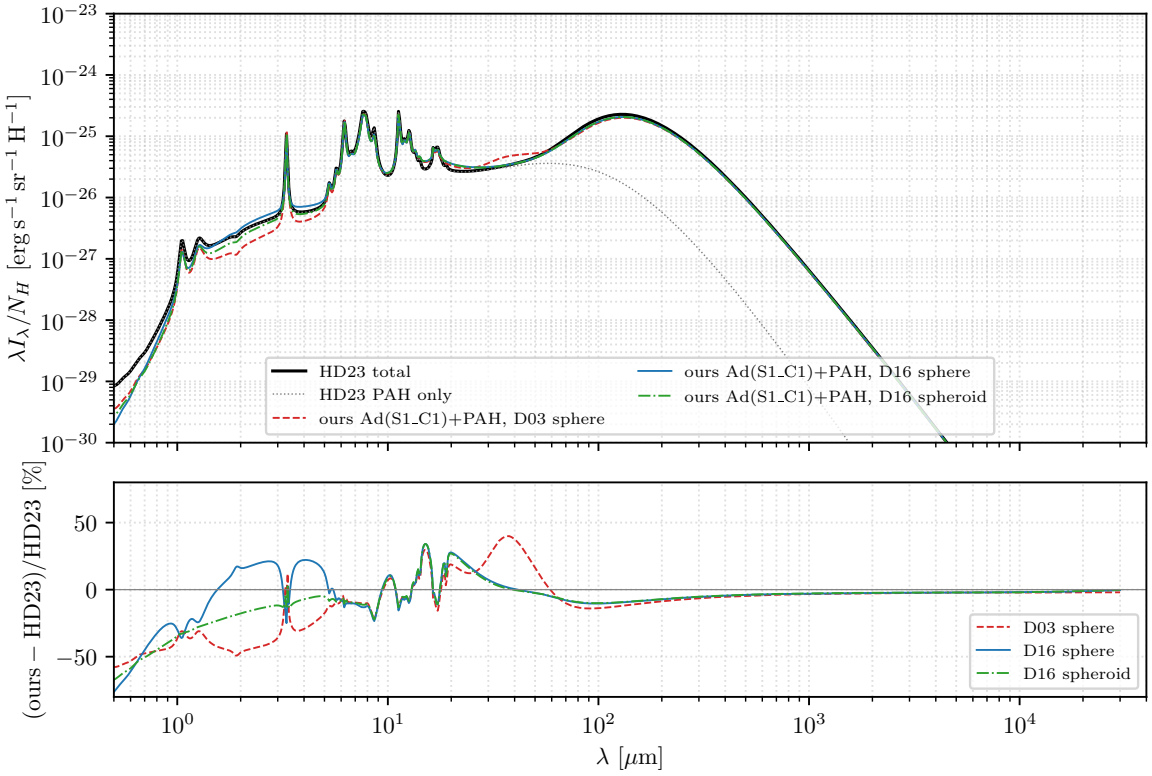


Figure 7: Three graphite choices for the ξ-blend continuum, all calculated with the corrected MMP83 field at $U = 1.585$. The HD23 released spectra at $U = 10^{0.2}$ are shown for comparison. Top: HD23 total (black solid) and HD23 PAH-only (gray dotted) on the release wavelength grid, plus our `S1_C1 + PAH` totals with three graphite-source choices for the ξ-blend (D03 sphere red dashed, D16 sphere blue solid, D16 spheroid green dash-dot). Bottom: relative residual of each total against HD23 `model_irem.dat`. What separates the three is the PAH sub-mm: the 300–3000 μm PAH-only band ratio is 0.61 for the D03 sphere, 0.99 for the D16 sphere and 1.05 for the D16 spheroid. On the total the choice matters much less, because the astrodust population carries that band — the residual at the 130 µm FIR peak moves only from −12.0% (D03 sphere) to −9.5% (D16 sphere) and −9.3% (D16 spheroid), and the +34% spike near 15 µm is in all three.

Band-integrated ratios. The table compares the production D16 sphere with the D03 sphere and D16 spheroid options. "ours / HD23" is the ratio of band-integrated $\int(\lambda I_\lambda)/\lambda d\lambda$ on the wavelength grid of the release file (`PAH_irem.dat`, $\log U = 0.20$):

band (μm)	D03 sphere	D16 sphere	D16 spheroid	notes
optical ($\lambda < 1\,\mu$m)	0.55	0.67	0.57	absolute $\sim 10^{-28}$, negligible vs total
NIR (1–5 μm)	0.76	1.02	0.88	PAH-feature region
PAH features (5–15 μm)	0.92	0.91	0.92	dominated by DL07 Drude profiles
mid-IR (15–60 μm)	1.18	1.06	1.05	
FIR (60–300 μm)	0.75	0.95	0.98	D03 is 25% below HD23
sub-mm (300–3000 μm)	0.61	**0.99**	1.05	D16 sphere essentially zero residual
TOTAL (0.1–3000 μm)	0.96	0.97	0.97	global agreement preserved

These ratios were remeasured with `docs/make_figs.py` on 2026-10-31. The branch default is the D16 sphere. Its PAH-only 300–3000 μm ratio is 0.9883; the D03 sphere gives 0.6080 and the D16 spheroid gives 1.0506. In the D16 sphere calculation, the 5–15 µm band integrates 9.2% below HD23 and the 1–5 µm band is 2.1% above it. The residual feature structure is examined in Sect. 7.4.

All three columns are measurements of the present code, taken in the same run of `docs/make_figs.py` that draws Figure 7: the production run supplies the D16-sphere column, and `./calc_sed.x astrodust gra_d03_sphere` and `./calc_sed.x astrodust gra_d16_spheroid` the other two. They may therefore be read digit for digit against each other. The D03-sphere column differs from the historical table in §7.2, which predates the enthalpy ghost-bin correction and the correction of the DL07 eq. 2 cation near-IR continuum to $\exp[-(0.1x)^2]$. The 40% sub-mm under-prediction of the D03 sphere is closed by the D16 swap (0.608 to 0.988). What is left is in the PAH features themselves — the 5–15 µm band integrates 8–9% low for all three graphite choices — a cross-section issue independent of the graphite source.

7.4 Residual PAH-feature strength mismatch

After the D16 sphere swap, the remaining $\lambda < 5$ µm PAH residual was first decomposed into (i) a smooth NIR continuum, which then carried a −15% NIR-band deficit (a value of that earlier code state, not remeasured), and (ii) a feature-by-feature strength mismatch independent of the graphite source. Part (i) no longer holds: with the D16 sphere, the PAH 1–5 µm band integrates to 1.02× HD23. The feature-by-feature differences remain and are independent of the graphite choice. The numbers below compare the current branch default with HD23 `PAH_irem.dat` at $\log U = 0.20$, as measured by `docs/make_figs.py`. A feature's band is $\lambda_j(1 \pm \gamma_j)$ — one Drude FWHM to either side of the DL07 Table 1 mode centre, since a profile written in $(\lambda/\lambda_j - \lambda_j/\lambda)$ reaches half maximum at $\pm\gamma_j$ and so has FWHM = $\gamma_j\lambda_j$ — united over the modes that make up a complex, and widened to the three nearest release grid points where the FWHM falls below the grid spacing (only the 3.3 µm stretch, FWHM 0.040 µm against a 0.042 µm spacing). The bands contain the underlying continuum, so these are band ratios and not continuum-subtracted feature strengths:

feature	band (μm)	ours / HD23	
3.3 µm aromatic C–H stretch	3.256–3.344	0.79	largest deficit
6.2 µm C–C stretch	6.033–6.407	0.89	
7.7 µm C–C complex	6.482–8.352	0.87	modes $j = 11$–13
8.6 µm C–H in-plane bend	8.274–8.946	0.79	
11.3 µm C–H out-of-plane	10.967–11.693	0.87	modes $j = 17, 18$
12.0–12.7 µm duo/trio	11.450–13.150	0.93	modes $j = 19$–21
14.19 µm quartet	13.835–14.545	1.14	the one band over-predicted
17 µm complex	15.932–18.156	0.98	modes $j = 26$–28

Every band except the 14.19 µm quartet comes out low, by 2–21%, and the deficit is deepest at the two narrowest C–H bands (3.3 and 8.6 µm, both −21%). The HD23 paper §3.2 states that its PAH cross sections are "identical to those of DL07 with the exceptions of the 14.19 µm feature and the 17 µm complex, which have both been increased in strength by 33%"; our residuals are not consistent with that single isolated change. Were those two bands the only difference, both would sit $\sim 25\%$ low here, whereas the 14.19 µm band is 14% *high* and the 17 µm complex agrees to 2%, while bands HD23 says it left alone are 10–20% low. A separate inspection of DL07 Table 1 itself (in the context of the SED_v5 `QPAH_DL07` fork from which `sed/src/qpah.f90` descends) found that the 1.260 µm ionized PAH strength is $\sigma_{\rm Ion}(4) = 0.078$ in the published table whereas a value of 7.8×10^3 (a factor of 10^5 larger) reproduces the publicly distributed DL07 SED; the SED_v5 fork uses the corrected value and this report inherits that correction. We document the feature-by-feature residuals here as a numerical record; a definitive resolution would require either an independent re-derivation of the DL07 Drude parameters or direct correspondence with the HD23 authors about which DL07 Table 1 entries went into their `PAH_irem.dat` construction.

Mode-by-mode Drude tuning to HD23 (record of a past fit). An empirical mode-by-mode rescaling of the DL07 Drude parameters was fitted in May 2026 (the namelist that records it is dated 2026-05-18 in the working tree that preceded this repository), against HD23 `PAH_irem.dat` at $\log U = 0.20$, and closed the feature-by-feature mismatch, as it was estimated then, to within $\pm 3\%$ at the dominant peaks (3.3, 6.2, 7.7-complex, 8.6, 11.3-complex, 12.0–12.7-complex, 17 complex). Mode-by-mode multiplicative factors $f_\sigma^{(j)}$ for `sigma_neu(j)` and `sigma_ion(j)` (and optionally $f_\gamma^{(j)}$ for the Drude width) live as module-level public arrays in `sed/src/qpah.f90` (switch `use_tuned_pah`, default `.false.`), and the fitted values are recorded in `sed/pah_tune.nml`. No driver in this tree reads that file and none takes a `pah_tuned` argument: switching the tuning on means setting `use_tuned_pah` and the two arrays from calling code, and there is no tuned production run. The convergence path was: (i) three iterations of single-mode adjustment using a feature-by-feature ratio diagnostic; (ii) a 2-mode least-squares fit on $j = 22$ and $j = 23$, which drove $f_\sigma^{(23)}$ to the lower clamp at 0.30 – diagnostic for graphite-continuum bleed into those weak features rather than a PAH-σ issue, so the iteration-(i) values were kept; (iii) a 130-point (σ_6, γ_6) brute-force grid search at the 3.3 µm peak to resolve the σ-vs-γ ambiguity inherent in that very narrow Drude profile ($\gamma_6 = 0.012$, FWHM $\approx 0.04\,\mu$m – comparable to HD23's grid spacing of $0.042\,\mu$m). The grid optimum was at $f_\sigma^{(6)} = 2.00$ with $f_\gamma^{(6)} = 1.00$ (i.e., HD23 boosted only the integrated strength of the 3.3 µm C–H stretch; the underlying Drude width matches DL07), with $\chi^2 = 5.0 \times 10^{-4}$ over the three nearest HD23 grid points.

Pure-γ-narrow ($f_\gamma = 0.45$) and equal-boost ($f_\sigma = f_\gamma = 1.45$) scenarios were 5–10$\times$ worse. Table 8 lists the final factors. The residual +5–15% over-prediction at 13–19 μm was attributed to the graphite continuum (it could not be closed by PAH σ adjustment alone, as the least-squares fit showed).

Two things about Table 8 need saying plainly. The factors were fitted against the code as it stood before the enthalpy ghost-bin correction and before the DL07 eq. 2 cation near-IR continuum was corrected to $\exp[-(0.1x)^2]$, so they are not a fit to the residuals tabulated above. And the estimator they were fitted to — the "optical-equivalent ratio" of each feature — was never recorded with a band definition or a continuum-subtraction rule, so the fit cannot be re-derived from what this tree carries. The table is kept as the record of that fit, not as a recommended setting.

j	λ_j (μm)	feature	$f_\sigma^{(j)}$
6	3.300	C–H stretch (3.3)	**2.00**
7	5.270	C–H/C–C combo	0.84
8	5.700	C–H/C–C combo	0.81
9	6.220	C–C stretch (6.2)	0.73
10	6.690	weak	0.86
11	7.417	7.7 complex (a)	0.79
12	7.598	7.7 complex (b)	0.79
13	7.850	7.7 complex (c)	0.76
14	8.330	weak	1.04
15	8.610	C–H in-plane bend (8.6)	1.02
17	11.23	C–H oop (solo narrow)	1.34
18	11.33	C–H oop (solo broad, 11.3)	1.33
19	11.99	duo	1.17
20	12.62	trio	1.24
21	12.69	trio	1.22
22	13.48	quartet	1.06
23	14.19	quartet (HD23 +33%)	1.03
25	16.45	C-C-C bend	1.15
26	17.04	C-C-C bend (HD23 +33%)	1.34
27	17.375	C-C-C bend	1.24
28	17.87	C-C-C bend	1.02
29	18.92	weak	0.87

Table 8: Mode-by-mode Drude σ multipliers of the May 2026 fit to HD23 `PAH_irem.dat` at $\log U = 0.20$, applied to DL07 Table 1 (all $f_\gamma^{(j)}$ left at 1.00). $f_\sigma^{(j)}$ is applied identically to `sigma_neu` and `sigma_ion`. Modes not listed are kept at $f = 1$. Recorded in `sed/pah_tune.nml`, which no driver reads; see the text for the code state the fit was made in.

The production calculation uses its baseline DL07-derived mode parameters and leaves `use_tuned_pah` disabled. The ratios in the table above therefore describe the actual branch default, without empirical feature adjustments.

8. Polarization

Per HD23 §3.2, the PAH track does *not* contribute to polarized extinction or emission (PAH grains are spherical in the model and not aligned). The polarization comparison below therefore uses the astrodust population only.

This section is a *verification of the T-matrix optics* against the HD23 release, carried out by a standalone utility. It is not a library capability: this tree is the scalar branch (v1.00), whose API (`dust_emission`, `dust_extinction`) is unpolarized throughout. The dichroic extinction matrix and the aligned-grain scattering matrix live in the separate polarized branch (v1.20).

For perfectly oblate spheroids with the short (symmetry) axis aligned perpendicular to a magnetic field in the plane of the sky (the "MPFA" geometry of HD23 §2.2), the maximum polarized extinction per H atom is

$$\left(\frac{p_\lambda}{N_H}\right)^{\max} = \sum_a \frac{dn_{\rm Ad}}{N_H}(a)\, C_{\rm pol}^{\rm ext}(\lambda, a)\, f_{\rm align}(a), \tag{13}$$

where

$$C_{\rm pol}^{\rm ext}(\lambda, a) = \frac{1}{2}\pi a^2 \left[Q_{\rm ext}^{E\perp a}(\lambda, a) - Q_{\rm ext}^{E\parallel a}(\lambda, a)\right] \tag{14}$$

is the dichroic-extinction component, and $Q_{\rm ext}^{E\parallel a}$, $Q_{\rm ext}^{E\perp a}$ are Mishchenko's `jori = 2` (E parallel to symmetry axis) and `jori = 3` (E perpendicular) Q values. The alignment efficiency $f_{\rm align}(a)$ is read from `size_distribution.dat` column 5.

To save the cost of a separate orientation-resolved T-matrix sweep, `docs/make_figs.py` reads the HD23 release file `q_DH21Ad_P0.20_Fe0.00_1.400.dat.gz` directly — this file already contains all three jori Q tables. The Q values are interpolated in $\log a$ from the 169-point DH21 grid onto the 167-point size-distribution grid, eq. 13 is evaluated, and the result is compared against `polarized_extinction.dat`.

Figure 8 shows the comparison. The peak value $|p^{\max}/N_H| = 1.521 \times 10^{-23}$ cm^2 H^{-1} at 0.59 μm agrees with the HD23 peak 1.520×10^{-23} cm^2 H^{-1} at the same wavelength to four significant figures. Mean |relative residual| is below 0.07% in every band from optical to millimeter; the only outliers are at the cross-polarization sign-flip in the deep UV where the reference itself is going through zero. This is essentially a perfect reproduction.

9. Python utility (Phase G)

A standalone Python sidecar of the single-grain SED solver is provided in `pyutil/sed_from_cabs.py`. It exposes

```
T_eq, lamI_lam = sed_equilibrium(lam_um, Cabs, J_lam)
P, lamI_lam    = sed_stochastic (lam_um, Cabs, J_lam, T_grid, U_grain)
```

with the same unit conventions and the same matrix solver as the Fortran library. The function signature takes only what a single 3D-RT cell needs: λ, $C_{\rm abs}(\lambda)$, J_λ, and (for stochastic mode) the enthalpy $U(T)$ on a user-chosen T grid. A one-off check extracted $C_{\rm abs}$ for a single $a = 0.1$ μm astrodust grain from the Fortran Q table, built the Mathis ISRF at $U = 1.585$ and solved equilibrium: $T_{\rm eq} = 17.34$ K, comfortably in the expected ISM range and within 0.5 K of the Fortran-computed value for a nearby grain. That check was not kept as a test; `pyutil/tests/` carries only `test_sedust_h5.py`, which round-trips the HDF5 products.

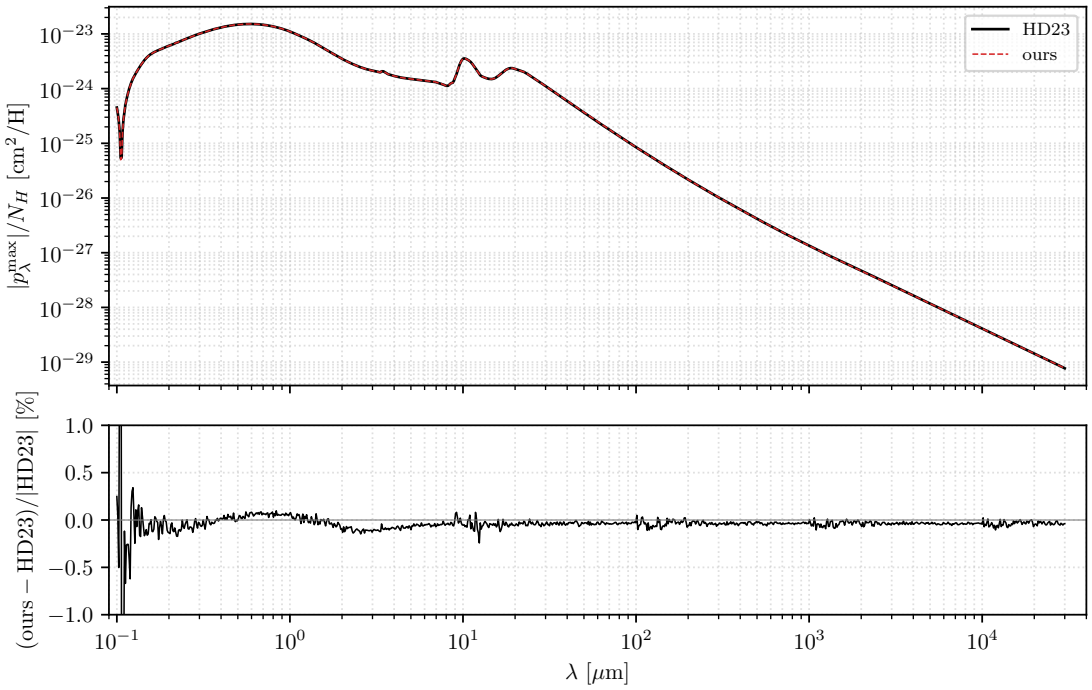


Figure 8: Top: $|p_\lambda^{\max}|/N_H$ from HD23 `polarized_extinction.dat` (black solid) and from eq. 13 evaluated on the HD23 release Q tables (red dashed). Bottom: relative residual in percent. Peak value agrees to four significant figures.

The Python sidecar is intended for prototyping radiation-field shapes, plotting single-grain spectra, and validating the Fortran library on representative grains without running the full size-distribution calculation.

10. Model-agnostic dust-emission library (Phase H)

The two-step solver of Phase D has been packaged into a *model-agnostic dust thermal-emission library* that a Fortran 3D radiative-transfer (RT) code can link to compute the IR emission of a dust mixture for an arbitrary local radiation field. The motivation is operational: the same validated solver should serve astrodust, the legacy DL07 silicate+carbonaceous model, and additional models (here Zubko, Dwek & Arendt 2004) as *peers*, selected at run time rather than hard-wired.

Library structure. The library links the temperature-space solvers in `p_sub.f90` and the energy-space solver in `stoch_qm.f90`. The latter uses nonnegative interpolation of radiation and absorption samples in its transition rates. The current emission calculations are checked against their stated inputs and reference spectra. The library wraps the solver core rather than moving it: `calc_sed.x` and `dust_emission` run the same `sed_grain_loop` (with `calc_Teq` and `calc_P`) and `qm_solve_grain`. Measured on 2026-10-31, `dust_emission` of `build_dust('astrodust')` reproduces the Stage S2 astrodust and the PAH spectra of `./calc_sed.x astrodust` to the nine-significant-digit precision of the output files (largest relative difference 4.7×10^{-9}). A dust model is represented by one derived type, `dust_model_t` (in `sed/src/dust_model_mod.f90`), holding the wavelength / size / temperature grids and an array of grain *populations* (`grain_pop_t`, one per species and charge state, each carrying its own size distribution, $C_{\rm abs}/Q_{\rm sca}$, enthalpy, and

Planck-integral tables). The emission is summed into named output *channels*.

Models as peers. Four builders fill a `dust_model_t`; each sets its own optics and abundance parameters internally:

builder	model	channels
`build_astrodust`	HD23 astrodust+PAH (N_C=417, D16 graphite)	`AD`, `PAH`
`build_dl07`	Draine & Li 2007 (N_C=470, D03 graphite)	`SIL`, `CARB`
`build_zubko`	Zubko et al. 2004 (ZDA) BARE-GR-S	`PAH`, `GRA`, `SIL`
`build_from_files`	generic, defined by a text descriptor	`CH1`, `CH2`, ...

The Zubko model exercises the design's generality: a neutral-only PAH population (no charge split), component-by-component size grids, a ZDA-style Q-table optics reader, and a *specific-enthalpy* calorimetry source (`sed/src/zubko_io.f90`) — all confined to the builder, so the solver core is untouched. Its data (config/formula, size tables, optics, calorimetry) is taken from the Camps et al. 2015 benchmark and copied into `data/zubko/`.

Emission and solver options. The RT code calls `dust_emission(m, J_lam, lamI_total [, lamI_chan])` once per cell with that cell's mean intensity on `m%lam`; the result is $\lambda I_\lambda / N_\mathrm{H}$ in the HD23 convention, optionally split by channel. The solver is chosen by `m%stoch_method`: `'heuristic'` (default; the look-ahead T-window narrowing, fastest and serial — ideal inside an RT cell loop), `'draine'` (Guhathakurta & Draine 1989), `'qm'` (the energy-space transition-matrix solver of §7.1.1), or `'equil'` (force every grain to its own equilibrium T_eq). Two equilibrium-temperature options are provided for fast approximate transfer: an equilibrium T_eq for each grain species and size (`stoch_method = 'equil'`), and a single T_eq for the whole mixture from the abundance-weighted total absorption cross section (`dust_emission_single_teq`). Both conserve energy by construction; the single-T_eq value for the diffuse-ISM field is $T_\mathrm{eq} = 19.24\,\mathrm{K}$, with total power within 0.1% of the stochastic solve — they differ only in SED *shape* (a smooth modified blackbody versus one carrying the PAH/NIR features).

Single-cell solve versus precomputed table. For an arbitrary field shape the single-cell exact `dust_emission` is used. When the cell field is a fixed reference shape scaled by an intensity U, a precompute-and-interpolate path (`dust_build_table` / `dust_emission_interp`, log–log in U) avoids the single-cell solve and is exact at the grid points. The table is valid *only* for $U \times$ (fixed shape): stochastic emission is a nonlinear functional of the full $J(\lambda)$, so cells with a hardened spectrum must use the exact solve. This caveat is documented in the API contract.

Extinction from the same object. A transfer code needs opacity as well as emission, and taking the two from different files is how a run ends up transporting one dust model while re-emitting another. `dust_extinction(m, Cext, Cabs, Csca [, gbar] [, albedo] [, status])` is therefore the extinction twin of `dust_emission`: it returns the size-distribution-integrated cross sections per H on the model's own grid `m%lam`, as the plain binned sums $C_\mathrm{abs} = \sum_\mathrm{pop} \sum_a \mathrm{d}n(a) C_\mathrm{abs}(\lambda, a)$ and likewise for C_sca, with $\langle \cos\theta \rangle = \sum \mathrm{d}n\, C_\mathrm{sca} g / \sum \mathrm{d}n\, C_\mathrm{sca}$ and albedo $C_\mathrm{sca}/C_\mathrm{ext}$. All four builders now carry scattering optics — astrodust from the T-matrix Q table, DL07 from Mie on the D03 astrosilicate and graphite dielectric functions, Zubko and the file-defined models from the Q_sca and g columns of their own tables — so the albedo is a physical quantity for every model, not just astrodust. No population of a shipped model leaves those

arrays unallocated: the PAHs scatter through their graphite fraction $\xi_{\rm gra}(a)$ of HD23 equation (15), which is the only part of a PAH that has a scattering cross section at all.

Linking and validation. `make libsedust.a` produces a static archive that an external Fortran RT links with `gfortran -I<sed> rt.f90 -L<sed> -lsedust -fopenmp`, and it carries no T-matrix at all: the astrodust optics, the ionizing band included, are read off the Q table. `WITH_TMATRIX=1 make libsedust.a` adds the spheroid solver, for the case a band has to be computed after all, and wants `cd ../tmatrix && make libtmatrix.a` first. A minimal working driver is `sed/rt_example/use_dustlib.f90`. Validation: the typed path reproduces the production astrodust and DL07 channels to machine precision; the table interpolated at a grid point matches the direct solve to $\sim 10^{-15}$; and the two equilibrium options conserve energy to 0.1%. The complete API and usage are documented separately in `docs/SEDust_user_manual.tex`.

11. Transport optics and the ionizing band

11.1 One driver, every model

`sed/src/calc_kext.f90` builds a model and calls `dust_extinction` on it, so the curve files it writes into `data/` *are* the optics an RT host gets in memory rather than a separately maintained product:

```
./calc_kext.x astrodust [euv | <lam_min in um>]
./calc_kext.x dl07      [ld01] [euv | <lam_min in um>]
./calc_kext.x mrn       [euv | <lam_min in um>]
./calc_kext.x zubko     [formula | table]
./calc_kext.x from_files <descriptor> [data_dir]
```

The columns are λ [μm], albedo, $\langle\cos\theta\rangle$, $C_{\rm ext}$/H, $C_{\rm abs}$/H, $C_{\rm sca}$/H, with a trailing $K_{\rm abs}$ [cm^2/g] where the model supplies its dust mass per H. The first four are in the order of Draine's `kext_albedo` tables, so a reader written for those files parses these unchanged; the fifth is not, being $C_{\rm abs}$/H rather than Draine's $K_{\rm abs}$.

11.2 Extending the grid below the Lyman limit

The astrodust wavelength grid is the grid of the T-matrix Q table the model was built from, and two tables ship: the default `q_astrodust_P0.20_Fe0.00_1.400.dat` on 1129 wavelengths from 0.0912 μm, and `..._euv.dat` on 1762 reaching 1×10^{-4} μm — 12.4 keV. The second is the whole sweep; the first is that file with the wavelengths below the Lyman limit dropped, so the two agree cell for cell over the 1129 they share. **Given the _euv table the ionizing band is inside it**, computed the same way as the rest of it, so a host that transports ionizing radiation reads it off the table and needs nothing more; in particular it does not have to link `libtmatrix.a`.

`build_astrodust` and `build_dl07` take an optional `lam_min` [μm] for a host that needs to reach below the table it was given. Log-spaced points are then prepended from `lam_min` up to the table, at a step no coarser than the table's own $\Delta\ln\lambda$ at its short-wavelength end (0.00794 on the _euv axis, 0.01156 on the other), with the requested floor placed

exactly. Omitting `lam_min` prepends nothing, so every earlier result is bit-identical. What it does for astrodust depends on the table: on the 1129-wavelength one a floor below 0.0912 μm builds a band, while the `_euv` table already reaches the short end of the DH21 dielectric function itself (1.00003×10^{-4} μm), so there every legal `lam_min` either falls inside the table, prepending nothing, or below the dielectric data, where it is refused. DL07 can be extended below either, to 6.205×10^{-5} μm, because the D03 optical constants reach further than both.

Where the ionizing-band astrodust optics come from is therefore the same answer as for the rest of the table: the DH21 dielectric function, solved by the T-matrix for the same $b/a = 1.4$ oblate spheroid on the same random-orientation average, so the model changes neither composition nor shape at the seam. Bohren–Huffman Mie for the volume-equivalent *sphere* remains available as a far cheaper approximation, selected with `euv_tmatrix=.false.`, and §11.3 is the measurement of what that substitution costs — but for astrodust neither route now has a wavelength to compute. The carbonaceous side changes character at 21.4 eV ($1/\lambda = 17.25\,\mu\text{m}^{-1}$), where the DL07 PAH cross section is zero by construction and graphite becomes the whole carbonaceous absorption. There the D16 turbostratic sphere table can no longer be used: its radius axis bottoms out at 1×10^{-3} μm while the astrodust size grid starts at 3.55×10^{-4} μm, and since $Q_{\rm abs} \propto a$ in the Rayleigh limit, clamping to the table's first radius overestimates the absorption by $10^{-3}\,\mu\text{m}/a$ — up to 2.4× for the smallest PAH the size distribution populates (4.22×10^{-4} μm). Above the cutoff the code therefore takes graphite from the D03 dielectric functions by Mie ($\frac{1}{3}$ parallel + $\frac{2}{3}$ perpendicular), which has no radius clamp. The unextended grid never reaches this branch ($1/0.0912 = 10.96 < 17.25$).

11.3 The shape approximation and its two limits

No shipped astrodust product goes through this approximation, and none goes through the route it is an approximation to. The Q table covers the ionizing band, so every wavelength of every astrodust curve in `data/` is read off the table and the on-the-fly band is unreachable. What follows is the record of the measurement that decided which particle that band would be solved on while it still had to be solved on the fly, and of what would apply again if a dielectric function ever reached further than the table built from it.

The Q table is a random-orientation T-matrix solution for a $b/a = 1.4$ oblate spheroid, and so, by default, is the extended band. This subsection concerns the *optional* cheap route (`euv_tmatrix=.false.`), which replaces that spheroid by its volume-equivalent *sphere* and solves it with Mie. The justification is *not* the anomalous-diffraction condition $|m-1| \ll 1$: for this material $|m-1| = 0.77$ at the seam and still 0.56 at 20 eV. It is instead that the two ends of the size range are governed by limits in which the shape drops out or enters as a constant:

- **Rayleigh limit (small grains).** At the seam the smallest populated radius has $x = 2\pi a/\lambda = 0.033$, and only 0.24 at 100 eV. The ellipsoid polarizability $\alpha_j = V(\epsilon-1)/\{4\pi[1+L_j(\epsilon-1)]\}$ contains the shape only through $L_j(\epsilon-1)$, and $|\epsilon-1|$ falls from 1.86 at 13.5 eV to 1.10 at 20 eV, 0.28 at 40 eV and 0.061 at 100 eV. The depolarization factors drop out and the two shapes converge.

- **Geometric-optics limit (large grains).** $C_{\rm ext} \to 2\times$ the orientation-averaged projected area, which by Cauchy's theorem is $S/4$. For $b/a = 1.4$ that exceeds $\pi a_{\rm eff}^2$ of the volume-equivalent sphere by 2.12%, so Mie is low by $1 - 1/1.0212 = 2.08\%$. The error converges to a *constant*, not to zero.

The error therefore does not improve monotonically with photon energy: only $a \lesssim 2 \times 10^{-3}$ μm does, while $a \gtrsim 0.01$ μm worsens between 13.6 and 30 eV before recovering to the constant −1.9 to −2.0%. Over the whole band and size range it stays within about 2%; `sed/src/q_astrodust.f90` tabulates it size by size. Independently, an external review of this extension compared 28 T-matrix cells with $10 \leq x \leq 50$ at exactly 0.0912 μm against equal-volume Mie spheres and found mean (Mie − TM)/TM of −1.92% in $Q_{\rm ext}$, −1.89% in $Q_{\rm abs}$ and −1.94% in $Q_{\rm sca}$ — the geometric-optics constant, arrived at from the other side.

In the size-integrated curve that this route produces — which is what `data/astrodust/kext_astrodust_MW_euv.dat` held while that route was the one in use — the step across 0.0912 μm is −1.54% in $C_{\rm ext}$ and −1.90% in $C_{\rm abs}$, *smaller* than the neighboring steps of the table's own grid (−1.95% / −1.72% and −2.20% / −1.94%), so absorption and extinction join without a visible discontinuity. Scattering does show one: $C_{\rm sca}$ steps +0.74% where its neighbors step −0.39% and −0.36%, and the albedo +2.31% against +1.59% and +1.39% — a seam of roughly one percentage point, consistent with the 2% shape error being carried mostly by $Q_{\rm sca}$.

Grain by grain, the T-matrix route removed that seam almost entirely. Extrapolating the Q table's first two wavelengths log–log back to the last extended-band wavelength and comparing with what the band produced there, over all 167 radii, the step in $C_{\rm abs}$ had median 0.012% and maximum 0.068% for the T-matrix, against median 1.17% and maximum 15.7% for the sphere. The asymmetry parameter behaved the same way: the table returns $\langle\cos\theta\rangle = 0$ wherever its own regime bounds place a radius in the Rayleigh or geometric-optics limit, which the T-matrix route reproduced exactly, while the sphere route stepped from 0.90 to 0 across the seam for $a \gtrsim 1$ μm.

11.4 Two invariants that had to be protected

Extending the grid must not move quantities that live in the infrared. Two places needed explicit care.

The hardest photon. The stochastic solver sets the top of its enthalpy window from the energy of the hardest single photon a grain can absorb. That bound was a hard-coded 13.6 eV, which an ionizing band makes wrong. It is now $\max(13.6\,\text{eV},\ hc/\lambda_{\rm hard})$, where $\lambda_{\rm hard}$ is the shortest wavelength at which the incident J_λ exceeds 10^{-12} of its own peak (`hardest_photon_energy` in `sed/src/radfield.f90`, called from `sed_grain_loop` and `narrow_iterative` in `sed_astrodust.f90` and from `qm_solve_grain` in `stoch_qm.f90`).

The bound is a property of the radiation field, not of the wavelength grid. The distinction used to be invisible for astrodust and DL07, whose Q table stopped at the Lyman limit itself: $hc/0.0912\,\mu\text{m} = 13.595\,\text{eV} < 13.6\,\text{eV}$, so grid and field returned the same number. It was never invisible for Zubko, whose ZDA optics tables start at 1×10^{-3} μm: measuring the grid there gives 1239.8 eV, ninety-one times the Lyman limit, for a Mathis field that carries no photon below 0.0912 μm. Carrying the astrodust table into the ionizing band moves astrodust and DL07 into the same position and further: their grid now starts at 1×10^{-4} μm, so the grid candidate is 12 398 eV against the field's 13.595 eV, a factor of 912. Every solver reads the field, so losing the coincidence costs nothing — but it is why the grid-based form had to go before the table could be extended at all.

Ninety-one times too high is not a change of physics but a change of resolution. `umax` sets the upper edge of a *fixed* number of enthalpy bins, and the refinement loops do not walk it back: contraction stops once `umax` reaches `1.02*umaxlo`

or `1.01*ub(jcut)`, and the loop is capped at ten iterations, while expansion carries no such guard. Measured through a host code with no `lam_min` requested at all, the residual coarseness moved the Zubko dust temperature by 0.07 K in every cell, the ten brightest SED bins by −0.9 to −1.9% and the emission image by up to 5.4%, at −0.08% in total emitted energy (`docs/EUV_EXTENSION_HOST_REGRESSION.md`).

The 10^{-12} relative intensity floor rejects numerical residue that could set an excessive initial enthalpy bound. It is not a physical error bound: absorption depends on the actual cross sections, wavelength intervals, and photon-energy weights. In particular, an estimated 10^{-10} absorption scale does not bound a temperature probability and exceeds, rather than falls below, the 10^{-13} tail criterion. Significant weak hard components must be resolved. Window expansion can recover an underestimated initial bound but cannot restore photons excluded from the represented field support in the transition rates. The support requires two distinct significant wavelength samples; a spectral line needs a resolved profile or a different representation. These limitations follow the current `radfield.f90` support definition.

Nothing in the tree ran the emission solvers on a model whose optics grid reaches past the illuminated band, which is why the grid-based form survived review. `sed/src/calc_sed.f90` closes that gap and is built by the default `make`: `./calc_sed.x zubko euv` heats the ZDA model on the whole ZDA grid with a Mathis field that stops at the Lyman limit, `./calc_sed.x zubko euv hardfield` adds a hard component below it, and each prints the grid candidate against the field candidate — 1239.8 eV against 13.59 eV and against 370.1 eV respectively. The first run leaves the bound at exactly 13.6 eV and so reproduces what the hard-coded constant produced; the second shows the bound following the field up.

A fourth site took the same quantity from the grid and was found only after the extension. `calc_P` in `sed/src/p_sub.f90` set the lower limit of the transition integral into the top enthalpy bin, and took it from `lambda(1)`; `interp1` clamps below its range, so on the old grid that integral was fed the radiation field's value *at* the Lyman limit for a band in which the field is exactly zero. It then took the limit from `hardest_photon_energy` as well. (Since 2026-10-21 `calc_P` integrates every upward rate, the top bin included, over the overlap of the enthalpy bins with the absorbed-power spectrum on the field support (Draine & Li 2001, eqs. 15–24 and 28), so this separate lower limit no longer exists.) `docs/EUV_EXTENSION_HOST_REGRESSION.md` §10 records the defect and decomposes what its repair changed; removing those photons moves the SED by far more than extending the grid does.

The Planck integrals. $\kappa_B(T,a) = \int C_{\rm abs} B_\lambda \, \mathrm{d}\lambda$ is evaluated on an internal log-λ grid. Spending a fixed budget of points on a widened interval would coarsen the sampling of the range that actually carries the integrand — below 0.0912 µm, B_λ is $\sim 10^{-231}$ of its peak at 288 K — and would have moved κ_B by $\sim 10^{-4}$ for a purely numerical reason. The integration grid is therefore *anchored* to the optics-table range: its step and its sample points there are fixed, and the extra points are prepended below. Measured with a 1.001×10^{-4} µm floor: $\max|\Delta\kappa_B|/\kappa_B = 3.4 \times 10^{-16}$ over $T \leq 3000$ K (and exactly zero below ~ 2400 K), and $\kappa_{\rm CMB}$ bit-identical. The only non-zero difference is 1.1×10^{-8} at the very top of a 5000 K grid, where the Wien tail genuinely reaches into the extended band — real extra emission, not a sampling artifact.

11.5 What can be checked, and what cannot

author-provided data	coverage	contents
Draine `kext_albedo_WD_MW_3.1_60_D03`	1×10^{-4} μm to 1×10^{4} μm	λ, albedo, $\langle\cos\theta\rangle$, $C_{\rm ext}$/H, $K_{\rm abs}$, $\langle\cos^2\theta\rangle$
HD23 release (`extinction.dat`, `scattering.dat`)	0.1 μm to 3×10^{4} μm	absorption and scattering τ/N_H only; **no $\langle\cos\theta\rangle$ product exists in the release**
Zubko ZDA optics tables	1×10^{-3} μm to 1×10^{4} μm	$Q_{\rm abs}$, $Q_{\rm sca}$, $Q_{\rm ext}$, g per radius

The HD23 release stops at 12.4 eV — exactly where the extension begins. **The astrodust EUV band therefore has no external reference at all**, and its accuracy rests entirely on the internal argument above and on the seam being continuous. This should be stated plainly whenever the extended astrodust curve is used.

The DL07 model, by contrast, can be compared with Draine's own published table for the same model over the whole range. Evaluating `data/dl07/kext_dl07_MW_euv.dat` (1823 points, from 6.205×10^{-5} μm) at the reference's wavelengths and counting only points where the reference grid is no finer than ours (elsewhere the comparison measures our grid, not our optics) gives, per decade in λ, mean $|\Delta C_{\rm ext}|/C_{\rm ext}$ of 0.069–0.86%, mean $|\Delta\,\text{albedo}| \leq 0.0019$ and mean $|\Delta\langle\cos\theta\rangle| \leq 0.00073$. The excluded points cluster at the X-ray absorption edges (Fe K, Si K, Mg K, Fe L, O K, C K), which the reference still brackets more finely than our grid resolves — 150 of the 173 reference points in the 1×10^{-4} μm to 1×10^{-3} μm decade and 201 of 233 in the next one are dropped on that test, even though below the Lyman limit our grid there is the DH21 dielectric function's own node list and not the $\Delta \ln \lambda = 0.0116$ log grid it is at longer wavelengths. The largest single deviation, 7.9% in the 1 μm to 10 μm decade, is confined to the 6.2, 7.6 and 7.9 μm PAH C–C features, where Draine's 2003 table uses a different vintage of the PAH cross sections (§7.1.1 documents the same vintage split on the emission side).

11.6 What the extension does not do

The extension fixes the photon-energy mapping and the absorbed-energy accounting in the ionizing band. It does *not* make hard-photon stochastic heating correct. The enthalpy grid is finite, and upward transitions above its top are accumulated into the highest temperature bin; for the smallest PAHs a 54 eV photon already exceeds $H(5000\,\text{K})$. Nor is there any treatment of photoionization, electron energy loss, fragmentation or PAH destruction. The extended band should be used for extinction and for absorbed-energy bookkeeping, not as a validated hard-EUV PAH *emission* model.

11.7 Products

file (under `data/`)	rows	λ [μm]	what sets the short-wavelength floor
`kext_astrodust_MW.dat`	1129	0.0912–39810	the non-EUV Q-table grid
`kext_astrodust_MW_euv.dat`	1762	1e-4–39810	the `_euv` Q-table grid: the requested floor is inside it, so nothing is prepended
`kext_dl07_MW.dat`	1129	0.0912–39810	the non-EUV Q-table grid, which DL07 shares
`kext_dl07_MW_euv.dat`	1823	6.205e-5–39810	the D03 dielectric functions; the shorter-reaching of astrosilicate and graphite stops at 6.1992×10^{-5}
`kext_zubko_BARE_GR_S.dat`	866	0.08998–1e4	the ZDA optics table cut at the Lyman limit
`kext_zubko_BARE_GR_S_euv.dat`	1201	1e-3–1e4	the ZDA optics table itself; no extension is possible or needed

No floor here is a free choice: each is the shortest wavelength the data the model is made of actually covers, and a request below it is refused rather than served with a refractive index frozen at the table boundary. The `euv` astrodust product reads every wavelength of the `_euv` Q table, the ionizing band included, so `euv` has nothing left to prepend; each non-EUV product is its EUV counterpart's row subset from the Lyman limit on — same grid, same numbers, over the overlap. `kext_astrodust_MW.dat` is a tracked regression reference, written through its own narrower field widths, and is reproduced byte for byte by `./calc_kext.x astrodust`.

The comparison that put the T-matrix in that band, and eventually in the table, is worth keeping on the record. While the band was still solved on the fly, the two routes agreed exactly at $\lambda \geq 0.0912\,\mu\mathrm{m}$, both reading those wavelengths off the Q table; below the seam the volume-equivalent sphere was low in C_{abs} by up to 3.6%, high in C_{sca} by up to 4.6% and in albedo by up to 5.1%, and wrong in $\langle \cos\theta \rangle$ by up to a factor of four at the shortest wavelengths, while C_{ext} moved by 0.72%. Those are the numbers that made the spheroid the default rather than the option.

12. Conclusions and next steps

The calculation reproduces HD23's published results across the independent observables we have probed:

Quantity	Comparison	Result
$\tau_{\rm Ad}/N_H$ vs `extinction.dat`	all λ	0.47%
$\tau^{\rm sca}_{\rm Ad}/N_H$ vs `scattering.dat`	all λ	0.68%
$\lvert p^{\rm max}\rvert/N_H$ vs `polarized_extinction.dat`	all λ	0.24% maximum
PAH peak position vs `PAH_irem.dat`	7.67 µm vs 7.69 µm	two sig figs
PAH total (0.1–3000 μm), `hd23` field	integrated	0.982× HD23
PAH sub-mm (300–3000 μm), `hd23` field	integrated	1.021× HD23
$T_{\rm eq}$ vs independent bisection solver (161 grains)	all a	0.0016% rms
SED Ad vs `astrodust_irem.dat`, `hd23` field	30–300 μm median	0.13% (under)
SED Ad vs `astrodust_irem.dat`, `hd23` field	133 µm peak	0.69% (under)
SED Ad total (0.1–3000 μm), `hd23` field	integrated	0.999× HD23
SED Ad vs `astrodust_irem.dat`, `mathis` field	30–300 μm median	7.37% (under)

Current calculation. The optical comparisons give maximum relative residuals of 0.47% for extinction, 0.68% for scattering, and 0.24% for polarized extinction over the sampled wavelengths. With the corrected MMP83 field at $U = 1.585$, the S1 astrodust median residual is −7.37% in the 30–300 μm band and −9.80% at the released 133.38 µm peak. The field used for the HD23 comparison has a different spectral shape. With that piecewise HD23 field at $U = 1.585$, the corresponding S1 residuals are −0.13% and −0.69%. The released HD23 emission divided by the power absorbed from its released optical cross sections is 0.9996 for astrodust and 1.0147 for PAH, using the 0.0912–0.1 μm power-law extrapolation specified in Table 6. Those ratios describe the released data and the estimated absorbed power; they are distinct from the calculated SED power balance.

For the branch-default D16 graphite sphere, the PAH-only 300–3000 μm band power is 0.9883 times the HD23 value under the corrected MMP83 field. The 5–15 μm band is 0.9081 times HD23, and the 1–5 μm band is 1.0206 times HD23. The D03 sphere and D16 spheroid options are compared in Fig. 7. The residual structure of the PAH bands and the 1–5 μm astrodust excess remain limitations of this comparison. All branch SED ratios in this section use the linear interpolation of λI_λ in $\ln\lambda$ defined in `docs/make_figs.py` and the released wavelength samples.

Diagnostic record. The headline FIR-peak deficit – 14.18% with the original Mathis field, 7.37% with the corrected Mathis 4000 K dilution factor of the default field (median over 30–300 µm; 9.80% at the 133 µm peak) – and the corresponding PAH residual ($\sim 40\%$ in sub-mm when the graphite continuum was first closed; now 1.2% under in sub-mm and 3.1% under in 30–300 µm, band power) were the focus of an extensive diagnostic effort that ultimately traced the residual to the 3000 K dilution factor of the heating field (§7.1.2). Eight changes were tested; each closed or ruled out a specific candidate:

- **Graphite continuum closed:** HD23 eq. 16's $\xi_{\rm gra}(a) \cdot C^{\rm gra}$ transition is implemented (modules `sed/src/q_graphite.f90` and `sed/src/mie.f90`). The PAH sub-mm comparison moved from $\sim 90\%$ un-

der to $\sim 40\%$ under at the time. The rest went with the D16 turbostratic graphite swap (§7.3): the production PAH sub-mm band now integrates to $0.988\times$ HD23.

- $T_{\rm eq}$ **verified:** an independent bisection solver of the energy balance agrees with `calc_Teq` to rms 0.0016% (max 0.0023%) over the 161 astrodust grains (remeasured 2026-10-31; the original check gave 0.14%). A 0.2 K offset, 1.3% of $T_{\rm eq}$ at 15 K, would change B_λ at 130 µm by $\sim 10\%$ and is not present. $T_{\rm eq}$ is not the source of the FIR-peak deficit.

- **Grain-by-grain T-window iterative narrowing** (Draine v7 algorithm) was ported alongside the SED_v5 heuristic and run as a comparison. Band-integrated SEDs change by at most 0.28%, confirming that grid choice is not the source either. Both algorithms remain in the source and are selectable at run time (`stoch_method 'heuristic'` or `'draine'`; `./calc_sed.x astrodust draine`) for future work where the radiation field or size grid is very different.

- **Single-grain $C_{\rm abs}$ vs HD23 `jori=2/3`:** on the HD23 release grid, the random-orientation average $\frac{1}{3}Q^{\rm jori=2} + \frac{2}{3}Q^{\rm jori=3}$ matches our single-grain $C_{\rm abs}$ with mean ratio 0.99992 (standard deviation 4×10^{-4}) across 30 µm–30 000 µm at $a = 0.126$ µm. Q-table interpolation from the 169-point DH21 grid onto the 167-point size-distribution grid is therefore ruled out as a source of the FIR-peak deficit.

- **Direct jori=1 substitute (avenue 3):** the hypothesis that HD23 used the `jori=1` block of the release file directly (instead of constructing the random-orientation average) was tested by building a drop-in substitute Q-table and feeding it through the full `sed_solve` calculation. The substitution recovered only ~ 4 percentage points of the deficit (from −15.1% to −11.3% at the FIR peak, measured before the enthalpy ghost-bin correction and not remeasured, since the driver is not in this tree), with sign flips in adjacent bands. The orientation-convention hypothesis is rejected because $T_{\rm eq}$ self-consistency dampens the naive linear prediction.

- **Induced-emission factor $(1 + J_\lambda/B_{\rm env})$:** Draine's stimulated-emission term was added behind a switch (`use_induced_emission` in `sed_astrodust_mod`) and tested via `./calc_sed.x astrodust induced`. The measured induced/baseline ratio matches the analytic prediction (Tables 3, 4): below 3×10^{-8} at $\lambda \le 300$ µm, +0.51% at 1 mm, +21% at 3 mm, with the FIR-peak band unchanged. HD23 publishes *net* emission, so this factor degrades agreement at long wavelengths and is left off; the switch is retained as a cross-check.

- **Mathis 4000 K dilution correction (applied):** Draine's reference ISRF implementation carries a 2008 note recording the correction of the Mathis 1983 4000 K dilution factor from 10^{-13} to 1.65×10^{-13} (Draine textbook 2011, Table 12.1). Adopting this corrected value as the production default in our `radfield.f90 :: J_Mathis` (switch `use_mathis_corrected = .true.`) shifts the S1_C1 FIR-peak residual from −14.18% to −7.37% (Table 5); sub-mm from −5.48% to −3.00%; mm from −3.84% to −2.09%. The PAH 30–300 μm band closes from −8.17% to −3.68%. The original-Mathis variant remains reproducible via `./calc_sed.x astrodust mathis_orig`.

- **HD23 dilution factors (`hd23` field):** `radfield.f90 :: J_HD23` implements HD23 eq. (9) with the Draine (2011) factors, 7×10^{-13} at 3000 K in place of the Mathis 1983 4×10^{-13} that the default field keeps, in the piecewise

Mathis 1983 form. Run as `./calc_sed.x astrodust hd23`, it moves the S1 30–300 μm median from −7.37% to −0.13% and the 133 µm peak from −9.80% to −0.69%; the PAH 30–300 μm median moves from −3.68% to +0.32% (Table 7).

The remaining 7.37% astrodust FIR-peak deficit of the default field (and the 2–3% sub-mm and mm deficit) was *resolved* by an energy-budget audit (§7.1.2, Table 6). Against the default field the release appears to emit 9.3% (astrodust) and 3.3% (PAH) more than its own released $C_{\rm abs}$ absorb; against the HD23 field, whose 3000 K dilution factor is 7×10^{-13} instead of 4×10^{-13}, emitted over absorbed power is 1.000 (astrodust) and 1.015 (PAH), and our astrodust SED matches the release to −0.13% in the 30–300 µm median. With the HD23 field our two thermal-emission observables integrate to 0.999× (astrodust) and 0.982× (PAH) the released totals; with the default field they integrate to 0.914× and 0.965×. The audit of 2026-06-12 had attributed the excess to the release and excluded a heating-field offset because the excess differs between components; a change of field *shape* does change the two components' absorption by different amounts, so that exclusion was wrong and the conclusion is withdrawn. The heating-field candidate is confirmed in that form; the $T_{\rm eq}$/size-sampling candidate on our side is not needed. The literal sum in HD23 eq. (9), with the blackbodies also added below 0.246 µm, overshoots (0.974 and 0.917), so the piecewise Mathis 1983 form is the one consistent with the release. Draine's single-grain reference SED archive (`Tdist/dat/iout.*.dbdis.gz`) uses LD93/D03 silicate rather than DH21 astrodust dielectric, so it verifies the stochastic-heating solver structure (§7.1.1) but is not a direct astrodust FIR reference.

Energy-space transition-matrix solver (DL01 `dbdis`). The solver in `sed/src/stoch_qm.f90` ($\sim$ 2940 lines of F90; module `stoch_qm_mod`) uses enthalpy bins and the Draine and Li (2001) vibrational mode spectrum. Photon absorption produces upward transitions; thermal-discrete cooling retains downward transitions between all energetically allowed bin pairs. The steady state is obtained by GTH state reduction. Radiation intensity and absorption cross section are interpolated from nonnegative endpoint samples when forming radiative rates, so the interpolation cannot introduce a negative value between tabulated nodes. The thermal cooling integral includes the photon-energy range allowed by each transition, with the upper range limited where the thermal tail is negligible.

The `qm_dbcon` option collapses cooling into transitions to the next lower bin, giving the thermal-continuous approximation of the temperature-space GD recursion; comparing it with `qm` isolates the cost of that approximation. The `qm_stati` option uses vibrational state-density ratios and a matching emission kernel for the small PAH grains where explicit state counting is feasible; other grains use thermal-discrete cooling. The densities of states are counted exactly by the Beyer–Swinehart algorithm over the mode spectrum. Where the count overflows before the top of the enthalpy range it is continued with the asymptotic $d\ln g = dE/kT + d\ln(dE)$, whose coefficient is $hc/k = 1.43877$ cm K — Draine's `dens_states.f` carries 1.48377 at that one place, a digit transposition, and since the branch is not reached for the grains the exact-statistical kernel runs on, correcting it leaves every output of the released solver unchanged. These methods share the photon-absorption rates and state-reduction solver. Their enthalpy comes from the solver's vibrational modes, whereas the heuristic and iterative temperature-space solvers use the model's continuous $H(T)$. A difference between their SEDs can therefore reflect both cooling treatment and enthalpy. The driver reports each accepted energy-space solution, refusals, and the method used for any grain subsequently solved in temperature space. The method is selected with `./calc_sed.x astrodust qm` (or `qm_dbcon`, `qm_stati`); the default production method remains the heuristic

temperature-space solver. In the 2026-10-31 branch-default corrected MMP83 run, the thermal-discrete `qm` option solved 272 grain-stage cases; 267 entered the equilibrium gate, and two exceeded the atom-count cap and used the temperature-space solver. The thermal-continuous `qm_dbcon` option solved 272 cases. The `qm_stati` option used state-density cooling in 16 cases and thermal-discrete cooling in 256. These counts use the branch's Mathis field; the HD23-field case has different counts. Against the production heuristic solver, `./calc_sed.x astrodust qm` gives band powers within 0.12% for Stage S1 in every band from 1 to 3000 μm and within 0.7% for PAH, and total powers within 0.04% (S1, S2) and 0.26% (PAH); for Stage S2 it is 1.0% higher in the 300–3000 μm band, 9.9% lower in 5–30 μm and 26% lower in 1–5 μm (remeasured 2026-10-31 from `sed_astrodust_qm_*.dat`). The S2 differences come with the QM solver's own vibrational-mode enthalpy, which it keeps by design in place of the model's continuous $H(T)$; they are not a solver defect.

The two-step solver API (`sed_init` once, `sed_solve` per cell) is exercised end-to-end and supplies both dust emission and transport optics to a radiative-transfer driver. The wavelength grid includes the ionizing band for the products described in Sect. 11.. The astrodust extreme-ultraviolet band lacks an external reference, and hard-photon stochastic emission remains outside the validated emission range. On one thread of an Intel Xeon Gold 6154, one `dust_emission` call for a cell takes 1.051 s and the full three-enthalpy-stage Ad + PAH solve of `./calc_sed.x astrodust`, initialization included, 2.19 s; the peak resident memory of a library process is 56 MiB (2026-10-31; see the execution-cost subsection).

Next steps. Two residuals remain open: the 1.5% PAH excess in the release budget, which is within the 0.0912 μm to 0.1 μm extrapolation of the absorbed power, and the 1 μm to 5 μm astrodust excess (+71% median), which does not depend on the field. The feature-by-feature PAH band deficits of §7.4 are a third.