

# Polarization in SEDust: Implementation Notes

## Abstract

These are maintenance notes for the polarization capability added to SEDust: where the polarized optics come from, how they are assembled, which implementation choices were made and why, and how far the numbers have been checked. The reader is assumed to be someone extending or repairing this code, not someone calling it. The calling interface is documented in `SEDust_user_manual.pdf`; the underlying physics is documented in `aligned_grain_polarization.pdf`. Neither is repeated here.

## Contents

|                                                                     |          |
|---------------------------------------------------------------------|----------|
| <b>1. Scope, and what lives elsewhere</b>                           | <b>2</b> |
| <b>2. The data source</b>                                           | <b>3</b> |
| <b>3. Derived quantities</b>                                        | <b>4</b> |
| <b>4. Module map</b>                                                | <b>4</b> |
| <b>5. Implementation decisions</b>                                  | <b>4</b> |
| 5.1 The accumulation sites, and the two that hide . . . . .         | 4        |
| 5.2 Why the 'qm' rescaling is exact . . . . .                       | 5        |
| 5.3 The $P(T)$ solver was not touched . . . . .                     | 5        |
| 5.4 Mixed orientation-average conventions . . . . .                 | 5        |
| 5.5 $f_{\text{align}}$ is the fit, not the release column . . . . . | 6        |
| 5.6 Names and surface area . . . . .                                | 6        |
| 5.7 Why the accessor exists now . . . . .                           | 7        |
| 5.8 Compression handling, and a scratch file that leaked . . . . .  | 7        |
| <b>6. Verification</b>                                              | <b>8</b> |
| 6.1 Polarized extinction against the release . . . . .              | 8        |
| 6.2 Polarized emission . . . . .                                    | 8        |
| <b>7. First-principles orientation-resolved optics</b>              | <b>9</b> |
| 7.1 Rayleigh regime, $x < 0.1$ . . . . .                            | 9        |
| 7.2 T-matrix regime, $0.1 \leq x \leq 50$ . . . . .                 | 9        |
| 7.3 Geometric-optics regime, $x > 50$ . . . . .                     | 10       |
| 7.4 Dispatch, generation, and drop-in . . . . .                     | 11       |
| 7.5 How far it has been checked . . . . .                           | 11       |
| 7.6 Computing named wavelengths . . . . .                           | 12       |
| 7.7 The polarized band below the Lyman limit . . . . .              | 12       |
| 7.8 Certifying the T-matrix above $x = 50$ . . . . .                | 13       |

|                                                                     |           |
|---------------------------------------------------------------------|-----------|
| <b>8. Aligned-grain scattering matrix and the extinction matrix</b> | <b>14</b> |
| 8.1 The engine . . . . .                                            | 14        |
| 8.2 Symmetries, verified not assumed . . . . .                      | 15        |
| 8.3 The generator: one code path . . . . .                          | 16        |
| 8.4 Products and the $\eta$ contract . . . . .                      | 17        |
| 8.5 The library: two strictly separated layers . . . . .            | 17        |
| 8.6 Birefringence through dust_extinction . . . . .                 | 19        |
| 8.7 Verification . . . . .                                          | 19        |
| 8.8 Which wavelengths the table carries . . . . .                   | 19        |
| 8.9 Why the library reads tables and does not solve . . . . .       | 20        |
| <b>9. Status codes, and the relation to v1.00</b>                   | <b>20</b> |
| 9.1 sed_init/build_astrodust status codes . . . . .                 | 20        |
| 9.2 v1.20 is v1.00 plus polarization . . . . .                      | 20        |
| <b>10. Limits and extension paths</b>                               | <b>21</b> |
| 10.1 Cost of build_Cpol . . . . .                                   | 22        |
| <b>11. Reproducing the numbers</b>                                  | <b>22</b> |
| <b>12. References</b>                                               | <b>23</b> |

## 1. Scope, and what lives elsewhere

Three documents cover this work and they do not overlap (Table 1). This one answers only the question *how was this built and why that way*. Where physics is needed to justify a coding choice, the relevant section of the theory report is cited rather than restated.

| Question                                                                                                                   | Document                                       |
|----------------------------------------------------------------------------------------------------------------------------|------------------------------------------------|
| Why polarization at all; alignment theory; the polarized transfer equation; sign and index conventions; validation targets | aligned_grain_polarization.pdf (§1–§5, §9–§11) |
| How an RT code calls the library: signatures, units, what is already integrated                                            | SEDust_user_manual.pdf, §5                     |
| How the code is built, why each choice was made, what was measured                                                         | this note                                      |

Table 1: Division of the polarization documentation.

**Where this is going.** The polarization capability is not an end in itself. The destination is the polarized radiative transfer of Seon (2018): dust scattering plus dichroic extinction by aligned grains, computed with grain optics derived from first principles, in the Monte Carlo formalism of Peest et al. (2017) as extended to aligned spheroids by Peest et al. (2023). Under that contract MoCafe supplies the alignment assumptions (the  $f_{\text{align}}$  profile and the local magnetic-field geometry) and SEDust returns every material quantity in matrix form: the  $4 \times 4$  extinction matrix  $K(\theta_i)$  with dichroism and birefringence, the phase matrix  $Z_{\text{al}}(\theta_i; \theta_s, \phi)$  of the aligned population, the random-orientation remainder, and the scattering cross sections. Dichroic extinction, polarized emission, birefringence, and the orientation-resolved cross-section table cover the extinction and emission side of that contract; the scattering side — the fixed-orientation Mueller matrix, the extinction matrix built from forward amplitudes, and their library interface — is now built and verified, and is documented in §8.. What remains is MoCafe’s consumption: the frame rotations, direction sampling, peel-off, and the  $\exp(-K\tau)$  transfer step, none of which is a SEDust responsibility.

## 2. The data source

Everything polarized derives from one file, the orientation-resolved Draine & Hensley (2021) astro dust spheroid table:

```
data/astrodust/q_DH21Ad_P0.20_Fe0.00_1.400.dat.gz
```

(porosity  $P = 0.20$ , iron fraction  $f_{\text{Fe}} = 0$ , axial ratio  $b/a = 1.400$ ). Its layout is not self-describing, so it is recorded here (Table 2).

| Property           | Value                                                                         |
|--------------------|-------------------------------------------------------------------------------|
| header             | 12 lines, skipped                                                             |
| payload order      | ((Q(jw,jr,jori), jw=0,1128), jr=0,168), jori=1,3                              |
| repeated           | three times: $Q_{\text{ext}}$ , then $Q_{\text{abs}}$ , then $Q_{\text{sca}}$ |
| one record on disk | the 169 sizes of one (jw, jori) pair                                          |
| records read       | 3 quantities $\times$ 3 orientations $\times$ 1129                            |
| total values       | $3 \times 3 \times 1129 \times 169 = 1,717,209$                               |
| wavelength axis    | data/astrodust/DH21_wave, 1129 nodes                                          |
| size axis          | data/astrodust/DH21_aeff, 169 nodes                                           |
| axis file format   | two title lines, then the values free-format                                  |

Table 2: Layout of the orientation-resolved  $Q$  table. The wavelength and size axes are *not* parsed out of the header; they come from the companion grid files, which list the nodes the table was computed on.

**Orientation index.** The convention is the trap most likely to bite a future reader, and it is worth stating in full (Table 3); see also §3.6 of the theory report.

| jori | geometry                                                 | meaning                             |
|------|----------------------------------------------------------|-------------------------------------|
| 1    | $\mathbf{k} \parallel \hat{a}$                           | propagation along the symmetry axis |
| 2    | $\mathbf{k} \perp \hat{a}, \mathbf{E} \parallel \hat{a}$ | field along the symmetry axis       |
| 3    | $\mathbf{k} \perp \hat{a}, \mathbf{E} \perp \hat{a}$     | field across the symmetry axis      |

Table 3: Orientation index of the DH21 table.  $\hat{a}$  is the spheroid symmetry axis.

### 3. Derived quantities

For a grain whose symmetry axis lies in the plane of the sky and which is perfectly aligned, the polarization cross section is the difference the grain presents to the two linear polarizations, halved:

$$Q_{\text{pol}} = \frac{1}{2} [Q(\text{jori}=3) - Q(\text{jori}=2)], \quad (1)$$

formed separately from  $Q_{\text{ext}}$  and from  $Q_{\text{abs}}$ , giving  $Q_{\text{pol,ext}}$  and  $Q_{\text{pol,abs}}$ . The random-orientation average of the same table is the trace,

$$Q_{\text{ran}} = \frac{1}{3} [Q_1 + Q_2 + Q_3], \quad (2)$$

which is the modified picket fence approximation (MPFA; §3.2 of the theory report). Cross sections follow from

$$C = Q \pi a_{\text{eff}}^2, \quad (3)$$

with  $a_{\text{eff}}$  converted to cm. The alignment efficiency is the Hensley & Draine (2023) Table 1 fit,

$$f_{\text{align}}(a) = \frac{f_{\text{max}}}{1 + (a_{\text{align}}/a)^\alpha}, \quad a_{\text{align}} = 0.0749 \mu\text{m}, \quad \alpha = 1.80, \quad f_{\text{max}} = 1.00. \quad (4)$$

### 4. Module map

Table 4 lists what each file is responsible for. Nothing else in the tree touches polarization.

### 5. Implementation decisions

This is the part worth reading again in six months.

#### 5.1 The accumulation sites, and the two that hide

`sed_grain_loop` branches on `stoch_method` into four cases, and the polarized channel has to be accumulated wherever the total is. There are eight such statements (Table 5). Missing any one of them produces a polarized spectrum that is silently too low rather than an error.

**The lesson.** The 'draine' and 'qm' branches are OpenMP-parallel and accumulate into a thread-local `Jout_local` that is reduced at the end. The 'heuristic' branch is serial and accumulates into `Jout` directly. A survey of the accumulation sites conducted by searching for `Jout_local` therefore returns only the parallel branches, and an initial

pass over this code identified four sites on exactly that basis. Since 'heuristic' is the *default* solver, that omission would have made the polarized emission come back identically zero in the default configuration while every other solver looked correct. Any future change that adds a term to the emission sum must be checked against Table 5 branch by branch, not by pattern search.

## 5.2 Why the 'qm' rescaling is exact

qm\_solve\_grain does not return a temperature distribution; it returns the emitted spectrum with  $C_{\text{abs}}$  already folded in. The polarized counterpart is therefore obtained by rescaling:

```
if (Cabs_pop(ii, ir) > 0.0_wp) then
  Jpol_local(ii) = Jpol_local(ii) + &
    wpol * emission_qm(ii) * (Cpol_pop(ii, ir) / Cabs_pop(ii, ir)) / &
    (4.0_wp * PI * lam(ii) * 1.0e-3_wp)
end if
```

This is not an approximation. The temperature distribution  $P(T)$  is set by the *heating*, which is governed by  $C_{\text{abs}}$  and is unchanged by the question of which polarization is being read out. Only the emission readout changes, so

$$C_{\text{pol}} \int P(T) B_{\lambda}(T) dT = \frac{C_{\text{pol}}}{C_{\text{abs}}} C_{\text{abs}} \int P(T) B_{\lambda}(T) dT = \frac{C_{\text{pol}}}{C_{\text{abs}}} \text{emission\_qm} \quad (5)$$

holds identically, wavelength by wavelength. Wavelengths with zero  $C_{\text{abs}}$  emit nothing and are skipped rather than divided by zero.

## 5.3 The $P(T)$ solver was not touched

No change was made to calc\_P, narrow\_iterative, or qm\_solve\_grain. This is deliberate and it is what makes stochastic heating automatically correct for polarization: at every site the same temperature weight multiplies both accumulators, one with  $dn C_{\text{abs}}$  and the other with  $dn f_{\text{align}} C_{\text{pol}}$ . The polarized spectrum is thus the same average over the same  $P(T)$ , and no separate stochastic treatment of the polarized channel is needed (see §5.7 of the theory report for why this is the physically right structure).

## 5.4 Mixed orientation-average conventions

This is the one open physical wart, and it is documented rather than hidden.

$C_{\text{abs}}$  in SEDust comes from our own T-matrix run, whose random-orientation average is computed exactly.  $C_{\text{pol}}$  comes from the release table, whose available average is the 1/3 trace, Eq. (2), i.e. the MPFA. The two are different approximations to the same quantity, and the code mixes them: build\_Cpol adds the polarized channel and deliberately does *not* recompute  $C_{\text{abs}}$ .

Table 6 gives the measured disagreement,  $|Q_{\text{abs}}^{\text{trace}}/Q_{\text{abs}}^{\text{exact}} - 1|$ , on the shared  $(\lambda, a_{\text{eff}})$  grid.

**Reading the far-infrared maximum.** The 3.46% worst case occurs at  $a = 4.2\,\mu\text{m}$ , a size carrying  $dn/n_H \sim 2.8 \times 10^{-36}$ , so it is numerically irrelevant. Restricting to  $a \leq 0.5\,\mu\text{m}$ , which covers 99.999% of the geometric cross section, the worst case drops to 1.98%. At  $a \ll \lambda$  both averages approach the same Rayleigh limit, which is why the far infrared is well behaved.

**Verdict: deliberately deferred in this release.** The mixture is left in place for v1.20, as a decision rather than an oversight. The reasoning is recorded here so that it is not rediscovered later. The mixture is harmless for polarized *emission*, which is a far-infrared phenomenon. It is **not** defensible for polarized *extinction* in the ultraviolet or the optical, and column 8 spans the whole wavelength grid. The `Cpol_ext` output of `dust_extinction` is the same quantity by the same route and carries the same caveat unchanged. Anyone using either shortward of the near infrared must resolve this first. The clean resolution is to recompute  $C_{\text{abs}}$  on the MPFA convention so that both channels share one average. That changes the existing SEDs, so it breaks their byte-identity with the current reference outputs and has to be done as a separate, deliberately validated step rather than folded into this one. The size of the change is not an argument for leaving it: the correctness argument is that two channels of one calculation should not use two different orientation averages.

What was decided for this release is only *when*, not *whether*. The planned use of these optics — redoing Seon (2018) in the  $V$  band, with one or two further optical bands possible — lies entirely in the range where the disagreement is 0.080% or below, while unifying the convention would change  $C_{\text{abs}}$  and break the byte-identical regression baseline that the rest of v1.20 was validated against. Spending that baseline buys nothing inside the present scope. It has to be spent before any ultraviolet polarized extinction is quoted, and that is the trigger for doing the work.

## 5.5 $f_{\text{align}}$ is the fit, not the release column

`data/release/size_distribution.dat` also ships an alignment column. It agrees with Eq. (4) to a median ratio of 1.0000 but differs by up to 0.32% at the small end. The analytic form is the canonical one in SEDust because it is the published fit, because it is defined on any radius grid rather than only the tabulated one, and because it keeps the alignment model in a single place. Do not mix the two sources: a quantity computed half from one and half from the other is neither.

## 5.6 Names and surface area

**Cpol was not renamed to Cpol\_abs.** The name appears at roughly 40 sites and is part of the public surface (`grain_pop_t%Cpol`, the optional `set_pop(..., Cpol_in=...)`). A rename would touch every one of them for no functional gain. The declaration carries an explicit note instead: `Cpol` is the absorption one, `Cpol_ext` its extinction twin.

**The polarized extinction was exported as a table column first.** A dedicated `dust_cpol_ext` accessor was considered at the time and rejected: `dust_lib` was then an emission-facing API with no extinction surface at all, so such an accessor would have been the module's only extinction entry point, a worse inconsistency than the inconvenience

it removed. The condition attached to that decision was that Cpol\_ext should join a full extinction API if one ever appeared. It has; see §5.7.

## 5.7 Why the accessor exists now

The table-only export was right for a library that emitted and did not extinguish. Rebuilding a radiative transfer code on top of SEDust changed the premise: the host would have obtained its emission from a call and its opacity from a parsed text file, and that asymmetry is the whole problem. The file is also written on whatever wavelength grid the driver used, so a host on any other grid has to interpolate optics that the model could have evaluated exactly.

dust\_extinction closes it. The model object already held everything required, so the accessor adds no physics and no new data source; it performs the same size integral the driver performs, over the populations already built. Two fields were added to grain\_pop\_t to make that possible: Cpol\_ext, previously a module array only, and gsca, the scattering asymmetry, which the emission path never needed and which sed\_init now interpolates onto the size grid beside  $C_{\text{abs}}$  and  $C_{\text{sca}}$ , from the same  $Q$  table and by the same routine the driver used. Agreement with the table is at the precision the file is written with (a median relative deviation of  $\sim 7 \times 10^{-9}$  in each cross section), which is the expected result for one integral evaluated twice.

The table remains, unchanged byte for byte, and remains the right route for a tool that wants the optics without linking the library.

**Later revision: the scalar optics went back to the table.** The argument above is about the *interface*, and it stands: a host must get its opacity from a call on the same model object that gives it emission, on its own wavelength grid. What changed afterwards is where that call gets its numbers. The scalar transport optics of a dust model do not depend on the transport, so recomputing the size integral inside an RT run buys nothing that reading the recorded integral does not. dust\_extinction therefore now serves  $C_{\text{abs}}$ ,  $C_{\text{sca}}$ ,  $C_{\text{ext}}$  and  $\langle \cos \theta \rangle$  from a data/<model>/kext\_\*.dat table, attached to the model by the builder (optional kext\_path; the defaults are the EUV products, the widest grid each model has) and interpolated onto  $m\lambda_{\text{m}}$ , exactly at the nodes the two grids share. The host's call site is unchanged, and so is the requirement that the file be opened while the working directory is still sed/ — which is why the read happens in build\_\* and not in dust\_extinction.

The polarized outputs were deliberately *not* moved. Cpol\_ext and Cbir\_ext carry the  $f_{\text{align}}$  weight, which a host resets cell by cell through dust\_set\_alignment; a column fixed at build time could not follow it. So the routine is now mixed on purpose: scalar optics from a table, polarized optics from the model. The size integral itself is unchanged and is still reachable, for all six outputs, as size\_integrated\_extinction — which is what calc\_kext.x calls to write the tables, and what the in-tree checks call when they must compare one build of a model against another.

## 5.8 Compression handling, and a scratch file that leaked

The table ships compressed and Fortran cannot read a deflate stream. Adding a zlib binding was rejected as too heavy for one file. The reader instead shells out once,

```
call execute_command_line('gzip -dc "'//trim(gz_file)//'" > "'// &
```

```
trim(out_file)/*', exitstat=estat, cmdstat=cstat)
```

decompresses to a scratch copy, reads it, and deletes it. A fresh clone works with no manual setup step, and neither a 17MB untracked sibling in `data/dielectric` nor a new build dependency is created. A path not ending in `.gz` is opened directly.

**The scratch target: TMPDIR, with the process id.** The scratch copy is written under TMPDIR (else `/tmp`), named with the process id: `q_jori_<pid>.dat` for the polarized  $Q$  reader (`q_table_jori`) and `scatmat_aligned_<pid>.dat` for the aligned scattering reader (`scatmat_aligned_mod`); both loaders share the scheme. Naming by process id means concurrent MPI ranks never write the same scratch file, and placing it in TMPDIR rather than the launch directory means a read-only launch directory is never written to.

**The bug this had.** An earlier version wrote a fixed `q_jori_scratch.dat` in the *working* directory. Shell redirection creates the target file *before* `gzip` runs, so when decompression failed the early return left a zero-length `q_jori_scratch.dat` behind there. The failure path now calls `discard_scratch_copy` before bailing out, like every other error exit — and, since the scratch file moved to `TMPDIR/q_jori_<pid>.dat`, a leaked remnant no longer lands in the user's directory at all. The general shape of the bug, an error path that skips the removal the success path performs, is worth watching for in any future edit to this reader: it has eight separate early returns.

## 6. Verification

### 6.1 Polarized extinction against the release

The comparison against the release is restricted to  $\lambda \geq 0.11 \mu\text{m}$  because the reference file turns negative below that. The 0.9422% maximum deviation is not a systematic offset: it sits at  $\lambda = 0.112 \mu\text{m}$ , right at the sign reversal of the polarized extinction, where the quantity passes through zero and a relative deviation is not a meaningful measure. Away from that point the agreement is at the level of the median.

The  $7 \times 10^{-8}$  figure is the precision the ASCII output format carries, so column 8 and `calc_pol_ext.x` agree as closely as the files allow. They are two independent routes to the same quantity: `calc_pol_ext.x` works straight from `qpol_ext` and the release size distribution, while column 8 goes through `sed_init`, the  $\log a$  interpolation onto the SED size grid, and `build_Cpol`.

### 6.2 Polarized emission

Intrinsic polarization fraction `lamI_pol/lamI_total` under a Mathis ISRF at  $U = 1.585$ :

The mid-infrared zero is correct rather than a failure: there the emission is dominated by stochastically heated PAHs, which HD23 take to be unaligned.

**19.15% versus Planck's 22% is a property of the model, not a bug.** Planck measures  $p_{\text{max}} = 22.0(+3.5, -1.4)\%$  at 353 GHz. The 19.15% here is the *maximum* attainable value, obtained before any geometric factor, and every real line



of sight reduces it, so the model as implemented cannot reach the observed maximum. This traces to the dust model itself: the axial ratio  $b/a = 1.4$  is the minimum non-sphericity DH21b found the starlight polarization efficiency to require, and it was not tuned to maximize the submillimeter polarization fraction. Do not close the gap by adjusting  $f_{\text{align}}$  or the axial ratio without recognizing that as a change to the dust model, requiring its own revalidation against the extinction and emission constraints HD23 fit.

## 7. First-principles orientation-resolved optics

The table of §2. is Draine & Hensley’s precomputed one. SEDust can now regenerate the same orientation-resolved cross sections from the astrodust dielectric function with its own T-matrix engine, so the polarized optics no longer have to be taken on trust from the release file. The random-orientation optics were already SEDust’s own (the `q_astrodust` table read by `q_table`); this closes the remaining gap, the orientation-resolved part read by `q_table_jori`.

Recall the three orientations of the incident wave relative to the spheroid symmetry axis  $\hat{a}$ : `jori = 1` is  $\mathbf{k} \parallel \hat{a}$ ; `jori = 2` is  $\mathbf{k} \perp \hat{a}$  with  $\mathbf{E} \parallel \hat{a}$ ; `jori = 3` is  $\mathbf{k} \perp \hat{a}$  with  $\mathbf{E} \perp \hat{a}$ . The polarized and mean combinations are  $Q_{\text{pol}} = \frac{1}{2}(Q_3 - Q_2)$  and  $Q_{\text{ran}} = \frac{1}{3}(Q_1 + Q_2 + Q_3)$ .

The optics are reached two ways. The `table` is loaded once and interpolated: this is the production path, fast and safe to call from many threads. The `direct` routine `oriented_cross_sections` computes one node from first principles; it is the shared core the table generator loops over, and it also serves an arbitrary  $(a, \lambda)$  or a shape or porosity for which no table exists. It must not be called from inside a loop over radiative-transfer cells: each call is a full T-matrix solve plus an angular integral, and it works through an initialized T-matrix workspace of 80 MiB. Thread safety is no longer the objection — the library holds every working array in a caller-owned `tmatrix_workspace_t`, with no COMMON, no EQUIVALENCE and no mutable SAVE (§8.3) — the cost is.

### 7.1 Rayleigh regime, $x < 0.1$

In the electrostatic limit the oblate spheroid has a diagonal polarizability  $\alpha = \text{diag}(\alpha_a, \alpha_b, \alpha_b)$ , with  $\alpha_a$  for  $\mathbf{E}$  along the symmetry axis and  $\alpha_b$  for  $\mathbf{E}$  transverse, set by the depolarization factors  $L_a, L_b$  of the spheroid through  $\alpha_i \propto (\epsilon - 1)/[1 + (\epsilon - 1)L_i]$ . Absorption and scattering follow as  $Q_{\text{abs}} \propto \text{Im } \alpha$  and  $Q_{\text{sca}} \propto |\alpha|^2$ . The three orientations sample exactly these two polarizabilities: `jori = 1` and `3` present the transverse  $\alpha_b$  (the field is transverse to  $\hat{a}$  in both), `jori = 2` presents  $\alpha_a$ . So the orientation-resolved cross sections are the two polarization components the random average  $(\alpha_a + 2\alpha_b)/3$  is already built from; extracting them before the average is analytic and exact, and  $Q(\text{jori} = 1)$  equals  $Q(\text{jori} = 3)$  identically. `rayleigh_limit` returns them through optional arguments, leaving its averaged outputs untouched.

### 7.2 T-matrix regime, $0.1 \leq x \leq 50$

The T-matrix  $\mathbf{T}$  is a property of the particle alone — size, shape, refractive index — independent of orientation and of the incidence and scattering directions. SEDust solves it once for each wavelength and radius with `tmatrix_eval_scattering_matrix`, the library entry point that descends from Mishchenko’s random-orientation solver; the converged  $\mathbf{T}$  stays in the caller’s workspace and the returned `tmatrix_scattermat_t` identifies it. Two

orientation-resolved quantities are then read off the same  $\mathbf{T}$ .

**Extinction, from the optical theorem.** Mishchenko's AMPL evaluates the fixed-orientation amplitude matrix  $\mathbf{S}(\mathbf{k} \rightarrow \mathbf{k}')$  from the stored  $\mathbf{T}$  for a particle placed at chosen Euler angles. Taking the scattering direction equal to the incidence direction (forward) and applying the optical theorem,  $C_{\text{ext}} = (4\pi/k) \text{Im} S_{\text{fwd}}$  for the incident polarization, gives  $C_{\text{ext}}$  for each orientation: the symmetry axis is placed along  $\mathbf{k}$  (jori = 1) or across it (jori = 2, 3) and the co-polar forward amplitude is read ( $S_{\theta\theta}$  for  $\mathbf{E}$  along  $\hat{\theta}$ ,  $S_{\phi\phi}$  for  $\hat{\phi}$ ).

**Scattering, from the phase-matrix integral.** The same AMPL, swept over all scattering directions at fixed incidence, gives the differential scattering cross section  $|\mathbf{S}|^2$ ; the scattering cross section is its integral over the sphere,  $C_{\text{sca}} = \int (|S_{\theta \leftarrow p}|^2 + |S_{\phi \leftarrow p}|^2) d\Omega$  for incident polarization  $p$ , and absorption follows by conservation,  $C_{\text{abs}} = C_{\text{ext}} - C_{\text{sca}}$ . The fixed-orientation scattered intensity is a trigonometric polynomial of degree  $\sim 2N_{\text{max}}$  in  $\cos\theta$  and in  $\phi$ , so Gauss–Legendre nodes in  $\cos\theta$  ( $N_{\text{max}} + 2$  of them) and a uniform azimuth grid ( $2N_{\text{max}} + 2$  points) integrate it to machine precision.

**An engine subtlety, and how it was retired.** In Mishchenko's f77 sources the random-orientation expansion GSP reuses the /TMAT/ storage as scratch through an EQUIVALENCE and so destroys the T-matrix it reads. The amplitude evaluation needs that same T-matrix afterward, so the f77 driver kept an intact copy in a second common block and restored /TMAT/ after GSP. Without that restore the oriented cross sections came out as garbage (negative  $C_{\text{ext}}$ ), which corrected a premise stated in the planning note, that the converged T-matrix simply remains available after the solve. The library removes the overlay rather than repairing it: the workspace carries the T-matrix (full\_tstore) and the expansion scratch (full\_gsp\_\*) as separate components, the expansion routine (GSP\_FULL\_DIRECT) only reads the former, and it therefore survives the solve intact with nothing to restore. The planning note's premise holds again, now by construction; the episode is kept on the record because the failure it caused was silent.

**Reusing a workspace is checked, not trusted.** A workspace holds one T-matrix at a time, so an amplitude evaluation made against a scatmat older than the T-matrix now in the workspace would silently answer for the wrong particle. The workspace carries a generation counter tm\_tag that advances on every evaluation, successful or failed, and tmatrix\_eval\_scattering\_matrix copies it into scatmat%state\_tag; AMPL compares the two and refuses with IERR = 4 when they differ. The remaining codes are 0 success, 1 an angular argument outside its range, 2 no converged T-matrix in the workspace, 3 a truncation order outside the stored block. Nothing here is a convenience: each code replaces a case that used to be undefined behavior or a STOP inside a library.

### 7.3 Geometric-optics regime, $x > 50$

Here the T-matrix does not converge and an asymptotic model is used. Extinction is the extinction paradox,  $C_{\text{ext}} \rightarrow 2A_{\text{proj}}$ , twice the geometric shadow. For an oblate spheroid of equatorial semi-axis  $a_s$  and symmetry semi-axis  $c$  ( $a_s/c = b/a$ ), the shadow is  $\pi a_s^2$  when  $\mathbf{k}$  is along the symmetry axis (jori = 1) and  $\pi a_s c$  when it is across (jori = 2, 3), giving

$Q_{\text{ext}}(1) = 2(b/a)^{2/3}$  and  $Q_{\text{ext}}(2) = Q_{\text{ext}}(3) = 2(b/a)^{-1/3}$ . The shadow depends on the direction of  $\mathbf{k}$  relative to the axis but not on the incident polarization, so extinction carries no polarization at this order.

The polarization in this regime is carried by absorption. In the opaque limit, where a refracted ray is fully absorbed — valid when  $\text{Im}(m)x \gg 1$  — the absorptivity of an illuminated surface element is one minus its Fresnel power reflectance, and that reflectance is polarization-dependent ( $R_s \neq R_p$  away from normal incidence). Integrating over the illuminated half of the spheroid,

$$C_{\text{abs},p} = \int_{\text{illum}} [1 - (|\hat{e} \cdot \hat{s}|^2 R_s(\theta_i) + |\hat{e} \cdot \hat{p}|^2 R_p(\theta_i))] dA_{\text{proj}},$$

where  $\theta_i$  is the local angle of incidence and  $\hat{s}, \hat{p}$  the local senkrecht and parallel directions. The two grain-frame polarizations ( $\hat{e} = \hat{a}$  versus  $\hat{e} \perp \hat{a}$ ) project differently onto the local  $\hat{s}, \hat{p}$  basis as they sweep the curved surface, and that is the whole source of the dichroism. A purely geometric ray treatment, with no Fresnel coefficients, would give no polarization at all; the surface reflectance is what makes the regime dichroic, and it is needed to match the release, which carries a real absorption dichroism there — about 3% of the mean, and *growing* with size parameter, the wave-optics signature a ray limit cannot reproduce. This is a deliberate approximation, documented at the code site with its opaque-limit domain of validity; it would matter for a size distribution with large-grain weight, which the astrodust distribution does not have.

## 7.4 Dispatch, generation, and drop-in

`oriented_cross_sections` selects the regime by  $x$  and returns the oriented ( $Q_{\text{ext}}, Q_{\text{abs}}, Q_{\text{sca}}$ ) for one node, with the same branch boundaries and fallbacks as the random-orientation driver. `run_q_jori.x` loops it over the DH21 grid and writes the table in the exact stream the release reader expects (§2.), so the result is a drop-in: `sed_init` and `build_astrodust` already take an optional `qpol_path` that defaults to the release file, and pointing it at the regenerated table switches the polarized optics to SEDust's own with no code change. A full sweep is dominated by the T-matrix nodes (of order 16 hours on one core), so the driver has a wavelength-window range mode for process-level parallelism and a merge mode that reassembles the windows — a window is not a contiguous slice of the stream, because the wavelength index runs innermost.

## 7.5 How far it has been checked

Each regime was compared node by node against the release table (Table 9). Across the Rayleigh and T-matrix regimes — every wavelength and radius that carries astrodust weight — the reconstruction matches the release to a median of a few parts in  $10^4$ , in both the mean and the polarized cross sections. The geometric-optics region agrees only coarsely, as an asymptotic model should: the mean absorption to about ten percent and the dichroism to about three-quarters of the release value, both with the correct sign and size-parameter trend. It joins the T-matrix continuously across  $x = 50$  in sign and ordering, and it carries negligible weight in the size integral.

End to end, feeding the regenerated table through `calc_pol_ext` reproduces the released polarized extinction to a median of 0.06% over  $\lambda \geq 0.11 \mu\text{m}$ , against 0.03% for the release table itself; the two agree with each other to a median of 0.02% at the level of the observable. So the few-parts-in- $10^4$  reconstruction error at the level of a single

cross section survives the size integral as a few parts in  $10^4$  in the polarized extinction, computed entirely from the astroduct dielectric function with no recourse to the release optics.

## 7.6 Computing named wavelengths

Polarized transfer is run at a handful of wavelengths, not on a whole axis. The range mode cannot serve that: its JW1/JW2 are *indices* into a fixed wavelength file, so it reaches only wavelengths that file already carries. Three modes take the physical wavelength instead:

```
./run_q_jori.x lam L1 [L2 ...] [ja=JA1:JA2] [tag=NAME] # wavelengths [um]
./run_q_jori.x lamfile PATH [ja=JA1:JA2] [tag=NAME] # one per line, '#' comments
./run_q_jori.x lammerge STEM FILE [FILE ...] # reassemble ja= windows
```

Each writes a pair, q\_astroduct\_jori\_P0.20\_Fe0.00\_1.400.TAG.dat and its wavelength axis ...TAG.wave; the default TAG is lamL for one wavelength and lamLMIN\_LMAX\_nN for several. That pair is what sed\_init / build\_astroduct take through qpol\_euv\_path and qpol\_euv\_wave\_path. The reader interpolates in  $\log \lambda$ , so at least *two* wavelengths are needed to fill a band; a one-wavelength file is a valid product to read directly but cannot be interpolated.

**Processes, not threads.** ja=JA1:JA2 restricts one invocation to radii JA1..JA2 of the 169-point size axis so that several *processes* share one wavelength, and lammerge puts the column slices back side by side. This was once forced by the engine: the f77 solver handed the converged T-matrix to AMPL through COMMON /TMAT/ and kept its working storage in further COMMON blocks, two of them blank, so the core was not re-entrant. That constraint is gone — the library owns no shared state and distinct workspaces may be evaluated concurrently (§8.3) — but the driver keeps the process windows, which are already measured and cost nothing to run. Separate processes each get their own copy, and a merged file has been checked byte for byte against a single-process one. Measured: three wavelengths (0.0124, 0.0602 and 0.0912  $\mu\text{m}$ ) over 57 processes of three radii each took 9 m 22 s of wall time for 34 m 48 s of processor time. The speedup is 3.7 rather than 57 because the cost is not spread over the radii — a T-matrix solve grows like  $N_{\text{max}}^4$  with  $N_{\text{max}} \sim x$ , so the few radii just below the  $x = 60$  ceiling carry nearly all of the work. One radius per invocation (169 processes) lets the operating system interleave them and brings the wall time down to the single most expensive node.

**euv is an axis switch, not a wavelength.** The leading euv keyword selects a wavelength *axis file* (data/astroduct/DH21\_wave\_euv), so it cannot be combined with lam, lamfile or lammerge, which carry their own axis. The driver rejects the combination explicitly.

## 7.7 The polarized band below the Lyman limit

$C_{\text{pol}}$ ,  $C_{\text{pol,ext}}$  and  $C_{\text{bir,ext}}$  are read from an orientation-resolved table on the full DH21 wavelength axis, all 1129 nodes from 0.0912 to 39810  $\mu\text{m}$ : by default the release file data/astroduct/q\_DH21Ad\_P0.20\_Fe0.00\_1.400.dat.gz, or, through qpol\_path, the regenerated tmatrix/output/q\_astroduct\_jori\_P0.20\_Fe0.00\_1.400.dat.gz, which adds the 4th block and so is the one that makes  $C_{\text{bir,ext}}$  nonzero. Either way the wavelength coverage is the whole axis.

Do not confuse this with the five-band aligned scattering table of §8. — the polarized *extinction* optics are on the full grid.

When a model grid is extended shortward of  $0.0912\mu\text{m}$  (through `lam_min`), `build_Cpol` looks for the companion table `q_astrodust_jori_euv_P0.20_Fe0.00_1.400.dat.gz`. **That table is not shipped.** The default state is therefore that the extreme-ultraviolet block stays at the zero it was initialized with, and the reader says so on `stderr`:

```
build_Cpol: no EUV polarized table (...)
           dichroic extinction and birefringence are zero below 9.120E-02 um
           (zero by omission, not by physics).
```

**That zero is a documented deficit, not the right value.** It is worth being blunt about this, because a zero is easy to mistake for an asymptote. First-principles size integrals (`tmatrix/driver/euv_polarized_optics.f90`) put the alignment-weighted  $|C_{\text{pol,ext}}|/C_{\text{ext}}$  at  $1.26 \times 10^{-3}$  at the  $0.0912\mu\text{m}$  seam, *rising* to  $3.64 \times 10^{-3}$  near  $20.6\text{ eV}$  — a factor 2.9 — before decaying to  $1.6 \times 10^{-4}$  at  $100\text{ eV}$ , with the sign *opposite* to the optical band; the reversal falls near  $0.106\mu\text{m}$ , between  $0.1072$  and  $0.1059\mu\text{m}$ .  $|C_{\text{bir,ext}}|/C_{\text{ext}}$  falls monotonically from  $7.2 \times 10^{-3}$  at the seam and changes sign near  $51\text{ eV}$ . Zero would be the correct limit only for  $x \gg 1$  and  $|m-1| \ll 1$ , and  $|m-1|$  is  $0.765$  at the seam and below  $0.1$  only above  $\sim 50\text{ eV}$  while the grains carrying the signal sit at  $x = 5\text{--}50$ . A sphere has *exactly* zero dichroic extinction and exactly zero birefringence, so the scalar extreme-ultraviolet optics (`q_astrodust_mod`, Mie on the volume-equivalent sphere) could not have supplied these entries either — the shape has to be kept to get them at all.

**When the companion table is present.** `build_Cpol` reads it and interpolates in  $\log \lambda$  and  $\log a$ , as  $C_{\text{abs}}$  and  $C_{\text{sca}}$  are interpolated. A grid running outside the wavelengths the table covers ( $0.0124\text{--}0.0912\mu\text{m}$ ) is rejected with `status = 9` rather than filled by extrapolation. Generate the table with the `euv` mode of `run_q_jori.x`, or a subset of it with the named-wavelength modes of §7.6.

**Why the axis stops at  $0.0124\mu\text{m}$  (100 eV).** Above  $x = 60$  the only description left is the opaque geometric-optics limit, whose absorption is the Fresnel surface integral described in §7. and which needs the chord optical depth  $4\text{Im}(m)x$  to be large. For DH21 astrodust that quantity is  $3.7$  at  $x = 50$  and  $100\text{ eV}$  (2.6% transmission), and it falls below  $1$  by  $200\text{ eV}$ , where the grain is becoming X-ray transparent and neither the extinction paradox nor the opaque Fresnel absorption applies. The floor is set where the physics of every branch used is still valid.

## 7.8 Certifying the T-matrix above $x = 50$

The extreme ultraviolet pushes size parameters up, so the  $x > 50$  cut of §7. would throw away most of the band. In the `euv`, `lam` and `lamfile` modes the T-matrix is therefore attempted up to  $x = 60$  (`X_TM_MAX`), under a two-setting certification: every node with  $x \in (50, 60]$  is solved twice, at the production (`DDELT, NDGS`) =  $(10^{-3}, 2)$  and at a coarser-quadrature  $(3 \times 10^{-3}, 3)$ , and is *kept only* if the two agree to `TOL_CHK` = 5% on every channel. A node that fails falls to geometric optics. Above  $x = 60$  geometric optics is used directly. The plain and range sweeps keep the older  $x > 50$  rule unchanged, so they still reproduce the shipped table byte for byte.

Three details are worth recording.

- **NDGS = 4 is not usable** as the second setting:  $\text{NGAUSS} = \text{NDGS} \times N_{\text{max}}$  then exceeds the  $\text{NPNG1} = 300$  storage limit, so the check would fail for want of array space rather than for want of convergence.  $\text{NDGS} = 3$  is the largest that leaves the comparison meaningful.
- **IERR = 0 is not sufficient at  $x \gtrsim 55$ .** Mishchenko's own convergence flag can report success and still return a wrong answer: at  $\lambda = 0.0912 \mu\text{m}$  and  $a = 0.9441 \mu\text{m}$  ( $x = 65.04$ ) it returns  $\text{IERR} = 0$  with a value 50% away from its neighbors. Measured at the seam,  $x = 51.7$  and  $54.7$  pass the two-setting check while  $57.9, 61.4, 65.0$  and  $68.9$  fail it. The certification exists because the engine's own flag does not carry this information.
- **The polarized channels need an absolute floor.**  $Q_{\text{pol}}$  is a difference of order-unity numbers, so a purely relative comparison becomes meaningless where it passes through zero.  $Q_{\text{POL\_FLOOR}} = 10^{-3}$  sets the absolute scale below which the polarized channels are not required to agree; after the size integral that corresponds to  $|C_{\text{pol}}|/C_{\text{ext}} \sim 5 \times 10^{-4}$ , about 1% of the  $V$ -band dichroism.

**What the certification costs.** Leaving the  $x > 60$  nodes at the geometric-optics value, which carries  $Q_{\text{pol,ext}} = Q_{\text{bir,ext}} = 0$  exactly, loses part of the *dichroic extinction*: nothing at the seam,  $\sim 8\%$  of the band total at  $0.031 \mu\text{m}$ ,  $\sim 20\%$  at  $0.022 \mu\text{m}$  and  $\sim 46\%$  at  $0.0124 \mu\text{m}$ . The true value lies between the certified figure (a lower bound,  $1.06 \times 10^{-4}$ ) and a  $1/x$  continuation (an upper bound,  $3.14 \times 10^{-4}$ ). No extrapolation is put in place of the missing part; the absolute level is already small there, about 1% of the optical-band signal.  $C_{\text{pol}}$  — the absorption dichroism, which drives polarized emission — has no such gap, because the geometric-optics limit obtains it from the opaque-grain Fresnel surface integral.

## 8. Aligned-grain scattering matrix and the extinction matrix

Section 7. regenerated the orientation-resolved *cross sections*. This section adds the angle-resolved object they left out: the fixed-orientation Mueller matrix  $Z(\theta_i; \theta_s, \phi)$  of the aligned astrodust spheroid, and the  $4 \times 4$  extinction matrix  $K(\theta_i)$  that goes with it. Between them they complete the material side of the Peest-formalism contract — SEDust returns every quantity in matrix form; the RT code (MoCafe, whose consumption side is not yet built) does the rotations, sampling, peel-off, and  $\exp(-K\tau)$  transfer. The design is set out in docs/tmatrix\_aligned\_scattering\_plan.md; this section records how each piece was built and why.

### 8.1 The engine

tmatrix/driver/scattering\_matrix\_oriented.f90 returns the fixed-orientation Mueller matrix  $Z$  of the DH21 oblate spheroid ( $b/a = 1.4$ ) for one incidence direction. Two regimes, one output convention.

**T-matrix regime, one AMPL call.** With the T-matrix already converged and held in the workspace by `tmatrix_eval_scattering_matrix` — the storage separation of §7. that made the oriented cross sections possible — a single call to Mishchenko's AMPL evaluates the  $2 \times 2$  complex amplitude matrix  $\mathbf{S} = [[S_{vv}, S_{vh}], [S_{hv}, S_{hh}]]$  for

the chosen incidence and scattering angles; the multipole truncation order and the wavelength come from the returned scatmat and from the workspace, never from the caller. The 16 phase-matrix elements then follow from the standard bilinears in  $\mathbf{S}$ , copied *verbatim* from Mishchenko's reference/upstream/ampld.lp.f (its Z11...Z44, Eqs. (13)–(29) of Appl. Opt. **39**, 1026, 2000) so the phase matrix can be checked line by line against the source — which is why that unmodified file is kept in the tree even though nothing compiles it. Copying rather than rederiving is deliberate: the bilinears carry sign and conjugation conventions that are easy to transcribe wrongly, and the whole point of anchor C below is that this transcription is exact.

**Rayleigh regime, analytic dipole.** Below  $x = 0.1$  the same  $Z$  is the electric-dipole matrix,  $S_{pq} = k^2 (\hat{e}_p^{\text{sca}} \cdot \alpha \cdot \hat{e}_q^{\text{inc}})$  with the spheroid polarizability  $\alpha = \text{diag}(\alpha_b, \alpha_b, \alpha_a)$ , then the same bilinears. The polarizability comes from spheroid\_dipole\_polarizability, a routine extracted from rayleigh\_limit (§7.) without changing an operation, so the dipole cross sections that §7. already validated and the dipole matrix here rest on one implementation of  $\alpha_a, \alpha_b$ .

**The Stokes basis, stated once.** Both regimes use Mishchenko's meridional basis of each propagation direction *in the grain frame*, whose  $z$  axis is the alignment axis:  $\hat{e}_1 = \hat{\theta}$  (call it  $v$ ),  $\hat{e}_2 = \hat{\phi}$  ( $h$ ), with  $Q = I_v - I_h$ . This is exactly the basis AMPL rotates its amplitudes into, so the analytic dipole and the T-matrix matrix share one convention with no adaptor between them. At  $\theta_i = 90^\circ$  the  $v$ -polarization is jori = 2 ( $\mathbf{E} \parallel \hat{a}$ ) and  $h$  is jori = 3 ( $\mathbf{E} \perp \hat{a}$ ), which is what ties the extinction matrix back to the jori cross sections of §2..

## 8.2 Symmetries, verified not assumed

Three exact symmetries reduce the stored ranges and, more usefully, are numerical checks on the engine (anchor E). Each was confirmed on a grid of  $(\theta_i, \theta_s, \phi)$  before being relied on.

- **Azimuth mirror.** The plane through the axis and the incidence direction is a symmetry plane, so  $\phi \rightarrow 360^\circ - \phi$  leaves  $Z$  unchanged except that the two off-diagonal  $2 \times 2$  blocks (elements 13, 14, 23, 24, 31, 32, 41, 42) flip sign. Store  $\phi \in [0, 180^\circ]$ ; residual  $3.5 \times 10^{-7}$ .
- **Equatorial mapping.** The oblate spheroid is symmetric under  $z \rightarrow -z$ , so  $\theta_i \rightarrow 180^\circ - \theta_i$  maps onto the stored range with  $\theta_s \rightarrow 180^\circ - \theta_s$ ,  $\phi$  *unchanged*, and the same off-diagonal-block sign flip,

$$Z(180^\circ - \theta_i; 180^\circ - \theta_s, \phi) = S \odot Z(\theta_i; \theta_s, \phi), \quad S = \begin{pmatrix} + & + & - & - \\ + & + & - & - \\ - & - & + & + \\ - & - & + & + \end{pmatrix},$$

the diagonal blocks keeping their sign. Store  $\theta_i \in [0, 90^\circ]$ ; residual  $8.8 \times 10^{-16}$ . Physically this also encodes that the alignment axis is headless.

- **Axisymmetry at  $\theta_i = 0$ .** There  $Z$  must be independent of  $\phi$  up to a frame rotation,  $Z(0; \theta_s, \phi) = Z(0; \theta_s, 0)R(\phi)$ , the  $\phi = 0$  matrix being six-element block-diagonal. This is a check, not a storage saving.

The library reconstructs the unstored octants from these relations, applying the  $\pm 1$  sign once to the interpolated matrix rather than to eight interpolation corners (§8.5); because each reflection contributes a constant that is the same for every corner of a cell, it factors out of the linear interpolation exactly.

### 8.3 The generator: one code path

The shared physics lives in `tmatrix/driver/aligned_population_optics.f90`, and both the production driver `run_scattermat_aligned.f90` and the verification program `compare_scattermat_aligned.f90` call it. This is on purpose: the numbers an anchor checks are produced by the exact code the production file is written from, so a passing anchor certifies the shipped path and not a parallel one.

**One solve, three products.** For each (band, size) in the T-matrix regime a single `tmatrix_eval_scattering_matrix` solve serves everything: the expansion coefficients it returns in the `scattermat` feed the random-orientation accumulation (the  $F$  block below), the T-matrix it leaves in the workspace feeds the AMPL sweep over the  $Z$  nodes, and its forward amplitudes give the extinction-matrix elements. No quantity is recomputed.

**Regime split and the  $x > 50$  guard.** The size-parameter regimes match `run_scattermat.f90`:  $x < 0.1$  analytic dipole,  $0.1 \leq x \leq 50$  T-matrix. Above  $x = 50$  no converged oriented matrix exists, so the bin is *omitted from the aligned products* and instead enters  $F_{\text{tot}}$  as its random-orientation geometric-optics matrix, i.e. scattered as unaligned. Its  $C_{\text{sca}}$  is tallied, and the driver stops hard if the omitted fraction of the scattering weight exceeds  $10^{-6}$ . It never does for the shipped bands: the measured omitted fraction is at most  $3.1 \times 10^{-15}$  (at  $0.36 \mu\text{m}$ ) and exactly zero at  $0.64$  and  $0.79 \mu\text{m}$ . The omission is thus provably negligible rather than assumed so.

**The threading contract: one active caller for each workspace.** The only parallel region is the loop over  $Z$  nodes inside one size. AMPL and VIGAMPL take the workspace `intent(in)`: they read the stored T-matrix and otherwise write local automatic arrays, so distinct nodes are independent and each grid slot is written once — which is why a 1-thread and an  $N$ -thread run are bitwise identical. What this rests on is no longer a compile flag but the shape of the library. There is no COMMON, no EQUIVALENCE and no mutable SAVE anywhere in it, and every working array lives in the caller-owned `tmatrix_workspace_t`, so the contract fits in one line: *one active caller for each workspace, and distinct workspaces may be used concurrently*. Reading is the exemption this loop depends on — an amplitude evaluation only reads, so several threads may evaluate against a single workspace at once, and that is exactly what happens here: the T-matrix is solved once, serially, and only the scattering directions are swept in parallel. Giving each thread a workspace of its own would not be the safer choice but the wrong one, a fresh workspace holding no T-matrix and the evaluation returning `IERR = 2`.

The whole library is compiled `-frecursive`, so large locals stay on the caller's stack. One file, `tmatrix_core.f90`, additionally carries `-fno-inline`, and that flag has nothing to do with threads: inlining moves where the compiler contracts multiply-add pairs, which moves the last bits of the T-matrix convergence test, and without it 2 of the 6258 rows of the shipped astrodust table shift by  $\sim 10^{-6}$ . Lifting the largest scratch arrays out of the call chain and into the



workspace also brought the stack requirement of the deepest path down from about 3.1 MB to about 190 kB; what now scales with the thread count is the workspace itself, 83 926 096 bytes (80.04 MiB) each.

The lesson worth carrying forward is the reverse of the one this paragraph used to record. Deciding thread-safety file by file through compile flags holds only as long as someone remembers which file is which. Removing the shared state and handing the storage to the caller makes the contract a property of the source instead, and one sentence long.

## 8.4 Products and the $\eta$ contract

Per band the file (output/scatmat\_aligned\_astrodust\_P0.20\_Fe0.00\_1.400.dat) carries three blocks, all per H.

- **K block**, on the  $\theta_i$  grid:  $C_{\text{ext,al}}$ , the dichroic  $C_{\text{pol,al}}$ , and the birefringent  $C_{\text{bir,al}}$  from the forward amplitudes, plus  $C_{\text{sca,al}}$  as the grid closure  $\int Z_{11} d\Omega$ . Because these come from the forward amplitude at each  $\theta_i$ , the extinction matrix is *exact at every propagation angle*. This replaces the  $\sin^2 \psi$  interpolation between the  $\psi = 0$  and  $90^\circ$  values that the three-orientation jori table forces, and reduces to that law analytically in the Rayleigh limit (anchor G).
- **F block**:  $F_{\text{tot}}$  (every grain, identical to run\_scatmat.f90) and  $F_{\text{ref}}$  ( $f_{\text{align}}$ -weighted), the six-element random-orientation matrices,  $\alpha_1$ -normalized so  $\frac{1}{2} \int F_{11} d\cos\Theta = 1$  (equivalently  $\int F_{11} d\Omega = 4\pi$ ); the header gives  $C_{\text{sca,tot}}$ ,  $C_{\text{sca,ref}}$ ,  $C_{\text{ext,tot}}$ ,  $C_{\text{ext,ref}}$ , from which the absolute differential matrix is  $C_{\text{sca}} F / (4\pi)$  ( $F_{\text{tot}}$  with  $C_{\text{sca,tot}}$ ,  $F_{\text{ref}}$  with  $C_{\text{sca,ref}}$ ).
- **Z block**:  $Z_{\text{al}} = \int da n(a) f(a) Z(a)$ , 16 elements,  $\mu\text{m}^2 \text{sr}^{-1}$  per H, on  $\theta_i = 0(5)90^\circ (19) \times \theta_s = 0(1)180^\circ (181) \times \phi = 0(5)180^\circ (37)$ .

**The  $\eta$  contract.** The file header documents, and the library enforces, the exact rule for varying the alignment degree cell to cell by a single scalar  $\eta$  (the local alignment scale,  $f_{\text{cell}} = \eta f_{\text{ref}}$ ). The aligned matrix scales linearly,  $Z_{\text{al,cell}} = \eta Z_{\text{al}}$ ; the extinction matrix likewise,  $K_{\text{al,cell}} = \eta K_{\text{al}}$ ; the unaligned remainder scattering matrix is  $Z_{\text{unal}} = [C_{\text{sca,tot}} F_{\text{tot}} - \eta C_{\text{sca,ref}} F_{\text{ref}}] / (4\pi)$  in absolute units, whose  $\int Z_{11} d\Omega = C_{\text{sca,tot}} - \eta C_{\text{sca,ref}}$ ; and the unaligned population adds an isotropic  $C_{\text{ext,tot}} - \eta C_{\text{ext,ref}}$  to the extinction (its polarized and birefringent parts vanish). All of this is exact by the linearity of the size integral in  $f_{\text{align}}$ , so three stored integrals give the scattering optics of any  $\eta$  with no re-integration.

## 8.5 The library: two strictly separated layers

sed/src/scatmat\_aligned.f90 (module scatmat\_aligned\_mod) serves the table to a polarized RT host with the same lifecycle as the rest of the library: initialize once, query along the path.

**Initialization (serial, once).** load\_scatmat\_aligned(path, status) parses the table into module storage and free\_scatmat\_aligned releases it. The resident cost is dominated by scm\_Z, whose size is  $n_{\theta_i} n_{\theta_s} n_{\phi} \times 16 \times n_{\text{band}}$  doubles: for the production grid ( $19 \times 181 \times 37$ , five bands) that is about 81 MB, linear in the number of bands, with the  $F$  and  $K$  arrays adding under 0.1 MB. Queries are trilinear interpolations — a few hundred floating-point operations per scattering event, with no wavelength search once scatmat\_band has mapped the wavelength to a band. It is wired into sed\_init / build\_astrodust through an optional scatmat\_path argument — exactly the qpol\_path pattern of

§7., except that a `scatmat_path` that fails to load is an error (`status = 3`), because a host that asked for the table depends on it, whereas a missing polarized-optics table only disables polarized emission. This layer is not thread-safe and is called from serial code.

**Path queries (pure reads, OpenMP-safe).** Six routines write only their own outputs and read the loaded arrays, so they are safe to call concurrently from photon threads:

- `scatmat_band(lambda_um, iband, exact)` — nearest stored band, so the host path takes an integer and no wavelength search runs per event;
- `extinction_matrix_aligned(iband, theta_i, eta, kmat)` — the  $4 \times 4$  matrix with diagonal  $\eta C_{\text{ext,al}}$ ,  $K(1,2) = K(2,1) = \eta C_{\text{pol,al}}$ ,  $K(3,4) = \eta C_{\text{bir,al}}$ ,  $K(4,3) = -\eta C_{\text{bir,al}}$ ,  $\theta_i$  folded into  $[0, 90^\circ]$  because  $K$  is even under the equatorial reflection;
- `mueller_matrix_aligned(iband, theta_i, theta_s, phi, z)` — trilinear interpolation plus the symmetry reconstruction of §8.2, at the reference alignment ( $\eta = 1$ ; the caller scales);
- `mueller_matrix_random(iband, big_theta, f_tot, f_ref)` — the two six-element random-orientation matrices;
- `mueller_matrix_total(iband, theta_i, theta_s, phi, eta, z)` — the recommended host call: the *absolute* combined matrix  $z = \eta Z_{\text{al}} + L(\pi - \sigma_2) F_{\text{unal}}(\Theta) L(-\sigma_1) / (4\pi)$  with  $F_{\text{unal}} = C_{\text{sca,tot}} F_{\text{tot}} - \eta C_{\text{sca,ref}} F_{\text{ref}}$ ,  $\cos \Theta = \cos \theta_i \cos \theta_s + \sin \theta_i \sin \theta_s \cos \phi$  and the meridional rotation angles  $\sigma_1, \sigma_2$  from `meridional_scattering_angles` (closed `atan2` forms, poles included), so a host need not reassemble the mixture from the arrays;
- `scattering_cross_sections(iband, theta_i, eta, csca_aligned, csca_unaligned [, csca_pol_aligned])` with  $\text{csca\_unaligned} = C_{\text{sca,tot}} - \eta C_{\text{sca,ref}}$  and the optional  $\text{csca\_pol\_aligned} = \eta C_{\text{sca,pol,al}}(\theta_i) = \eta \int Z_{12} d\Omega$ , the aligned matrix's polarized scattering cross section (Peest et al. 2023; the random remainder integrates to zero), on the same  $[0, 90^\circ]$  fold.

`scm_csca_pol_al(nti, nband)`, the polarized scattering cross section  $\int Z_{12} d\Omega$  that `csca_pol_aligned` returns, is formed at load time with the generator's own closure quadrature — the same rule that closes  $\int Z_{11} d\Omega$  to  $C_{\text{sca,al}}$  — so it needs no extra file column. The loaded grids and arrays are also exposed public, protected, so a host that prefers to inline its own lookups can take them once at startup and never call back. The whole cell dependence enters through just two scalars the host already holds:  $\eta$ , and  $\theta_i = \arccos(\hat{k} \cdot \hat{B})$  from one dot product.

**The alignment guard (status 4).** The  $K$  and  $Z$  arrays were integrated over size once, under the profile recorded in the header; the sanctioned runtime variation is  $\eta$ , not a change of profile. So when a loaded scattering table is present and the host calls `dust_set_alignment` or `dust_set_alignment_profile` with a profile that differs from the recorded one, the setter returns a *new* status code 4. It is non-fatal: the  $f_{\text{align}}$  update still completes (the emission and  $C_{\text{pol,ext}}$  channels re-integrate live and honor the new profile), and the code only signals that the loaded scattering matrices no longer correspond to the model's alignment. A different profile requires regenerating the table (`run_scatmat_aligned.x`

profile=FILE), which takes minutes. A tabulated profile never matches the recorded analytic one and so always raises the flag.

## 8.6 Birefringence through dust\_extinction

dust\_extinction gained an optional Cbir\_ext output, the  $f_{\text{align}}$ -weighted size integral of the jori table's 4th-block birefringence,

$$C_{\text{bir,ext}}(\lambda) = \sum_a dn_{\text{Ad}}(a) C_{\text{bir,ext}}(\lambda, a) f_{\text{align}}(a),$$

mirroring Cpol\_ext exactly. It is zero when the loaded table carries no 4th block (as the release table does not), so the default path is unchanged, and it is the same birefringence the aligned  $K$  block carries at  $\theta_i = 90^\circ$ , reached here without loading the aligned scattering table. This closes the gap noted in the previous release, that dust\_extinction did not yet return the birefringence.

## 8.7 Verification

Every anchor of the plan passes (Table 10). The engine anchors (A–H) run at single sizes through compare\_scmat\_aligned.x — anchor H adds the orientation-averaged amplitude engine at general geometry ( $\theta_i = 40^\circ$ ,  $\theta_s = 70^\circ$ ,  $\phi \in \{30, 120, 250\}^\circ$ ) against  $L(\pi - \sigma_2)F(\Theta)L(-\sigma_1)$ , which a wrong  $\sigma$  sign fails at order unity to order 30. The library checks, including the failing-tolerance certification tests [3a]–[3e] of test\_scmat\_aligned.x — the absolute  $4\pi$  closure, the  $\int Z_{12}$  polarized cross section, the mueller\_matrix\_total closure, and the  $\theta_i = 0$  rotation sign — run on the production table.

**Cost of the production run.** The five UBVRI bands (0.36, 0.44, 0.55, 0.64,  $0.79\mu\text{m}$ ) take about 8 minutes at 32 threads (about 26 minutes at 8); the file is 138 MB, 38 MB gzipped. The existing regressions are unchanged: calc\_pol\_ext reproduces the release to a median of 0.0296% as before, and every sed, mc and tmatrix target builds clean.

## 8.8 Which wavelengths the table carries

The shipped table holds five bands, 0.36, 0.44, 0.55, 0.64 and  $0.79\mu\text{m}$  (BANDS in run\_scmat\_aligned.f90), on the grid  $\theta_i = 0(5)90^\circ$  (19)  $\times$   $\theta_s = 0(1)180^\circ$  (181)  $\times$   $\phi = 0(5)180^\circ$  (37). That is the intended scope: polarized transfer is essentially never run on a whole wavelength grid, and five optical bands cover the update this work is for. What matters is that a wavelength one later needs can be computed, not that everything is precomputed now.

run\_scmat\_aligned.x does already take wavelengths on its command line (`./run_scmat_aligned.x 0.44 0.55 0.79`), so a different set is one run away. Three limitations of that path should be stated plainly rather than discovered:

- **The output stem is fixed** (`f_stem//'.dat'`), so a custom run *overwrites* the shipped five-band table. Move or rename the shipped file first.

- **There is no merge mode.** A band cannot be added to an existing table; every requested band is recomputed from scratch into a new file (the five together are the 8-minute run above).
- **There is no size-window split and no large- $x$  certification.** The driver has no  $ja=$  equivalent of §7.6, and it applies only the plain  $X\_LARGE = 50$  cut, without the two-setting cross-check of §7.8.

None of these is hard to lift; none of them is done, and the wavelengths in hand have not required it.

## 8.9 Why the library reads tables and does not solve

The library never runs a T-matrix solve at run time. The reason is not the size of the code — the engine is already in the tree and the drivers link it — but that convergence cannot be certified at  $x \gtrsim 55$  from inside a transport loop: as §7.8 records, Mishchenko's IERR can report success on a node that is 50% wrong, and the check that catches it costs a second solve at different settings. A silently wrong optic is worse than a missing one, and the transport loop has nowhere to report a failure to. The tables are generated offline, where a node can be rejected and the rejection counted in the file header. This is the current policy rather than a permanent decision: the T-matrix core itself is expected to be revised, and a certification cheap enough to run inline would change the trade.

## 9. Status codes, and the relation to v1.00

### 9.1 `sed_init/build_astrodust` status codes

Both take an optional status; 0 is success, and when the argument is absent the readers keep their message-and-stop behavior instead. The v1.20 codes are those of Table 11.

An *absent* extreme-ultraviolet companion table is not an error: `build_Cpol` reports it on `stderr` and leaves that band zero (§7.7), and the build succeeds. The two codes that band can raise are 9, when a companion table that *was* read does not cover the requested grid, and 8, when the main polarized table and the model grid disagree node for node. An implicitly defaulted main polarized table that cannot be read degrades to zero polarized optics; only one named through `qpol_path` raises 4.

`dust_set_alignment` and `dust_set_alignment_profile` have their own independent numbering: 1, 2, 3 for an out-of-range `a_align`, `alpha_align` or `f_max`, and 4 for the non-fatal alignment guard of §8.5. Do not read a 4 from a setter as the `qpol_path` failure above.

### 9.2 v1.20 is v1.00 plus polarization

The two maintained branches are a superset relation, not a fork. Everything scalar is identical: the same builders (`build_astrodust`, `build_dl07`, `build_zubko`, `build_from_files`), the same `dust_emission` and `dust_extinction`, the same `lam_min` extreme-ultraviolet extension, and the same `calc_kext.x` products. What v1.20 adds is the polarized material: `Cpol`, `Cpol_ext`, `Cbir_ext`,  $f_{\text{align}}$ , the aligned scattering table and its library layer. The only incidental difference is the status numbering — what v1.00 calls 3 and 4 (the dielectric function load and the `lam_min` floor) are 6 and 7 here, because 3, 4 and 5 were taken by the polarized paths. A scalar-only v1.20 run

(load\_polarized\_optics = .false.) returns the scalar  $C_{\text{ext}}$ ,  $C_{\text{abs}}$ ,  $C_{\text{sca}}$ ,  $g_{\text{bar}}$  and the total SED bit-identical to the polarized build at zero alignment.

## 10. Limits and extension paths

These are limits of the published optics, not of the code.

**Birefringence is now computed; the release table has none.** The release table stores  $Q_{\text{ext}}$ ,  $Q_{\text{abs}}$  and  $Q_{\text{sca}}$  only, no phase retardation, so from it  $C_{\text{circ}} = 0$ . The first-principles route of §7. already evaluates the fixed-orientation amplitude matrix (Mishchenko’s AMPL, in `tmatrix/src/fixed_orientation_amplitude.f90`) and used only its imaginary forward part, for extinction. Keeping the real part as well costs nothing — it comes from the same amplitude call — and gives the birefringence, the phase retardation that couples  $U$  and  $V$ . It is stored as an optional 4th block of the regenerated table,  $Q_{\text{re}}(\text{jori}) = (4\pi/k) \text{Re}[S_{\text{fwd}}]/(\pi a^2)$ , the real twin of  $Q_{\text{ext}}$ , from which  $C_{\text{bir}} = \frac{1}{2}(Q_{\text{re},3} - Q_{\text{re},2})\pi a^2$ . The reader exposes it as `qbir_ext` with a `has_bir` flag that stays `.false.` for an older 3-block table, so nothing about the default path changes. No astrodust reference for the birefringence exists, so it is certified internally by Kramers-Kronig:  $\text{Re}$  and  $\text{Im}$  of the forward-amplitude difference are a Hilbert pair, and the directly computed birefringence agrees with the Hilbert transform of the already-validated dichroism to a median  $\sim 0.1\%$  over the interior of the grid (the transform itself self-tests to  $2 \times 10^{-5}$  on a Lorentz oscillator; the residual is the finite-range truncation). What remains is to return it through `dust_extinction` and to consume it in the  $U \leftrightarrow V$  transfer term, both on the host side.

**The aligned scattering matrix is now computed.** The release gives the integrated  $Q_{\text{sca}}$  with no angular information, so scattering by an aligned spheroid could not be modeled from it. It is now modeled from first principles: the fixed-orientation Mueller matrix  $Z(\theta_i; \theta_s, \phi)$  of the aligned astrodust spheroid is built from the same fixed-orientation amplitude (AMPL) used above for extinction, size-integrated over the astrodust distribution with the alignment weight, and paired with the  $4 \times 4$  extinction matrix  $K(\theta_i)$  from the forward amplitudes. It is written for five optical bands by `run_scattermat_aligned.x` and served to a polarized RT host through a two-layer library API. §8. is the full account, and §8.8 covers the wavelength coverage and what computing a different band involves.

**The random-orientation scattering matrix does exist, in five optical bands.** This is a different object from the previous paragraph and should not be confused with it. Dropping the alignment, a randomly oriented spheroid with a plane of symmetry has six independent scattering-matrix elements rather than sixteen, and those *are* computed in the tree: `tmatrix/driver/run_scattermat.f90` accumulates the single-size expansion coefficients over the HD23 size distribution with  $C_{\text{sca}} dn/n_{\text{H}}$  weighting, expands them onto the scattering-angle grid with `SCATMAT_FROM_MOMENTS` (Mishchenko’s MATR, in module `scattering_matrix_expansion`), and writes `tmatrix/output/scattermat_astrodust_P0.20_Fe0.00_1.400.dat`: one block per wavelength, 181 rows at  $\theta = 0, 1, \dots, 180^\circ$ , with columns  $F_{11}, F_{22}, F_{33}, F_{44}, F_{12}, F_{34}$ . It answers what an unaligned grain population scatters; it does not answer the aligned question above.

The stored wavelengths are 0.36, 0.44, 0.55, 0.64 and  $0.79 \mu\text{m}$ , roughly *UBVRI* — the same five bands the aligned table carries, and for the same reason (§8.8). Five bands are the deliverable, not a sweep left half-finished: polarized transfer

is run at a few wavelengths, the  $V$  band first with one or two further optical bands possible, and the whole 1129-point grid would cost about 1.7 hours spent almost entirely on wavelengths nothing is waiting for. What the design owes is a way to compute a wavelength when one is wanted, and that costs a run rather than a code change: `./run_scattermat.x` all does the full grid, or a list of wavelengths does a subset, at about 5.5s per wavelength.

**Far infrared needs neither.** In that regime the albedo is essentially zero, so the transfer requires only the extinction matrix and the emission vector, both fully determined by what is exported. The HAWC+  $154\mu\text{m}$  comparison lies inside this regime, so the far-infrared milestone is reachable with the present data and should be attempted before any scattering work.

**Astroduct only.** DL07 and Zubko have no polarized optics. `build_dl07` releases any polarized arrays left by a previous astroduct initialization, so `lamI_pol` returns zeros rather than stale values.

## 10.1 Cost of build\_Cpol

Loading the orientation-resolved table is the most expensive part of model construction (Table 12).

The orientation-resolved arrays `qext_j`, `qabs_j` and `qsca_j` are deallocated as soon as the combinations of Eqs. (1)–(2) have been formed. They are intermediate quantities: nothing downstream of `build_Cpol` reads them, and every later step uses the five derived arrays only. Releasing them removes 13.7 MB, about two-thirds of what the table would otherwise hold, and leaves 7.6 MB resident for the rest of the run. The three arrays are private to `q_table_jori`, so no consumer can depend on their surviving the build.

Measured on `calc_sed.x` astroduct, the maximum resident set size falls from 51232 kB to 40088 kB, a reduction of about 10.9 MB. The peak does not fall by the full 13.7 MB because it is a high-water mark: once the orientation-resolved arrays are gone, the later SED solution stages set the peak instead. The steady-state occupancy just after `sed_init` does drop by the full 13.7 MB.

## 11. Reproducing the numbers

From SEDust/sed/:

```
make calc_polext.x # standalone polarized-extinction check
./calc_polext.x # -> output/polarized_extinction_ours.dat
                  # lambda, ours, reference, ratio

make calc_kext.x # full-pipeline extinction table, all dust models
./calc_kext.x astroduct # -> ../data/astroduct/kext_astroduct_MW.dat, 8 columns
                  # ./calc_kext.x astroduct euv (EUV grid)
                  # ./calc_kext.x astroduct 0.03 (lam_min [um])
                  # ./calc_kext.x dl07 | zubko | from_files ...

make libsedust.a # library, for the emission path
```

`calc_polext.x` prints the median and maximum deviation against the release directly, restricted to  $\lambda \geq 0.11 \mu\text{m}$ . Column 8 of the extinction table should reproduce it to the ASCII precision quoted in Table 7; if the two disagree by more than that, the  $\log a$  interpolation in `build_Cpol` or the size grid it interpolates onto is the place to look, since that is the only step separating the two routes.

For the emission fractions of Table 8, build a model with `build_astrodust`, evaluate `J_Mathis` at  $U = 1.585$ , and call `dust_emission` with `lamI_pol` requested; the example in §5.4 of the user manual is the complete recipe.

## 12. References

- Draine, B. T., & Hensley, B. S. 2021b, ApJ, 919, 65 (DH21b). 2021ApJ...919...65D
- Hensley, B. S., & Draine, B. T. 2023, ApJ, 948, 55 (HD23). 2023ApJ...948...55H
- Peest, C., Camps, P., Stalevski, M., Baes, M., & Siebenmorgen, R. 2017, A&A, 601, A92. 2017A&A...601A..92P
- Peest, C., et al. 2023, A&A, 673, A112. 2023A&A...673A.112P
- Planck Collaboration 2020, A&A, 641, A12. 2020A&A...641A..12P
- Reissl, S., Wolf, S., & Brauer, R. 2016, A&A, 593, A87 (POLARIS I). 2016A&A...593A..87R
- Seon, K.-I. 2018, ApJ, 862, 87. 2018ApJ...862...87S

| File / entity                                | Responsibility                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|----------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| sed/src/q_table_jori.f90                     | The reader. Decompresses, parses the three quantity blocks, validates (monotone axes, finite values, no trailing data), forms Eqs. (1)–(2), and supplies $f_{\text{align\_hd23}}(a)$ , Eq. (4). Publishes $q_{\text{pol\_ext}}$ , $q_{\text{pol\_abs}}$ and $q_{\text{ran\_}}$ ; the orientation-resolved $q_{\text{ext\_j}}/q_{\text{abs\_j}}/q_{\text{sca\_j}}$ are private and are deallocated once the combinations have been formed.                                                                                                                                                                |
| sed_astrodust.f90: build_Cpol                | Interpolates $Q_{\text{pol}}$ from the 169-node table axis onto the size-distribution grid in $\log a$ , exactly as $C_{\text{abs}}$ and $C_{\text{sca}}$ are, multiplies by $\pi a_{\text{eff}}^2$ , and fills $f_{\text{align\_ad}}$ . Checks that the polarized and $Q$ wavelength grids agree node for node.                                                                                                                                                                                                                                                                                         |
| sed_astrodust.f90: Cpol, Cpol_ext, falign_ad | Module-level arrays, (NLAM, NA) and (NA). Cpol is the <i>absorption</i> one and drives polarized emission; Cpol_ext is its extinction twin and drives dichroic extinction.                                                                                                                                                                                                                                                                                                                                                                                                                               |
| dust_model_mod.f90: grain_pop_t              | Carries $C_{\text{pol}}(:, :)$ and $f_{\text{align}}(:)$ . Both unallocated is how the solver recognizes an unpolarized population. Also carries the extinction-side $C_{\text{pol\_ext}}(:, :)$ and $g_{\text{sca}}(:, :)$ , read only by dust_extinction.                                                                                                                                                                                                                                                                                                                                              |
| sed_astrodust.f90: set_pop                   | Optional Cpol_in, falign_in. Both must be supplied together; supplying neither leaves the population unpolarized. The extinction-side C_sca_in, Cpol_ext_in and g_sca_in are optional and independent; a population that omits them contributes zero to those terms of the size integral.                                                                                                                                                                                                                                                                                                                |
| sed_astrodust.f90: sed_grain_loop            | Optional Cpol_pop, falign_pop, Jpol, again all three or none (do_pol). Accumulates the polarized channel alongside Jout in every solver branch (§5.1).                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| sed_astrodust.f90: dust_emission             | Optional lamI_pol, last argument. Sums Jpol over the populations that carry polarized optics and applies the same unit convention as lamI_total.                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| sed_astrodust.f90: dust_extinction           | The extinction accessor, per H atom on m%lam. It integrates $C_{\text{pol,ext}} f_{\text{align}}$ (and the birefringent twin) over the size distribution of every population — these are the same numbers column 8 of the table carries, without the file. $C_{\text{abs}}$ , $C_{\text{sca}}$ and $\langle \cos \theta \rangle$ it later stopped integrating; it reads them from the precomputed kext table instead, while the polarized outputs stay computed because their $f_{\text{align}}$ weight is runtime state (§5.7). The first-principles integral of all six is size_integrated_extinction. |
| sed/src/calc_polext.f90                      | Standalone check: computes the polarized extinction directly from $q_{\text{pol\_ext}}$ and the release size distribution, and compares against the published polarized_extinction.dat. Independent of sed_init.                                                                                                                                                                                                                                                                                                                                                                                         |
| sed/src/calc_kext.f90                        | The extinction-table driver, one executable for every dust model (./calc_kext.x astrodust   dl07   zubko   from_files). In its astrodust mode it writes column 8, C_polext/H, of data/astrodust/kext_astrodust_MW.dat as $\sum_a dn_{\text{Ad}}(a) C_{\text{pol,ext}}(\lambda, a) f_{\text{align}}(a)$ . Only that mode writes the column, because only the astrodust builder fills Cpol_ext.                                                                                                                                                                                                            |

Table 4: Responsibility of each file touched by the polarization update.



| branch      | context                      | accumulates into | pattern       |
|-------------|------------------------------|------------------|---------------|
| 'draine'    | equilibrium grain            | Jpol_local       | thread-local  |
| 'draine'    | stochastic, sum over $P(T)$  | Jpol_local       | thread-local  |
| 'heuristic' | equilibrium grain            | Jpol             | <b>direct</b> |
| 'heuristic' | stochastic, sum over $P(T)$  | Jpol             | <b>direct</b> |
| 'qm'        | QM solve, rescaled           | Jpol_local       | thread-local  |
| 'qm'        | GD fallback after QM failure | Jpol_local       | thread-local  |
| 'qm'        | equilibrium grain            | Jpol_local       | thread-local  |
| 'equil'     | every grain                  | Jpol             | direct        |

Table 5: Where the polarized emission is accumulated. The two boldface rows are the ones a search for the Jout\_local pattern does not find.

| band                                        | median | maximum |
|---------------------------------------------|--------|---------|
| far infrared ( $\lambda > 30 \mu\text{m}$ ) | 0.022% | 3.46%   |
| restricted to $a \leq 0.5 \mu\text{m}$      |        | 1.98%   |
| mid infrared                                | 0.033% |         |
| near infrared                               | 0.038% |         |
| optical                                     | 0.080% |         |
| ultraviolet                                 | 0.21%  | 9.53%   |

Table 6: Disagreement between the trace average and the exact orientation average of  $Q_{\text{abs}}$ .

| Check                                          | Result                                  |
|------------------------------------------------|-----------------------------------------|
| Column 8 vs. released polarized_extinction.dat | median 0.0296%, max 0.9422%, 992 points |
| Column 8 vs. calc_polext.x                     | agree to $7 \times 10^{-8}$ relative    |
| Columns 1–7 of kext_astrodust_MW.dat           | byte-identical to the pre-change file   |
| astrodust_irem_ours_S1/S2/PAH.dat              | byte-identical                          |
| dl07_sed_ours_mw31_60.dat                      | byte-identical                          |
| Build                                          | zero warnings                           |

Table 7: Verification of the polarized extinction path and regression of everything else.

| $\lambda$                                   | fraction |
|---------------------------------------------|----------|
| $12\,\mu\text{m}$                           | 0.00%    |
| $100\,\mu\text{m}$                          | 13.67%   |
| $154\,\mu\text{m}$                          | 17.15%   |
| $350\,\mu\text{m}$                          | 18.79%   |
| $850\,\mu\text{m}$                          | 19.15%   |
| median over $300\text{--}3000\,\mu\text{m}$ | 19.18%   |

Table 8: Intrinsic polarization fraction of the emitted spectrum, before any geometric factor.

| Regime   | $x$    | Quantity                     | median  ours/HD23 – 1 |
|----------|--------|------------------------------|-----------------------|
| Rayleigh | < 0.1  | $Q_{\text{ran}}(\text{abs})$ | $2.3 \times 10^{-4}$  |
|          |        | $Q_{\text{pol}}(\text{abs})$ | $3.7 \times 10^{-4}$  |
| T-matrix | 0.1–50 | $Q_{\text{ran}}(\text{ext})$ | $2.1 \times 10^{-4}$  |
|          |        | $Q_{\text{pol}}(\text{ext})$ | $8.4 \times 10^{-4}$  |
|          |        | $Q_{\text{ran}}(\text{abs})$ | $1.9 \times 10^{-4}$  |
|          |        | $Q_{\text{pol}}(\text{abs})$ | $7.6 \times 10^{-4}$  |
|          |        |                              |                       |
| GO       | > 50   | $Q_{\text{ran}}(\text{abs})$ | $9.3 \times 10^{-2}$  |
|          |        | $Q_{\text{pol}}(\text{abs})$ | $3.3 \times 10^{-1}$  |

Table 9: Node-by-node reconstruction against the HD23 release table. The geometric-optics dichroism is at  $\sim 0.75$  of the release with matching sign and trend; its size-parameter region carries no astrodust weight.

|                                    | Check                                                          | What it certifies                                                                                                                                           | Result                                    |
|------------------------------------|----------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------|
| <i>Engine (single size)</i>        |                                                                |                                                                                                                                                             |                                           |
| A                                  | closure $\int Z_{11} d\Omega$ vs $C_{\text{sca}}(\text{jori})$ | phase-matrix normalization                                                                                                                                  | $2.7 \times 10^{-7}$                      |
| B                                  | optical theorem vs $C_{\text{ext}}(\text{jori})$               | forward-amplitude geometry                                                                                                                                  | $2 \times 10^{-14}$                       |
| C                                  | orientation average vs GSP $F(\Theta)$                         | the bilinears, incl. $F_{12}, F_{34}$ signs                                                                                                                 | $1.4 \times 10^{-6}$                      |
| D                                  | dipole vs closed form                                          | Rayleigh implementation                                                                                                                                     | $2 \times 10^{-16}$                       |
|                                    | dipole vs T-matrix at $x \sim 0.1$                             | regime continuity, $ \delta Z /Z_{11}$                                                                                                                      | 0.5–1.1%                                  |
| E                                  | $\phi$ mirror                                                  | azimuth symmetry                                                                                                                                            | $3.5 \times 10^{-7}$                      |
|                                    | equatorial mapping                                             | $z \rightarrow -z$ symmetry                                                                                                                                 | $8.8 \times 10^{-16}$                     |
| G                                  | Rayleigh $K(\theta_i)$ vs $\sin^2$ law                         | extinction-matrix angle law (exact)                                                                                                                         | $3.3 \times 10^{-16}$                     |
|                                    | T-matrix $K(\theta_i)$ vs $\sin^2$ law                         | at $x = 0.137$                                                                                                                                              | $6.3 \times 10^{-4}$                      |
| H                                  | general geometry                                               | amplitude engine vs $L(\pi - \sigma_2)F(\Theta)L(-\sigma_1)$ , $\theta_i=40^\circ, \theta_s=70^\circ$ ;<br>wrong $\sigma$ sign fails O(1)–O(30)             | $4.1 \times 10^{-5}$                      |
| <hr/>                              |                                                                |                                                                                                                                                             |                                           |
| <i>Generator (size-integrated)</i> |                                                                |                                                                                                                                                             |                                           |
| F                                  | $K$ vs independent jori integrals                              | $C_{\text{ext}}$                                                                                                                                            | $1.6 \times 10^{-6}$                      |
|                                    |                                                                | $C_{\text{pol}}$                                                                                                                                            | $2.0 \times 10^{-5}$                      |
|                                    |                                                                | $C_{\text{bir}}$                                                                                                                                            | $1.3 \times 10^{-7}$                      |
| <hr/>                              |                                                                |                                                                                                                                                             |                                           |
| <i>Library (production table)</i>  |                                                                |                                                                                                                                                             |                                           |
|                                    | reader round-trip                                              | parse vs re-emit                                                                                                                                            | bitwise 0                                 |
|                                    | symmetry reconstruction                                        | unstored octants                                                                                                                                            | bitwise 0                                 |
|                                    | $\eta$ algebra                                                 | linear scaling                                                                                                                                              | residual 0                                |
|                                    | $K(90^\circ)$ vs jori vs dust_extinction                       | $C_{\text{pol}} / C_{\text{bir}}$                                                                                                                           | $1.6 \times 10^{-5} / 8.0 \times 10^{-5}$ |
|                                    | rt_example closures                                            | end-to-end consumption                                                                                                                                      | 0.997–1.001                               |
| [3a]                               | random absolute closure                                        | $(4\pi)^{-1} \int C_{\text{sca,tot}} F_{11} d\Omega = C_{\text{sca,tot}}$                                                                                   | $3.5 \times 10^{-4}$                      |
| [3b]                               | stored $Z_{11}$ re-integration                                 | vs the table’s $C_{\text{sca,al}}$ column (ASCII rounding)                                                                                                  | $3.3 \times 10^{-6}$                      |
| [3c]                               | $C_{\text{sca,pol}}$ vs $\int Z_{12} d\Omega$                  | routine vs direct; $ C_{\text{sca,pol}} /C_{\text{sca,al}} \leq 0.101$                                                                                      | exact                                     |
| [3d]                               | muellder_matrix_total closure                                  | $\int z_{11} = \eta C_{\text{sca,al}} + C_{\text{sca,tot}} - \eta C_{\text{sca,ref}}$ ; $\int z_{12} = \eta C_{\text{sca,pol}}$ ; $\eta \in \{0, 0.37, 1\}$ | $2.6 \times 10^{-4} / 7 \times 10^{-7}$   |
| [3e]                               | rotation sign at $\theta_i = 0$                                | $Z(0; \theta_s, \phi) = Z(0; \theta_s, 0)L(-\phi)$                                                                                                          | $6.8 \times 10^{-5}$                      |

Table 10: Verification of the aligned scattering matrix, its extinction matrix, and the library API. Engine and generator anchors are those of the plan's §5, plus anchor H (the amplitude engine at general geometry) and the failing-tolerance library certification tests [3a]–[3e] of `test_scattermat_aligned.x`; the library checks run on the shipped production table.

| status | Meaning                                                                                                                                                                                 |
|--------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1      | $Q$ -table load failed                                                                                                                                                                  |
| 2      | size-distribution load failed                                                                                                                                                           |
| 3      | aligned scattering table load failed (only when <code>scatmat_path</code> was supplied — a host that asked for it depends on it)                                                        |
| 4      | a polarized $Q$ table requested explicitly through <code>qpol_path</code> could not be read (an implicit default degrades gracefully instead)                                           |
| 5      | <code>load_polarized_optics = .false.</code> combined with an explicit polarized-optics path — a contradictory request                                                                  |
| 6      | astrodust dielectric function load failed (extreme-ultraviolet band only)                                                                                                               |
| 7      | <code>lam_min</code> below the astrodust dielectric function's own shortest wavelength (extreme-ultraviolet band only)                                                                  |
| 8      | the polarized $Q$ table's wavelength grid does not sit node for node on the long-wavelength end of the model's grid (wrong count, or a node that differs)                               |
| 9      | the extreme-ultraviolet band of the grid runs outside the wavelengths the companion polarized table covers (0.0124–0.0912 $\mu\text{m}$ )                                               |
| 10     | an explicitly named extinction table ( <code>kext_path</code> ) could not be read. It is 10 here and 5 elsewhere because this builder's polarized codes already occupy 3, 4, 5, 8 and 9 |
| 11     | <code>euv_tmatrix = .true.</code> but the spheroid optics of a band below the $Q$ table are not available — unreachable with the shipped table, which leaves no such band               |

Table 11: Status codes of `sed_init` and `build_astrodust` in v1.20.

| Item                                                                       | Size                               |
|----------------------------------------------------------------------------|------------------------------------|
| compressed file                                                            | 4.1 MB                             |
| scratch file when expanded                                                 | 17.2 MB                            |
| raw arrays <code>qext_j</code> , <code>qabs_j</code> , <code>qsca_j</code> | 13.7 MB, transient                 |
| derived <code>qpol_*</code> , <code>qran_*</code>                          | 7.6 MB, kept                       |
| resident after <code>sed_init</code>                                       | 7.6 MB                             |
| time                                                                       | about 1 s of the 3.1 s model build |

Table 12: Cost of loading the orientation-resolved table. The orientation-resolved arrays exist only between the read and the formation of the derived combinations.